

Physics-Informed Joint-Embedding Predictive Architectures for Multivariate Time-Series Anomaly Detection: Spectral, Transport, and Field-Theoretic Formulations

Andre Williams	Imadeldin Mahgoub	Hamid Akbarian
<i>Florida Atlantic University</i>	<i>Florida Atlantic University</i>	<i>Florida Atlantic University</i>
<i>Department of EECS</i>	<i>Department of EECS</i>	<i>Department of EECS</i>
<i>Boca Raton, US</i>	<i>Boca Raton, US</i>	<i>Boca Raton, US</i>
<i>andrewwilliam2018@fau.edu</i>	<i>mahgoubi@fau.edu</i>	<i>hakbaria@fau.edu</i>

Abstract—Unsupervised anomaly detection in multivariate telemetry is central to operational safety across aerospace systems, industrial processes, cyber-physical infrastructure, and smart environments. While reconstructive architectures such as TimesNet and TranAD achieve competitive continuous ranking performance (PR-AUC 0.2585–0.2612), their observation-space residuals destabilize automated streaming threshold calibration, generating over 1.05×10^6 false positives. To address this calibration failure, we evaluate Joint-Embedding Predictive Architectures (JEPAs) that forecast dynamics in regularized latent representation space. We introduce three physics-informed inductive regularizers grounded in dynamical systems theory, continuum mechanics, and field theory: (1) Operator-Entropy JEPA, which parameterizes a latent Koopman transition matrix and regularizes dynamic spectral entropy to prevent modal collapse; (2) Reynolds-Stress JEPA, which adapts hydrodynamic Reynolds decomposition to align second-order cross-channel fluctuation covariances; and (3) Potential-Flow JEPA, which parameterizes latent trajectory transitions via an irrotational velocity potential ($\hat{v} = \nabla\Phi$). To handle heterogeneous sensor arrays, the framework incorporates coordinate saliency gating for active channel isolation, dyadic wavelet shells for multi-scale frequency decomposition, and kurtosis-adapted Extreme Value Theory (EVT) for streaming threshold calibration. Evaluated on a matched benchmark of 45 continuous telemetry channels across 11 datasets, latent prediction suppresses false alarms while capturing true physical anomalies. Operator-Entropy JEPA achieves the highest Dataset-Macro Point-F1 (0.2066, a +43.2% relative improvement over TimesNet and +28.4% over TranAD) and is the only architecture to exceed the analytical All-Positive baseline (0.1893). In Pooled Micro Point-F1, Potential-Flow JEPA achieves 0.3458 (3.2 $\times$ higher than TranAD and 5.3 $\times$ higher than TimesNet) while reducing false positives by 8.0-fold against TimesNet (130,237 vs 1,047,144) and 3.5-fold against TranAD (455,128). These results show that latent predictive architectures yield better-conditioned residual distributions that substantially mitigate streaming threshold calibration collapse, with physical inductive priors providing targeted performance gains in specialized dynamical regimes.

Index Terms—Multivariate Time-Series, Anomaly Detection, Joint-Embedding Predictive Architecture (JEPA), Physics-Informed Machine Learning, Koopman Operator Theory, Reynolds Stress, Potential Flow, Extreme Value Theory, Deep Learning.

I. INTRODUCTION

Autonomous anomaly detection in multi-sensor telemetry monitors operational health across aerospace systems, Earth observation, cyber-physical infrastructure, and wearable digital healthcare [1], [2]. In remote sensing deployments (such as the NASA Soil Moisture Active Passive (SMAP) satellite, Mars Science Laboratory (MSL) rovers, industrial IoT facilities, and wearable biometric sensor arrays), dozens to hundreds of continuous telemetry channels concurrently track electrical power buses, thermal gradients, mechanical vibrations, and kinematics [3]. When sensor faults or operational breakdowns emerge at the edge, monitoring systems must identify dynamical anomalies before hardware damage occurs [4].

A. Challenges in Modeling Multi-Sensor Telemetry

From a physical perspective, multi-sensor telemetry reflects the state trajectory of an underlying continuous physical system $\dot{x}(t) = f(x(t))$. However, attempting to model dynamical equations directly in raw observation space presents four main challenges on edge devices:

- 1) *Numerical Instability in Empirical Differentiation*: Estimating instantaneous time derivatives $\dot{x} \approx \frac{\Delta x}{\Delta t}$ directly on discrete, noise-corrupted sensor streams amplifies high-frequency measurement noise, destabilizing numerical trajectory tracking.
- 2) *Partial Observability & Unmeasured Internal States*: Remote sensing telemetry captures only a noisy surface projection of unmeasured thermodynamic, electrical, and environmental states, violating full-state observability.
- 3) *High-Dimensional Cross-Sensor Coupling*: In multi-sensor arrays ($K \geq 50$), cross-channel dependencies and non-linear interactions cannot be adequately represented by rigid functional forms in raw signal space.
- 4) *Non-Stationary Downlink & Operational Drift*: Edge IoT and satellite systems undergo discrete operational mode transitions (e.g., orbital day/night thermal cycling, thruster firings, Parkinsonian gait freezing, cloud server garbage collection) that cause continuous distribution drift.

B. Limitations of Deep Anomaly Detectors

Deep learning architectures provide data-driven alternatives to explicit equation modeling, yet they face three primary limitations:

- 1) *Signal-Level Noise Memorization*: Generative models (e.g., autoencoders, VAEs, masked Transformers) minimize point-level L_1/L_2 errors in raw signal space [3], [5], [6]. Because reconstruction loss penalizes every discrepancy equally, autoencoders allocate capacity to fitting high-frequency sensor jitter rather than learning nominal dynamical transitions.
- 2) *Limitations of Synthetic Contrastive Perturbations*: Contrastive frameworks such as Neural Contextual Anomaly Detection (NCAD) [7] train encoders via Contextual Outlier Exposure (COE) using heuristic synthetic perturbations (e.g., Gaussian noise spikes, sequence splicing). However, real physical faults manifest as subtle deviations from nominal dynamics. Synthetic injections distort embedding geometry, inducing artificial decision boundaries that generate false alarms on continuous telemetry.
- 3) *Scale Distortion and Threshold Fragility in Observation Space*: In high-dimensional observation space, even after per-channel standardization, Euclidean reconstruction errors ($\|x - \hat{x}\|_2^2 = \sum_k (x_k - \hat{x}_k)^2$) accumulate stochastic residual variance across dozens of irrelevant or noisy sensors. Under automated streaming thresholding such as Peaks-Over-Threshold (SPOT), this residual accumulation yields heavy, unstable tail distributions that distort extreme-value calibrators, triggering false alarm flooding in operational telemetry.

C. Our Approach: Physics-Informed Joint-Embedding Predictive Architectures

To eliminate signal-space noise memorization without relying on synthetic contrastive perturbations, we formulate a family of *Physics-Informed Joint-Embedding Predictive Architectures* (PI-JEPA) for multi-sensor telemetry. Instead of reconstructing raw, jitter-corrupted sensor observations $\hat{x} \approx x$, a JEPA maps past observations x_{ctx} into an abstract representation space and predicts target representations z_{tgt} directly [8], [9]. By operating in latent space, the architecture discards high-frequency stochastic sensor noise that lacks long-range predictability.

However, unconstrained latent prediction alone can be underspecified in multi-sensor telemetry. To incorporate physical domain priors into self-supervised representation learning, we introduce three physics-informed JEPA architectures featuring inductive dynamical regularizers:

- 1) **Operator-Entropy JEPA (Koopman Spectral Mechanics)**: Parameterizes a latent linear Koopman transition matrix regularized by dynamic spectral entropy $\mathcal{H}_{\text{spectral}}(\mathbf{K})$ to prevent dimensional collapse and preserve linear modal expressivity in multi-sensor dynamics.
- 2) **Reynolds-Stress JEPA (Continuum Mechanics & Covariance Alignment)**: Inspired by hydrodynamic Reynolds decomposition ($z = \bar{z} + z'$), this architecture aligns the second-order cross-channel fluctuation covariance tensors

$\mathbf{R}_{ij} = \mathbb{E}[z'_i z'_j]$. Structural anomalies disrupt multi-sensor covariance balance.

- 3) **Potential-Flow JEPA (Conservative Field Theory)**: Enforces conservative latent trajectory transitions by parameterizing the transition field as the gradient of a scalar potential ($\hat{v} = \nabla \Phi$). Because gradient fields are irrotational by construction, non-conservative dynamical anomalies produce large predictive discrepancies.

To stabilize learning across heterogeneous high-dimensional sensor arrays, all three architectures incorporate Coordinate Saliency Gating to isolate active channels, Dyadic Wavelet Shells to decouple high-frequency transients from baseline drift, and a Kurtosis-Adaptive Extreme Value Theory (EVT / SPOT) calibrator that modulates tail thresholds against non-Gaussian kurtosis.

D. Contributions

The contributions of this work are as follows:

- *Physics-Informed Joint-Embedding Framework*: We formulate an unsupervised anomaly detection framework that embeds physical regularizers (spectral dissipation, turbulent momentum transfer, conservative flow) directly into latent representations, avoiding raw-space noise reconstruction.
- *Physics-Informed JEPA Formulations*: We derive and implement three physics-informed architectures (Operator-Entropy JEPA, Reynolds-Stress JEPA, and Potential-Flow JEPA) and benchmark them alongside an unconstrained control (TS-JEPA).
- *Inductive Architectural Components*: We introduce Coordinate Saliency Gating for cross-channel fault localization, Dyadic Wavelet decomposition for multi-resolution temporal analysis, and Kurtosis-Adaptive EVT tail calibration to provide automated streaming alarms without transductive test tuning.
- *Multi-Tier Benchmark Evaluation*: We evaluate performance across 11 telemetry datasets (45 matched channels), reporting Dataset-Macro, Channel-Macro, Pooled Micro-F1, and Area Under the Precision-Recall Curve (PR-AUC) alongside an analytical All-Positive reference baseline. We show that latent prediction mitigates threshold calibration collapse, achieving up to **+43.2% relative improvement** in Dataset-Macro F1 and a **3.2× to 5.3× gain** in Pooled Micro-F1 while reducing false positives by 8.0-fold against TimesNet (130,237 vs 1,047,144) and 3.5-fold against TranAD (455,128).

The remainder of this paper is structured as follows. Section II reviews related work in reconstructive, contrastive, and joint-embedding anomaly detection. Section III details the 11-dataset multi-sensor telemetry benchmark suite. Section IV formalizes the PI-JEPA Triad, coordinate saliency gating, dyadic wavelets, and kurtosis-adaptive EVT. Section V establishes the experimental protocol and multi-tier evaluation framework. Section VI evaluates architectural ablations and inductive layer contributions. Section VII provides domain-by-domain empirical analysis and qualitative physical case studies. Finally, Section VIII concludes the paper.

II. RELATED WORK

Unsupervised time-series anomaly detection spans dynamical equation discovery, reconstructive generative modeling, contrastive representation learning, and predictive world models. This section reviews these approaches and positions our Physics-Informed JEPA framework.

A. Dynamical Systems Modeling and Koopman Operators

The physical modeling of multi-sensor telemetry has long sought to capture system behavior as continuous dynamical trajectories governed by differential equations $\dot{x}(t) = \mathbf{f}(x(t))$ [10].

- *Neural Ordinary Differential Equations (Neural ODEs)* [11]: Parameterize continuous velocity fields using neural networks $\frac{dx}{dt} = f_\theta(x(t), t)$ and solve trajectories via numerical ODE integrators. While Neural ODEs learn flexible vector fields, fitting them directly in raw sensor space forces numerical solvers to integrate through high-frequency stochastic sensor noise, leading to numerical stiffness, slow training, and extreme sensitivity to measurement noise.
- *Koopman Operator Theory and Dynamic Mode Decomposition* [12]–[14]: Analyzes non-linear dynamical systems by learning linear transition operators acting on observable functions. Modern deep Koopman frameworks (e.g., Koopa [14]) learn observable embeddings for non-stationary forecasting. In anomaly detection, physical disruptions manifest as spectral decoherence or sudden changes in modal dissipation. Rather than operating on uncalibrated raw signals, our framework parameterizes latent linear transitions and monitors dynamic spectral entropy in representation space.

B. Reconstruction-Based Detection

The predominant deep learning approach in time-series anomaly detection relies on learning nominal patterns through raw signal reconstruction. Hundman et al. [3] introduced LSTM-based prediction coupled with non-parametric dynamic thresholding for NASA spacecraft telemetry (SMAP/MSL). Modern architectures incorporate specialized attention and frequency-domain transformations:

- *Anomaly Transformer* [5]: Introduces Anomaly-Attention to quantify the *Association Discrepancy* between local prior Gaussian kernel associations and global series self-attention. While effective for isolated point anomalies, it couples discrepancy with point-level waveform reconstruction.
- *TimesNet* [6]: Decomposes 1D multivariate time series into 2D temporal variations using Fast Fourier Transform (FFT) periodicity discovery and 2D Inception blocks. However, optimizing point-level 2D reconstruction makes it susceptible to over-fitting high-frequency sensor noise.
- *TranAD* [15]: Employs a two-phase adversarial Transformer with sub-series attention and focus scoring to amplify subtle reconstruction errors, but remains vulnerable to phase-drift in continuous dynamical systems.

Reconstructive methods share a common vulnerability: *signal-level reconstruction memorization*. Because loss functions optimize point-level L_1 or L_2 differences in raw signal space, reconstructive networks are biased toward fitting

high-frequency stochastic noise rather than learning invariant dynamical laws. In addition, during sustained anomalies, corrupted observations contaminate the context window, causing the reconstructor to adapt and reconstruct the anomaly as the "new normal."

C. Contrastive Learning and Representation-Based Detection

To avoid reconstructing raw noise, contrastive methods learn discriminative representations by separating nominal sequences from anomalous perturbations:

- *Neural Contextual Anomaly Detection (NCAD)* [7]: Uses Temporal Convolutional Networks (TCNs) trained with Contextual Outlier Exposure (COE). Artificial synthetic anomalies (Gaussian noise, point spikes, trend shifts) are injected into training windows to create negative pairs. However, synthetic injections distort the physical dynamical manifold and fail to generalize to subtle real-world physical faults (e.g., gradual thermal dissipation or gait freezing).
- *DCdetector* [16]: Replaces reconstruction with multi-scale dual-attention (patch-wise and channel-wise) contrastive representation learning, maximizing cross-scale representation consistency. While DCdetector eliminates artificial injection heuristics, contrastive objectives require careful scale selection and representation pooling.

D. Non-Contrastive Learning and Time-Series JEPAs

Joint-Embedding Predictive Architectures (JEPA), originally proposed by LeCun [8] and demonstrated in computer vision by I-JEPA [9] and V-JEPA [17], predict target representations from context in latent space. Non-contrastive regularizers such as Barlow Twins [18], W-MSE [19], and VICReg [20] prevent representation collapse without negative mining.

Recent works have extended joint-embedding predictive principles to time-series sequences:

- *Time-Series JEPAs*: SC-JEPA / MTS-JEPA [21] stabilizes latent predictive learning for time-series anomaly prediction using a multi-scale soft-codebook dictionary. T-SAR-JEPA [22] applies self-supervised latent prediction to Synthetic Aperture Radar (SAR) amplitude time series. CF-JEPA [23] employs mask-free forward prediction with asymmetric encoders for time-series representation learning. In cyber-physical and automotive domains, Girgis et al. [24] explore time-series JEPAs for predictive remote control over capacity-constrained IoT channels, while Fertig et al. [25] investigate JEPA embeddings for online automotive sensor monitoring.

However, existing time-series JEPA adaptations rely on generic unconstrained Euclidean loss or discrete codebooks without physical domain regularizers, leaving latent trajectories unanchored to physical conservation principles.

E. Physics-Informed Deep Learning and Structured Inductive Biases

Physics-Informed Neural Networks (PINNs) [26] integrate differential equations into deep models by penalizing residual PDE errors via automatic differentiation. While effective for

forward and inverse problems in fluid and solid mechanics, standard PINNs face two primary limitations in high-dimensional sensor telemetry:

- 1) *Gradient Pathologies in Multi-Objective Optimization*: Jointly optimizing data fidelity and PDE residual loss leads to stiff, competing gradient dynamics, often causing training instability or failure to converge in high noise regimes [27].
- 2) *Unobservable System Dynamics*: Complex IoT and aerospace telemetry (e.g., SMAP, MSL, and cyber-physical water networks) rarely follow a single closed-form PDE whose analytical derivatives can be evaluated across dozens of uncalibrated sensors.

Rather than enforcing analytical differential equations through auxiliary training penalties, our framework incorporates physical conservation, turbulent transport, and spectral dissipation directly into the latent representations and coordinate projections. This embeds physical inductive constraints directly, avoiding multi-loss balancing conflicts.

F. Positioning of the Proposed Framework

The proposed Physics-Informed JEPA framework bridges joint-embedding predictive learning and physical domain mechanics through four distinct structural designs:

- 1) *Latent Predictive Forward Learning Without Reconstruction Traps*: By forecasting future representations z_{tgt} rather than raw observations $\hat{x}$, our models circumvent high-frequency noise drowning and sensor jitter memorization that hobble reconstructive baselines (TimesNet, TranAD).
- 2) *Physics-Informed Inductive Formulations*: We introduce three physics-informed inductive formulations: Operator-Entropy (Koopman spectral dissipation), Reynolds-Stress (turbulent momentum transport), and Potential-Flow (conservative irrotational field theory), addressing cyclic, turbulent, and conservative physical dynamics.
- 3) *Multi-Resolution and Saliency Stabilizers*: We incorporate Coordinate Saliency Gating to prevent active channel dilution in high-dimensional arrays, and Littlewood-Paley Dyadic Wavelets to isolate transient impulses from low-frequency operational drift.
- 4) *Kurtosis-Adaptive Statistical Calibration*: We deploy Extreme Value Theory (EVT / SPOT) [28] modulated by dynamic sample kurtosis, establishing decision boundaries on operational telemetry without transductive test-label tuning.

III. MULTI-SENSOR IoT AND REMOTE SENSING BENCHMARK SUITE

To evaluate our Physics-Informed JEPA framework across diverse edge, cyber-physical, and orbital sensing dynamics, our benchmark suite comprises 11 datasets spanning 45 strictly matched continuous telemetry channels sourced from NASA spacecraft archives [3], the PhysioNet biomedical database [29], TimeEval [30], and mTSBench [31]. We group these channels into four IoT and remote sensing domains:

A. Domain 1: Satellite and Spacecraft Telemetry (18 Channels)

- 1) *NASA SMAP Earth Observation Satellite [3] (8 evaluated channels)*: Continuous power, thermal, and attitude telemetry from the Soil Moisture Active Passive (SMAP) spacecraft tracking solar array currents, bus voltages, battery temperatures, and thruster status subject to radiation upsets and operational mode shifts.
- 2) *NASA MSL Curiosity Rover [3] (9 evaluated channels)*: Environmental and robotic actuator telemetry from the Mars Science Laboratory (MSL) rover monitoring drive motor currents, joint angles, arm kinematics, and power subsystem faults during planetary exploration.
- 3) *SWAN-SF Space Weather Telemetry [32] (1 channel)*: Solar photospheric vector magnetogram telemetry from NASA SDO/HMI tracking active region magnetic field perturbations during solar eruptive events.

B. Domain 2: Wearable Biometrics and Kinematics (9 Channels)

- 1) *Daphnet Wearable Body Sensor Network [29] (9 channels)*: High-frequency tri-axial accelerometry from wearable body sensor nodes (placed on the shank, thigh, and trunk; 9 continuous dimensions per subject) tracking Freezing-of-Gait (FoG) breakdown events in Parkinson's patients.

C. Domain 3: Industrial IoT and Cyber-Physical Systems (8 Channels)

- 1) *GHL Industrial Separation Plant (6 channels)*: Multi-phase gas-oil-water separation industrial process telemetry monitoring tank levels, flow rates, and valve control faults.
- 2) *Genesis Automation Testbed (1 channel)*: Industrial manufacturing testbed telemetry tracking automated handling machinery faults.
- 3) *GECCO Water Distribution IoT (1 channel)*: Physical water distribution process telemetry monitoring chemical and hydrological sensor measurements.

D. Domain 4: Smart Infrastructure and Environmental Sensing (10 Channels)

- 1) *Room-Occupancy Environmental IoT [33] (2 channels)*: Multi-sensor ambient environmental IoT sensing (CO₂, light, temperature, humidity) tracking room-occupancy state switches.
- 2) *CICIDS Edge Network Telemetry (6 channels)*: High-throughput network flow telemetry capturing edge gateway protocol dynamics and security anomalies (DDoS bursts and port scans).
- 3) *Callt2 Building People-Flow Telemetry [34] (1 channel)*: Continuous pedestrian counts via optical beam sensors at the UC Irvine Calit2 facility entrance.
- 4) *Metro Highway Traffic (1 channel)*: Interstate highway traffic volume and environmental condition telemetry.

Each benchmark channel provides pre-split training sequences (consisting exclusively of nominal operational behavior) and testing sequences containing labeled ground-truth

anomalies. Continuous numeric features are standardized via training statistics ($\mu_{\text{train}}, \sigma_{\text{train}}$), flooring standard deviations at $\max(\sigma, 10^{-2})$ to prevent division by zero during flat-line operational phases. While full mission repositories contain larger sensor arrays (e.g., SMAP records 51 telemetry streams and MSL records 26 streams across diverse subsystems), our benchmark evaluates the standardized 8- and 9-channel subsets commonly evaluated in the recent benchmark literature [3], [30] to ensure strict matching and reproducible runtime conditions across all baselines.

IV. METHODOLOGY: THE PHYSICS-INFORMED GEOMETRIC JEPA TRIAD

This section formalizes the mathematical formulation and inductive biases of the proposed framework.

A. Problem Formulation and Latent Prediction Paradigm

Let $\mathbf{X} \in \mathbb{R}^{T \times K}$ denote a continuous multivariate time-series sequence consisting of K observed sensor variables over T timesteps. The sequence is partitioned into historical context $x_{\text{ctx}} = \mathbf{X}[t - C : t, :] \in \mathbb{R}^{C \times K}$ and suspect/target horizon $x_{\text{tgt}} = \mathbf{X}[t : t + S, :] \in \mathbb{R}^{S \times K}$, where C and S denote the respective temporal horizons.

Existing deep learning baselines (e.g., TimesNet, TranAD, Anomaly Transformer) frame anomaly detection as pointwise reconstruction:

$$\min_{\theta} \mathbb{E} [\|\mathbf{X} - \mathcal{D}_{\theta}(\mathcal{E}_{\theta}(\mathbf{X}))\|_F^2]. \quad (1)$$

In complex multi-sensor telemetry, ambient high-frequency stochastic jitter dominates the Frobenius norm, causing models to waste representational bandwidth reconstructing sensor noise rather than modeling system state dynamics (leading to noise drowning).

To overcome this, we adopt the Joint-Embedding Predictive Architecture (JEPA) paradigm [8]. A context encoder E_{θ} maps $x_{\text{ctx}} \rightarrow z_{\text{ctx}} \in \mathbb{R}^D$, while an Exponential Moving Average (EMA) momentum target encoder E_{ϕ} maps $x_{\text{tgt}} \rightarrow z_{\text{tgt}} \in \mathbb{R}^D$, updated via:

$$\phi_k \leftarrow m_k \phi_{k-1} + (1 - m_k) \theta_k, \quad m_k \in [0.996, 0.9995]. \quad (2)$$

The core objective is to predict target representations in latent space under physical dynamical constraints.

B. The Physics-Informed JEPA Triad

Rather than relying on unconstrained Euclidean regression, we incorporate physical domain regularizers through three complementary geometric formulations:

1) *1. Operator-Entropy JEPA (Koopman Spectral Dynamics)*: In dynamical systems theory, non-linear trajectories can be analyzed via linear operators acting on observable functions. In our architecture, the encoder maps the observation window to a finite-dimensional latent observable vector $z \in \mathbb{R}^D$. We define the transition between context and target latent states via a learned linear transition matrix $\mathbf{K} \in \mathbb{R}^{D \times D}$:

$$\hat{z}_{\text{tgt}} = \mathbf{K} z_{\text{ctx}}. \quad (3)$$

To regularize modal stability across latent coordinates, we evaluate the singular value spectrum $\sigma_i(\mathbf{K})$, which governs the energy amplification and contraction of the linear map. We define the normalized spectral entropy as:

$$\mathcal{H}_{\text{spectral}}(\mathbf{K}) = - \sum_{i=1}^D \tilde{\sigma}_i \log \tilde{\sigma}_i, \quad \tilde{\sigma}_i = \frac{\sigma_i}{\sum_{j=1}^D \sigma_j + \epsilon}. \quad (4)$$

A unitary, energy-conserving linear map distributes energy evenly across all orthogonal modes ($\sigma_i \approx 1$), producing maximal spectral entropy $\mathcal{H}_{\text{max}} = \log D$. Conversely, severe damping or modal collapse contracts dynamics onto a low-rank subspace (low spectral entropy). To prevent representation collapse toward low-rank degeneration and encourage non-degenerate linear transitions during training, the Operator-Entropy training objective regularizes the predictive transition while maximizing spectral dispersion:

$$\mathcal{L}_{\text{Op-Entropy}} = \|\mathbf{K} z_{\text{ctx}} - z_{\text{tgt}}\|_2^2 - \lambda_{\text{ent}} \mathcal{H}_{\text{spectral}}(\mathbf{K}), \quad (5)$$

where $\lambda_{\text{ent}} = 10^{-3}$ penalizes dimensional collapse across latent modes.

2) *2. Reynolds-Stress JEPA (Continuum Mechanics & Covariance Alignment)*: In multi-channel telemetry representing coupled flows (thermal, convective, hydraulic), channel interactions are governed by transport phenomena. Inspired by hydrodynamic *Reynolds decomposition*, we decompose the temporal sequence of target representations $\mathbf{Z}_{\text{tgt}} = [z_{t+1}, \dots, z_{t+S}]^T \in \mathbb{R}^{S \times D}$ across horizon S into a rolling baseline trajectory and dynamic fluctuations:

$$\mathbf{Z}_{\text{tgt}} = \bar{\mathbf{z}} \mathbf{1}^T + \mathbf{Z}', \quad (6)$$

where $\bar{\mathbf{z}} = \frac{1}{S} \sum_{s=1}^S z_s \in \mathbb{R}^D$ denotes the window temporal mean, and $\mathbf{Z}' \in \mathbb{R}^{S \times D}$ denotes the zero-mean temporal fluctuations.

The inter-coordinate momentum transport across the temporal window is characterized by the second-moment fluctuation covariance matrix:

$$\mathbf{R}_{\text{tgt}} = \frac{1}{S} \mathbf{Z}'^T \mathbf{Z}' \in \mathbb{R}^{D \times D}. \quad (7)$$

Anomalies disrupt the nominal multi-sensor fluctuation covariance balance. The Reynolds-Stress JEPA loss aligns the predicted mean trajectory and target fluctuation covariance:

$$\mathcal{L}_{\text{Reynolds}} = \|\bar{z}_{\text{tgt}} - \hat{z}_{\text{tgt}}\|_2^2 + \lambda_R \|\mathbf{R}_{\text{pred}} - \mathbf{R}_{\text{tgt}}\|_F^2, \quad (8)$$

with $\lambda_R = 10^{-2}$.

3) *3. Potential-Flow JEPA (Conservative Field Theory)*: Under the Helmholtz-Hodge decomposition theorem, any sufficiently smooth vector field can be uniquely resolved into an irrotational (curl-free) conservative potential gradient and a solenoidal (divergence-free) rotational field.

In nominal physical equilibria, conservative dynamical systems conserve potential energy along phase trajectories. We parameterize the latent transition dynamics through the gradient of a scalar potential network $\Phi_{\psi}(z)$:

$$\hat{v}(z) = \nabla_z \Phi_{\psi}(z), \quad \hat{z}_{\text{tgt}} = z_{\text{ctx}} + \hat{v}(z_{\text{ctx}}). \quad (9)$$

Because the curl of the gradient of any C^2 scalar field is identically zero ($\nabla \times (\nabla \Phi) \equiv \mathbf{0}$), parameterizing the velocity




Fig. 1. Architectural overview of the proposed Physics-Informed Joint-Embedding Predictive Architecture (JEPA) framework. The historical context window $\mathbf{x}_{\text{ctx}}$ and suspect target window $\mathbf{x}_{\text{tgt}}$ are projected through multi-scale dyadic TCN encoders into latent representations $\mathbf{z}_{\text{ctx}}$ and $\mathbf{z}_{\text{tgt}}$. Rather than reconstructing raw observation-space waveforms (which causes false alarm flooding under ambient sensor jitter), the predictive core operates strictly in latent representation space, governed by three complementary physical inductive biases: Koopman operator spectral entropy, hydrodynamic Reynolds stress transport, and conservative potential flow field theory. Prediction discrepancy is calibrated online via streaming Extreme Value Theory (SPOT).

field as $\hat{v} = \nabla\Phi$ guarantees that the predicted trajectory field is *strictly irrotational by construction*. Physical anomalies that introduce non-conservative work, dissipative friction, or unmodeled circulation diverge from any integrable potential field, producing large residual discrepancies:

$$\mathcal{L}_{\text{Pot-Flow}} = \|\nabla_z \Phi_\psi(\mathbf{z}_{\text{ctx}}) - (\mathbf{z}_{\text{tgt}} - \mathbf{z}_{\text{ctx}})\|_2^2. \quad (10)$$

C. Anomaly Scoring and Operational Detection Latency

During inference, sliding context windows $\mathbf{x}_{\text{ctx}} = \mathbf{X}[t-C:t, :]$ and suspect target windows $\mathbf{x}_{\text{tgt}} = \mathbf{X}[t:t+S, :]$ are processed causally with a stride of 1 timestep. For all JEPA architectures, the anomaly score is computed as the latent predictive discrepancy between the predicted target representation $\hat{\mathbf{z}}_{\text{tgt}}$ and the target encoder representation $\mathbf{z}_{\text{tgt}}$:

$$S(t) = \|\hat{\mathbf{z}}_{\text{tgt}}(t) - \mathbf{z}_{\text{tgt}}(t)\|_2^2. \quad (11)$$

Because the target horizon spans $\mathbf{x}_{\text{tgt}} = \mathbf{X}[t:t+S, :]$, evaluating $S(t)$ operates with a fixed lookahead latency buffer of $S = 64$ timesteps. The physical inductive regularizers enter as structural inductive constraints during training and architectural parameterization, shaping the geometry of the learned latent manifold without requiring ad-hoc test-time heuristic score terms. Each timestep t is mapped directly to its calibrated anomaly score $S(t)$.

D. Shared Internal Inductive Biases

To stabilize learning across high-dimensional, noisy telemetry, all three architectures incorporate three inductive mechanisms:

1) *Coordinate Saliency Gating*: On high-dimensional streams ($K \geq 50$), anomalous signals frequently manifest in only a small subset of channels while remaining channels exhibit ambient noise. To prevent dimensional drowning, we introduce a parameter-free *Coordinate Saliency Gate* within the encoder:

$$s_k = \frac{\text{Var}_t(x_k)}{\sum_{j=1}^K \text{Var}_t(x_j) + \epsilon}, \quad \tilde{x}_k = x_k \cdot \sigma\left(\frac{s_k - \bar{s}}{\sigma_s}\right), \quad (12)$$

where $\sigma(\cdot)$ is the logistic sigmoid, and $\bar{s}, \sigma_s$ denote the mean and standard deviation of channel variances. This dynamically prioritizes informative non-stationary coordinates without discarding quiet sensor channels.

2) *Littlewood-Paley Dyadic Wavelet Shells*: Real-world physical anomalies operate across disparate temporal scales, ranging from high-frequency electrical spikes to gradual thermal drifts. We integrate Littlewood-Paley dyadic frequency decomposition into the encoder front-end. The input sequence is decomposed into dyadic octave bands via bandpass filters $\psi_j(\omega) = \psi(2^j\omega)$ satisfying the partition of unity:

$$|\phi(\omega)|^2 + \sum_{j=0}^{J-1} |\psi(2^j\omega)|^2 = 1, \quad \forall \omega \in [-\pi, \pi]. \quad (13)$$

This multi-scale frequency decomposition enables the model to isolate high-frequency transient anomalies from low-frequency baseline diurnal oscillations without expanding parameter count.

3) 3. Kurtosis-Adaptive Extreme Value Theory (EVT):

Following scoring, decision alarms $\hat{y}_t = \mathbb{I}(\mathcal{S}(t) > \tau_t)$ are generated via Streaming Peaks-Over-Threshold (SPOT). In standard SPOT, the extreme threshold τ is determined by fitting a Generalized Pareto Distribution (GPD) with scale σ and shape ξ to excesses over an initial threshold u at risk probability q :

$$\tau = u + \frac{\sigma}{\xi} \left[\left(\frac{qN_u}{N} \right)^{-\xi} - 1 \right]. \quad (14)$$

To account for varying tail heaviness across diverse sensor regimes, we modulate the threshold via the empirical excess kurtosis γ_2 of calibration residuals:

$$\tau_{\text{eff}} = \tau \cdot \left[1 + \frac{1}{6} \tanh(\gamma_2) \right]. \quad (15)$$

On heavy-tailed distributions ($\gamma_2 > 0$), τ_{eff} increases, raising the threshold to suppress false alarms from volatile noise bursts; on light-tailed distributions ($\gamma_2 < 0$), the threshold tightens to maximize detection sensitivity.

V. EXPERIMENTAL PROTOCOL AND EVALUATION METRICS

This section outlines the experimental setup, baseline architectures, and the multi-tier statistical evaluation protocol designed to prevent metric inflation.

A. Experimental Setup and Hyperparameters

To ensure rigorous, fair comparison across all models, all architectures are trained and evaluated under an identical streaming time-series pipeline:

- **Sliding Window Geometry:** Past context duration $C = 256$, suspect/target horizon $S = 64$ (total window length $L = 320$), evaluated with a stride of 1 timestep across testing streams.
- **Encoder Backbone:** Context representations $z_{\text{ctx}} \in \mathbb{R}^D$ ($D = 32$) are extracted using a causal dilated Temporal Convolutional Network (TCN) with 3 residual blocks, channel depth $F = 48$, kernel size $K = 5$, and exponentially increasing dilation factors $d \in \{1, 2, 4\}$ regularized by Layer Normalization and dropout ($p = 0.20$).
- **Target Encoder Update:** The target encoder parameters θ_{tgt} are updated via Exponential Moving Average (EMA) momentum tracking: $\theta_{\text{tgt}} \leftarrow m\theta_{\text{tgt}} + (1 - m)\theta_{\text{ctx}}$, with m following a cosine schedule from $m_{\text{base}} = 0.996$ to $m_{\text{max}} = 0.9995$.
- **Optimization Protocol:** AdamW optimizer with initial learning rate $\eta = 10^{-3}$ decayed to $\eta_{\text{min}} = 10^{-5}$ via Cosine Annealing, weight decay 10^{-4} , and mini-batch size $B = 32$ across 20 training epochs on nominal telemetry sequences.
- **Statistical Calibration:** Extreme Value Theory (EVT / SPOT) calibration fits Generalized Pareto Distribution (GPD) tails to the top 2.0% ($u = 0.98$) of nominal

training residuals under extreme risk parameter $q = 10^{-3}$, dynamically modulated by the empirical sample kurtosis of the channel anomaly stream.

B. Baseline Architectures and Unified Reimplementation

To isolate algorithmic principles from external codebase discrepancies, all baselines are executed within a unified PyTorch benchmarking pipeline under identical windowing, standardization, and EVT calibration:

- 1) **TimesNet [6]:** Multi-periodic reconstructive architecture that transforms 1D time series into 2D variation spaces via 1D-FFT periodic frequency discovery and 2D Inception blocks (574,529 trainable parameters).
- 2) **TranAD [15]:** Adversarial reconstructive Transformer featuring two-stage self-conditioning and sub-series attention to amplify subtle reconstruction errors (201,090 trainable parameters).
- 3) **TS-JEPA (Unconstrained Control):** Standard Joint-Embedding Predictive Architecture predicting target representations z_{tgt} from context z_{ctx} without physical inductive priors (217,925 trainable parameters; 431,466 total including the EMA momentum target encoder).
- 4) **Physics-Informed Formulations:** Three physics-informed architectures incorporating spectral Koopman dissipation (Operator-Entropy JEPA, 217,957 trainable parameters), turbulent momentum transport (Reynolds-Stress JEPA, 224,817 trainable parameters), and conservative curl-free potential constraints (Potential-Flow JEPA, 266,981 trainable parameters). The physical heads introduce negligible parameter overhead (under 1% to 22% over the base encoder), demonstrating that performance differences arise from inductive geometry rather than parameter capacity confounds.

C. Multi-Tier Statistical Evaluation Metrics

Recent critical surveys in time-series anomaly detection [30], [35], [36] demonstrate that standard evaluation protocols are plagued by severe reporting vulnerabilities, including point-adjustment inflation, channel-averaging imbalances, and arbitrary threshold gaming. To ensure transparency and reproducibility, we report evaluation metrics across multiple complementary tiers:

- 1) **Dataset-Macro Point-F1 (Primary Metric):** We compute the unweighted arithmetic mean of Point-F1 scores across datasets:

$$\text{Macro-F1}_{\text{DS}} = \frac{1}{|\mathcal{D}|} \sum_{d \in \mathcal{D}} \left(\frac{1}{|C_d|} \sum_{c \in C_d} \text{F1}_c \right) \quad (16)$$

where $\mathcal{D}$ is the set of benchmark datasets and C_d is the set of channels in dataset d . This prevents datasets with numerous channels from dominating the aggregate benchmark score.

- 2) **Pooled Micro Point-F1:** We pool all raw True Positives, False Positives, and False Negatives globally across all test timesteps:

$$\text{Micro-F1} = \frac{2 \sum_c \text{TP}_c}{2 \sum_c \text{TP}_c + \sum_c \text{FP}_c + \sum_c \text{FN}_c} \quad (17)$$

Micro-F1 directly penalizes models that flood long streams with hundreds of thousands of false alarms, exposing reconstructive noise drowning.

- 3) **Channel-Macro Point-F1:** The standard per-channel average $\frac{1}{N} \sum_{i=1}^N \text{F1}_i$ across all evaluated channels.
- 4) **Precision-Recall Area Under the Curve (PR-AUC):** Average Precision computed over all possible continuous anomaly score thresholds:

$$\text{PR-AUC} = \sum_k (R_k - R_{k-1}) P_k \quad (18)$$

PR-AUC evaluates raw score ranking capability entirely independent of the threshold calibration mechanism.

- 5) **Trivial All-Positive Reference Baseline:** We report the analytical performance of a constant all-positive predictor ($\hat{y}_t = 1, \forall t$):

$$\text{F1}_{\text{all-pos}} = \frac{2 \cdot P}{2 \cdot P + N} = \frac{2\pi}{1 + \pi} \quad (19)$$

where $\pi = P/(P+N)$ is the anomaly base rate. Any model failing to exceed $\text{F1}_{\text{all-pos}}$ provides less discriminative value than an uncalibrated constant alarm.

- 6) **Contiguous Point-Adjusted F1 (PA-F1):** Standard segment adjustment where detecting any single anomaly point expands to the entire contiguous anomaly segment. To prevent boundary bleeding, point adjustment is computed strictly within contiguous ground-truth segments before metric calculation. We also report the *Test-Sweep PA-F1* (best post-hoc threshold scan) to quantify the remaining calibration gap.

VI. ABLATION STUDIES AND ANALYSIS

To isolate the individual contributions of each mathematical and architectural component in our Physics-Informed JEPA framework, we conduct systematic ablation experiments across representative multi-domain telemetry types:

- 1) *High-Dimensional Aerospace Telemetry (NASA SMAP / MSL, $K \geq 50$):* High-dimensional sensor pods with strong inter-channel thermal and electrical cross-coupling.
- 2) *Periodic Biomechanical Kinematics (Daphnet FoG):* Fast multi-axial accelerometry exhibiting repeating limit-cycle dynamics and abrupt gait freezing.
- 3) *Transient Cyber-Physical Hydraulics (GECCO / SWAN):* Continuous environmental and hydrodynamic streams with sharp, turbulent anomalies.

A. Ablation 1: Physics-Informed Inductive Biases vs. Unconstrained Prediction

Table II and Table I(a) isolate the impact of embedding physical domain constraints into latent representations compared to unconstrained latent forecasting (TS-JEPA) and observation-space reconstruction (TimesNet, TranAD):

- *Reconstructive Observation Space:* Reconstructive models suffer from severe false alarm flooding (TimesNet accumulates over 10^6 false positives, yielding a Pooled Micro-F1 of only 0.0647).
- *Unconstrained Latent Prediction (TS-JEPA):* Moving prediction to latent space eliminates noise memorization, lifting

Micro-F1 to 0.2860 and achieving 0.1838 Dataset-Macro F1 with 0.3449 PA-F1.

- *Physics-Informed Formulations:* Operator-Entropy JEPA attains the highest Dataset-Macro F1 (0.2066) and Channel-Macro F1 (0.1743), while Potential-Flow JEPA achieves the highest overall Micro-F1 (0.3458) and reduces false alarms by 8-fold (130,237 vs 1,047,144).

B. Ablation 2: Coordinate Saliency Gating in High-Dimensional Pods

In high-dimensional telemetry ($K \geq 50$), anomalous faults typically originate in a localized subsystem (e.g., a single valve or battery cell). In standard latent pooling:

$$z_{\text{pool}} = \frac{1}{K} \sum_{k=1}^K z_k \quad (20)$$

the localized anomalous signal is diluted by ambient normal noise across the remaining 49 quiescent channels.

As shown in Table I(b), disabling Coordinate Saliency Gating ($\tilde{x}_k = x_k$) causes a substantial drop in Point-F1 on SMAP ($0.285 \rightarrow 0.198$) and MSL ($0.265 \rightarrow 0.174$). The parameter-free saliency gate $\tilde{x}_k = x_k \cdot \sigma\left(\frac{s_k - \bar{s}}{\sigma_s}\right)$ dynamically isolates high-variance, non-stationary coordinates, suppressing ambient quiescent dimensions and maintaining high signal-to-noise ratio.

C. Ablation 3: Littlewood-Paley Dyadic Wavelet Shells

Real-world industrial and environmental telemetry contains compound temporal dynamics operating simultaneously across disjoint frequency bands. For instance, in municipal water infrastructure (GECCO), sensors record 24-hour diurnal consumption cycles alongside sub-second hydraulic valve transients.

Standard monolithic convolutions either over-smooth high-frequency transients or fail to capture long-period cycles. As reported in Table I(c), replacing standard convolutions with Littlewood-Paley dyadic wavelet projections:

$$z^{(j)} = \Psi_j * x, \quad j \in \{0, 1, \dots, J-1\} \quad (21)$$

improves Point-F1 on GECCO from 0.1429 to 0.2904 on Operator-Entropy JEPA, while lifting Point-Precision on Daphnet accelerometry from 0.245 to 0.342. Decoupling frequency octaves prevents low-frequency power from masking subtle high-frequency phase disruptions.

D. Ablation 4: Kurtosis-Adaptive vs. Static Extreme Value Calibration

Standard Extreme Value Theory (SPOT) fits a Generalized Pareto Distribution to residuals exceeding an initial quantile u . However, telemetry noise profiles exhibit substantial variation in tail heaviness across sensor channels. In leptokurtic distributions ($\gamma_2 > 0$, heavy tails), static thresholds trigger frequent false alarms; in platykurtic distributions ($\gamma_2 < 0$), static thresholds miss low-amplitude structural faults.

Table I(d) compares our kurtosis-adaptive EVT calibrator against static calibration baselines:

- 1) *Static 3-Sigma Gaussian Threshold*: Yields poor detection (0.082 Macro Pt-F1) due to pervasive non-Gaussianity and heavy tails in raw telemetry residuals.
- 2) *Standard EVT (Unweighted SPOT)*: Achieves 0.174 Macro Pt-F1, but suffers from false positive bursts during non-stationary operational phases.
- 3) *Kurtosis-Adaptive EVT (Ours)*: By scaling the effective tail threshold $\tau_{\text{eff}} = \tau \cdot (1 + \frac{1}{6} \tanh(\gamma_2))$, the calibrator raises decision margins during heavy-tailed regimes while sharpening sensitivity on calm channels, achieving 0.2066 Macro Pt-F1.

VII. RESULTS AND DISCUSSION

This section reports quantitative benchmarking and dynamical analyses across 45 matched continuous telemetry channels evaluated across 11 benchmark datasets. We evaluate the proposed Physics-Informed JEPA framework against reconstructive and contrastive baselines.

A. Benchmark Performance on Core Multi-Sensor Telemetry

Table II presents the head-to-head empirical comparison across the 11 core benchmark datasets (45 matched channels) evaluated under an identical protocol. Figure 2 illustrates the comparative bar chart across strict Point-F1 and Point-Adjusted F1 (PA-F1).

Key findings from Table II include:

- 1) **Macro and Micro Point-F1 Performance**: In unweighted Dataset-Macro Point-F1, Operator-Entropy JEPA secures Rank 1 (0.2066), representing a **+43.2% relative improvement** over TimesNet (0.1443) and **+28.4%** over TranAD (0.1609). In Channel-Macro Point-F1, Operator-Entropy JEPA (0.1743) and Reynolds-Stress JEPA (0.1532) capture the top two ranks. In Pooled Micro-F1, Potential-Flow JEPA achieves the highest score of **0.3458**, more than **3.2× higher than TranAD** (0.1072) and **5.3× higher than TimesNet** (0.0647).
- 2) **Precision and False Alarm Suppression**: An operational limitation of conventional baselines is false alarm flooding in continuous telemetry. Reconstructive baselines suffer precision degradation under sensor jitter: TimesNet produces 1,047,144 false positives across the evaluated channels (precision 4.42%), while TranAD accumulates 452,621 false positives (precision 9.64%). In contrast, Potential-Flow JEPA achieves a precision of **45.98%** with 130,237 false positives, representing an 8-fold reduction in false alarms compared to TimesNet, while detecting 110,866 true positives.
- 3) **Continuous Ranking (PR-AUC) vs. Calibrated Decisions (Point-F1)**: Comparing threshold-free continuous ranking against calibrated decisions reveals an important empirical trade-off. In PR-AUC (Average Precision), TranAD (0.2612) and TimesNet (0.2585) achieve top aggregate ranking scores, outperforming JEPA models on 5 of the 11 datasets (Daphnet, Genesis, CICIDS, Metro, and Room Occupancy; see Figure 4). This demonstrates

that reconstructive representations retain strong continuous anomaly ranking fidelity. However, when automated streaming extreme-value thresholding (SPOT) is applied, observation-space reconstruction errors trigger widespread false alarm flooding (over 1.04×10^6 false positives for TimesNet, flagging nearly 40% of nominal points). This occurs because raw observation residuals exhibit high temporal variance and localized baseline drift, causing test-time reconstruction residuals to systematically exceed the scale of the nominal training GPD fit. In contrast, predicting in a normalized latent representation space ($\mathbb{R}^{32}$) bounds residual variance, yielding well-conditioned error distributions that stabilize streaming threshold calibration.

- 4) **Role of Unconstrained Latent Prediction (TS-JEPA Control)**: Unconstrained TS-JEPA achieves 0.1838 Dataset-Macro F1, 0.2860 Micro-F1, and the highest Dataset-Macro PA-F1 (0.3449). Across all metric comparisons in Table II, physics-informed variants outperform the unconstrained TS-JEPA control in 9 of 18 instances. This demonstrates that predicting in latent representation space accounts for the foundational improvement in threshold calibration stability. The physical inductive priors do not provide uniform dominance across all metrics, but instead act as targeted regularizers offering domain-aligned advantages in specialized regimes (e.g., Operator-Entropy in cyclic and spectral telemetry, Potential-Flow in conservative mechanical subsystems).
- 5) **Comparison to Analytical All-Positive Reference**: In Table II, the analytical All-Positive baseline achieves a Dataset-Macro F1 of 0.1893 and a Channel-Macro F1 of 0.1942. On Dataset-Macro, Operator-Entropy JEPA (0.2066) is the sole model to surpass this trivial reference limit, while on Channel-Macro, all deep models fall below the constant alarm baseline due to severe channel-level class imbalance.

B. Domain-Specialized Strengths of the Flagship Triad

Figure 3 illustrates the radar profile of the flagship models across five distinct telemetry domains, and Figure 4 details the per-dataset PR-AUC matrix across all tested benchmarks.

Each model in the triad exhibits specialized mathematical inductive advantages:

1) *Operator-Entropy JEPA (Koopman Spectral Mechanics)*: By monitoring dynamic spectral entropy within a latent Koopman transition representation, `operator_entropy_jepa` is sensitive to disruptions in non-linear cyclic and mixing dynamics. It achieves high performance across NASA spacecraft telemetry (MSL: 0.478 PR-AUC; SMAP: 0.467 PR-AUC), industrial process loops (GECCO: 0.351 PR-AUC vs Potential-Flow's 0.015), and space weather solar flare dynamics (`swan`: 0.436 PR-AUC).

2) *Reynolds-Stress JEPA (Continuum Mechanics & Turbulent Transport)*: Grounded in hydrodynamic Reynolds decomposition ($z = \bar{z} + z'$), `reynolds_stress_jepa` aligns cross-channel fluctuation covariance between context and target horizons. This formulation is designed for systems characterized by coupled flows. On the SMAP satellite telemetry

TABLE I

SYSTEMATIC ABLATION STUDIES OF CORE INDUCTIVE COMPONENTS ACROSS MULTI-DOMAIN TELEMETRY BENCHMARKS. PANEL (A) EVALUATES THE IMPACT OF PHYSICAL INDUCTIVE PRIORS; PANEL (B) VALIDATES COORDINATE SALIENCY GATING IN HIGH-DIMENSIONAL SENSOR ARRAYS ($K \geq 50$); PANEL (C) TESTS LITTLEWOOD-PALEY DYADIC WAVELET SHELLS ON MULTI-SCALE DYNAMICS; PANEL (D) EVALUATES KURTOSIS-ADAPTIVE EXTREME VALUE CALIBRATION AGAINST STATIC THRESHOLDING BASELINES.

(a) Physics-Informed Priors vs. Unconstrained Latent Prediction					(b) Coordinate Saliency Gating in High Dimensions			
Architecture	DS-Macro F1	Micro F1	PR-AUC	False Positives	Channel / Domain	Uniform Pooling	Saliency Gated	Relative Gain
TimesNet (2023)	0.1443	0.0647	0.2585	1,047,144	NASA SMAP ($K = 51$)	0.1982	0.2854	+44.0%
TranAD (2022)	0.1609	0.1072	0.2612	452,621	NASA MSL ($K = 26$)	0.1741	0.2651	+52.3%
TS-JEPA (Unconstrained)	0.1838	0.2860	0.2344	176,231	Daphnet Kinematics ($K = 9$)	0.2114	0.2340	+10.7%
Operator-Entropy JEPA	0.2066	0.2806	0.2541	161,469	GHL Separation ($K = 14$)	0.0782	0.1042	+33.2%
Reynolds-Stress JEPA	0.1544	0.1175	0.2406	188,198	High-Dim Macro Avg	0.1655	0.2222	+34.3%
Potential-Flow JEPA	0.1147	0.3458	0.1978	130,237	(d) Kurtosis-Adaptive Extreme Value Calibration			
(c) Littlewood-Paley Dyadic Wavelet Shells					Threshold Calibration	Macro Pt-F1	Macro PA-F1	Precision
Temporal Encoder	GECCO F1	SWAN F1	Daphnet F1		Static $3\text{-}\sigma$ Gaussian	0.0821	0.1432	0.0910
Standard Monolithic Conv	0.1429	0.2210	0.2015		Standard EVT (Unweighted)	0.1740	0.2845	0.2120
Fixed FFT Filterbank	0.1840	0.2845	0.2180		Transductive Test Sweep	0.3541	0.4820	0.4105
Dyadic Wavelet Shells (Ours)	0.2904	0.5048	0.2340		Kurtosis-Adaptive EVT (Ours)	0.2066	0.3320	0.3621

TABLE II

COMPREHENSIVE MULTI-TIER BENCHMARK ON CORE MULTI-SENSOR DATASETS (45 STRICTLY MATCHED CHANNELS ACROSS 11 DIVERSE MODALITIES). WE REPORT UNADJUSTED POINT-F1 ACROSS THREE AGGREGATION REGIMES: UNWEIGHTED DATASET-MACRO (EQUAL $1/N$ DATASET WEIGHT), POOLED MICRO-F1 ($\sum TP, \sum FP, \sum FN$), AND CHANNEL-MACRO F1, ALONGSIDE THRESHOLD-INDEPENDENT PR-AUC AND POINT-ADJUSTED F1 (PA-F1). THE BEST RESULT IN EACH COLUMN IS IN **BOLD**, AND THE SECOND BEST IS UNDERLINED. ALL MODELS EVALUATED UNDER IDENTICAL PURE INDUCTIVE EVT CALIBRATION WITHOUT TRANSDUCTIVE TEST-TAIL CLAMPS.

Architecture	Physical Inductive Bias / Paradigm	Strict Point-Level Evaluation (Unadjusted)			Ranking	Point-Adjusted (Literature Ref.)	
		DS-Macro $\uparrow$	Micro-F1 $\uparrow$	Chan-Macro $\uparrow$	PR-AUC $\uparrow$	DS-Macro $\uparrow$	Chan-Macro $\uparrow$
<i>Proposed Physics-Informed Geometric JEPA Triad</i>							
Operator-Entropy JEPA	Koopman Spectral Entropy ($\mathcal{H}_K$)	0.2066	<u>0.2806</u>	0.1743	0.2541	0.3320	0.3205
Reynolds-Stress JEPA	Hydrodynamic Stress Tensor ($\mathbf{R}$)	0.1544	0.1175	<u>0.1532</u>	0.2406	0.3153	<u>0.3226</u>
Potential-Flow JEPA	Curl-Free Potential Field ($\nabla\Phi$)	0.1147	0.3458	0.1259	0.1978	0.3053	0.3436
<i>Ablation Control (Unconstrained Latent Prediction)</i>							
Vanilla TS-JEPA	Latent Prediction (Unconstrained)	<u>0.1838</u>	0.2860	0.1442	0.2344	0.3449	0.3066
<i>Published State-of-the-Art Deep Baselines (Unified Reimplementations)</i>							
TranAD (VLDB 2022)	Reconstructive Transformer + Adversarial	0.1609	0.1072	0.1319	0.2612	0.2689	0.2537
TimesNet (ICLR 2023)	Multi-Period 2D Temporal Inception	0.1443	0.0647	0.1331	<u>0.2585</u>	0.2373	0.2072
<i>Theoretical Reference Limit</i>							
All-Positive Trivial Predictor	Constant Positive ($\hat{y} \equiv 1$)	0.1893	0.2372	0.1942	—	—	—

dataset, it achieves Rank 1 in PR-AUC (**0.512** PR-AUC), outperforming TimesNet (0.479) and TranAD (0.345). On SWAN solar telemetry, it achieves **0.522** PR-AUC and strong precision.

3) *Potential-Flow JEPA (Conservative Field Theory)*: By constraining nominal dynamics to curl-free potential gradients ($\hat{v} = \nabla\Phi$), `potential_flow_jepa` enforces conservative trajectory tracking. It achieves the **highest aggregate Micro Point-F1 (0.3458)** and **highest Channel PA-F1 (0.3436)** across evaluated models, with a precision of **45.98%**. On NASA spacecraft subsystems (MSL), it scores the top PR-AUC of **0.551** (vs TimesNet’s 0.380 and TranAD’s 0.332).

4) *Role of Vanilla TS-JEPA (Ablation Control)*: To isolate the contribution of the physical inductive biases from the predictive latent paradigm itself, we evaluate `ts_jepa` as an unconstrained control. `ts_jepa` achieves 0.1838 Dataset-Macro Point-F1, 0.1442 Channel-Macro Point-F1, and 0.3066 Channel PA-F1. While substantially outperforming reconstructive baselines (+27.4% in Dataset-Macro F1 over TimesNet), it trails `operator_entropy_jepa` in Dataset-Macro Point-F1 (0.1838 vs 0.2066) and Channel Point-F1 (0.1442 vs 0.1743), and trails `potential_flow_jepa` in Micro-F1

(0.2860 vs 0.3458) and Channel PA-F1 (0.3066 vs 0.3436). This indicates that while latent representation prediction mitigates noise drowning, domain-aligned physical conservation laws provide beneficial inductive structure.

C. Impact of Architectural Enhancements

Ablation of the three internal inductive layers confirms their collective role in stabilizing high-dimensional inference:

- 1) *Coordinate Saliency Gating*: On high-dimensional multi-sensor streams ($K \geq 50$ in NASA SMAP and MSL telemetry), unweighted latent pooling causes informative anomalous dimensions to drown in ambient channel noise. Saliency gating dynamically prioritizes high-variation informative channels, enabling `reynolds_stress_jepa` and `potential_flow_jepa` to detect localized subsystem faults across complex sensor arrays without manual channel sub-selection.
- 2) *Littlewood-Paley Dyadic Wavelet Shells*: Multi-scale temporal decomposition within the encoder allows the networks to isolate fast transient spikes from low-frequency baseline drifts. On GECCO (water infrastruc-




Fig. 2. Benchmark evaluation of the proposed Physics-Informed JEPA Triad (operator_entropy_jepa, reynolds_stress_jepa, potential_flow_jepa) and unconstrained control (ts_jepa) against published deep learning baselines across matched telemetry channels.




Fig. 3. PR-AUC radar comparison across five specialized physical domains: Spacecraft Telemetry (MSL, SMAP, swan), Biomechanics (Daphnet), Cyber-Physical Systems (GHL, GECCO, Genesis), Traffic & Environmental (CalIt2, metro, room-occupancy), and Network Cyber (CICIDS). The physics-informed triad provides broad multi-domain coverage.

ture with simultaneous daily consumption cycles and high-frequency valve faults), dyadic filtering doubled operator_entropy_jepa performance from 0.1429 to 0.2904 Point-F1.

3) *Kurtosis-Adaptive Extreme Value Theory*: Real-world anomaly score tails frequently deviate from standard Gaussianity. By dynamically scaling the EVT threshold via excess kurtosis γ_2 , the calibrator suppresses false alarms during high-entropy volatile bursts while stabilizing detection margins across heterogeneous sensor regimes.

D. Operational Incident Evaluation: Event Metrics and Streaming Adaptation

While unadjusted Point-F1 prevents score inflation from naive point adjustment, point-level scoring can penalize slight temporal boundary offsets in operational telemetry. In physical systems, an anomaly represents a continuous operational incident; if an algorithm detects the failure with a sub-second phase delay or slight boundary misalignment, point-level metrics penalize every misaligned point as a false positive or false negative, depressing nominal precision even when incidents are detected.

To evaluate operational incident detection, we analyze model performance under three complementary criteria: (i) strict unadjusted Point-F1, (ii) incident-level Event-F1 ($P_{\text{event}}, R_{\text{event}}$), and (iii) Tatbul et al. (NeurIPS 2018) Range-Aware F1 ($F1_{\text{range}}$) evaluated with existence reward ($\alpha = 0$) and uniform positional overlap reward. Because sliding window prediction evaluates target horizon $x_{\text{tgt}} \in \mathbb{R}^{S \times K}$ ($S = 64$), detections operate with an operational lookahead buffer of 64 timesteps. Table III examines the operational trade-offs across representative domains:

As shown in Table III, evaluating anomaly detection at the incident level indicates:

1) **Incident Sensitivity vs. False Alarm Volume**: On Room Occupancy and SMAP, static EVT thresholding achieves 100% Event Recall across all tested models, with dense




Fig. 4. PR-AUC performance matrix across all 11 multi-sensor telemetry datasets. Values represent average precision across matched channels, with asterisks (*) denoting the best-performing model for each dataset.

scoring reaching **0.9167 Event-F1** and **84.6% Event Precision**. However, static calibration generates over 4,000 point-level false alarms on Room Occupancy due to unmodeled diurnal baseline drift.

- 2) **Streaming Adaptation Pareto Trade-off:** When dynamic rolling SPOT thresholding is applied to track baseline drift, false alarms decrease from 4,705 to 223 points on Room Occupancy (a $21\times$ reduction). This dramatic false alarm reduction introduces an operational trade-off: event recall drops from 100% to 63.6% (7 of 11 events detected), demonstrating the classic operating curve between alarm volume and event sensitivity.
- 3) **Boundary Alignment in Operational Telemetry:** On challenging channels such as SMAP P-3 and Callt2, low point-level F1 scores ($F1 \leq 0.13$) are driven by boundary temporal offset rather than complete failure; range-aware F1 reaches 0.3634 on SMAP. On Callt2, rapid intermittent pedestrian bursts result in lower event recall (4.8%–23.8%), demonstrating that fast transient pulses remain challenging for fixed-horizon window architectures.

VIII. CONCLUSION AND FUTURE WORK

In this work, we presented a Physics-Informed Joint-Embedding Predictive Architecture (PI-JEPA) framework for unsupervised anomaly detection in multivariate sensor telemetry. By replacing raw observation-space reconstruction with forward prediction in latent representation space, our approach

eliminates the threshold calibration collapse that causes reconstructive baselines (TimesNet, TranAD) to flood continuous streams with hundreds of thousands of false alarms under ambient sensor noise.

We embedded physical domain principles directly into latent representations through three formulations:

- 1) **Operator-Entropy JEPA:** Quantifying spectral stability and dissipation via dynamic spectral entropy across latent Koopman transition modes.
- 2) **Reynolds-Stress JEPA:** Aligning cross-channel second-moment fluctuation covariance tensors inspired by hydrodynamic Reynolds decomposition.
- 3) **Potential-Flow JEPA:** Enforcing conservative, irrotational trajectory dynamics by parameterizing the transition field as a scalar potential gradient.

Together with Coordinate Saliency Gating, Littlewood-Paley Dyadic Wavelet Shells, and Kurtosis-Adaptive Extreme Value Theory tail calibration, the proposed framework stabilizes streaming threshold calibration. Across 11 telemetry datasets (45 matched channels) evaluated under identical protocols, Operator-Entropy JEPA achieves the highest Dataset-Macro Point-F1 (0.2066, a **+43.2% relative improvement** over TimesNet and **+28.4%** over TranAD) and clears the analytical all-positive baseline (0.1893). In Pooled Micro-F1, Potential-Flow JEPA attains **0.3458** (over $3.2\times$ higher than TranAD and $5.3\times$ higher than TimesNet) while reducing false alarms by **8-fold**.

TABLE III

OPERATIONAL INCIDENT DETECTION COMPARISON ACROSS REPRESENTATIVE DOMAINS. DENSE SCORING PROVIDES HIGH INCIDENT RECALL (UP TO 100% ON ROOM OCCUPANCY AND SMAP), WHEREAS STREAMING DYNAMIC SPOT EXPOSES AN OPERATIONAL TRADE-OFF: IT REDUCES FALSE ALARM COUNTS BY UP TO $21\times$ (FROM 4,705 TO 223 POINTS) WHILE MODERATELY REDUCING EVENT RECALL TO 63.6%.

Dataset	Configuration	Pt-F1	Pt-Prec	Evt-F1	Evt-Prec	Evt-Rec	Rng-F1	False Pos.
<i>Room Occupancy (Environmental Sensing, 11 Ground-Truth Incidents)</i>								
operator_entropy	Standard (Middle, Static EVT)	0.3728	22.9%	0.5238	35.5%	100%	0.3708	3,842
operator_entropy	Sharp (Dense, Static EVT)	0.3329	20.0%	0.9167	84.6%	100%	0.3329	4,705
operator_entropy	Adaptive (Dense, Dynamic SPOT)	0.1273	29.9%	0.6364	63.6%	63.6%	0.3699	223 ($\downarrow 21\times$)
potential_flow	Sharp (Dense, Static EVT)	0.3431	20.7%	0.9167	84.6%	100%	0.3431	4,496
potential_flow	Adaptive (Dense, Dynamic SPOT)	0.0441	10.2%	0.5714	60.0%	54.5%	0.1636	291 ($\downarrow 15\times$)
<i>SMAP (Spacecraft Propulsion & Telemetry, Continuous Fault Segment)</i>								
operator_entropy	Standard (Middle, Static EVT)	0.0087	14.6%	0.1818	10.0%	100%	0.2266	35
operator_entropy	Sharp (Dense, Static EVT)	0.1274	27.3%	0.1538	8.3%	100%	0.3634	295
potential_flow	Adaptive (Dense, Dynamic SPOT)	0.0110	6.9%	0.3333	20.0%	100%	0.1213	108
<i>CallIt2 (Pedestrian Entrance Flow, 21 Transient Event Clusters)</i>								
operator_entropy	Standard (Middle, Static EVT)	0.0131	1.9%	0.0588	7.7%	4.8%	0.0216	53
operator_entropy	Adaptive (Dense, Dynamic SPOT)	0.1368	8.7%	0.3226	50.0%	23.8%	0.1274	357
potential_flow	Adaptive (Dense, Dynamic SPOT)	0.0903	14.3%	0.0909	100.0%	4.8%	0.0714	42

Future work includes:

- *Non-Conservative Field Extensions:* Relaxing strict irrotational constraints with parameter-controlled solenoidal vorticity components for dissipative systems with limit cycles.
- *Continuous Online Streaming Adaptation:* Developing memory-efficient, test-time momentum adaptation to track orbital thermal cycles and environmental seasonal drift without catastrophic forgetting.
- *Multi-Seed Uncertainty Quantification:* Scaling evaluations across extensive multi-seed benchmarks with confidence intervals and range-based evaluation metrics (e.g., VUS-PR).
- *Edge Hardware Quantization and Acceleration:* Compiling PI-JEPA encoders into 8-bit integer (INT8) formats for deployment on radiation-hardened spacecraft microcontrollers and ultra-low-power wearable edge devices.

REFERENCES

- [1] F. Meng, B. Yuan, and F. Gao, "Spacecraft telemetry anomaly detection based on deep learning: A review and benchmark," *IEEE Transactions on Aerospace and Electronic Systems*, vol. 56, no. 5, pp. 3863–3878, 2020.
- [2] R. Chalapathy and S. Chawla, "Deep learning for anomaly detection: A survey," *arXiv preprint arXiv:1901.03407*, 2019.
- [3] K. Hundman, V. Constantinou, C. Laporte, I. Colwell, and T. Soderstrom, "Detecting spacecraft anomalies using lstms and nonparametric dynamic thresholding," *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pp. 387–395, 2018.
- [4] M. P. Cuéllar, M. Santos, J. L. Alonso, M. Fabregas, and R. Farías, "Explainable anomaly detection in spacecraft telemetry," *Engineering Applications of Artificial Intelligence*, vol. 133, p. 108119, 2024.
- [5] J. Xu, H. Wu, J. Wang, and M. Long, "Anomaly transformer: Time series anomaly detection with association discrepancy," in *International Conference on Learning Representations (ICLR)*, 2022.
- [6] H. Wu, T. Hu, Y. Liu, H. Zhou, J. Wang, and M. Long, "Timesnet: Temporal 2d-variation modeling for general time series analysis," in *International Conference on Learning Representations (ICLR)*, 2023.
- [7] C. U. Carmona, F.-X. Aubet, V. Flunkert, and J. Gasthaus, "Neural contextual anomaly detection for time series," in *International Conference on Learning Representations (ICLR)*, 2022.
- [8] Y. LeCun, "A path towards autonomous machine intelligence," *OpenReview Preprint*, 2022.
- [9] M. Assran, Q. Duval, I. Misra, P. Bojanowski, P. Vincent, M. Rabbat, Y. LeCun, and N. Ballas, "Self-supervised learning from images with a joint-embedding predictive architecture," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023, pp. 15 619–15 629.
- [10] S. L. Brunton, J. L. Proctor, and J. N. Kutz, "Discovering governing equations from data by sparse identification of nonlinear dynamical systems," *Proceedings of the National Academy of Sciences*, vol. 113, no. 15, pp. 3932–3937, 2016.
- [11] R. T. Chen, Y. Rubanova, J. Bettencourt, and D. K. Duvenaud, "Neural ordinary differential equations," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 31, 2018, pp. 6571–6583.
- [12] S. L. Brunton, B. W. Brunton, J. L. Proctor, and J. N. Kutz, "Koopman invariant subspaces and finite-dimensional approximations of nonlinear dynamical systems," *Journal of Nonlinear Science*, vol. 26, pp. 1131–1162, 2016.
- [13] B. Lusch, J. N. Kutz, and S. L. Brunton, "Deep learning for universal linear embeddings of nonlinear dynamics," *Nature Communications*, vol. 9, no. 1, p. 4950, 2018.
- [14] Y. Liu, C. Li, J. Wang, and M. Long, "Koopman: Learning non-stationary time series dynamics with Koopman predictors," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 36, 2023, pp. 13 480–13 492.
- [15] S. Tuli, G. Casale, and N. R. Jennings, "Tranad: Deep transformer networks for anomaly detection in multivariate time series data," *Proceedings of the VLDB Endowment*, vol. 15, no. 6, pp. 1201–1214, 2022.
- [16] Y. Yang, C. Zhang, T. Zhou, Q. Gao, Q. Wen, and Y. Liang, "Dcdetector: Dual attention contrastive representation learning for time series anomaly detection," in *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2023, pp. 3033–3045.
- [17] A. Bardes, Q. Duval, C. Feichtenhofer, I. Laptev, D. Haziza, V. Khalidov, N. Ballas, and Y. LeCun, "V-jepa: Video joint-embedding predictive architecture," *arXiv preprint arXiv:2404.08471*, 2024.
- [18] J. Zbontar, L. Jing, I. Misra, Y. LeCun, and S. Deny, "Barlow twins: Self-supervised learning via redundancy reduction," in *International Conference on Machine Learning (ICML)*, 2021, pp. 12 310–12 320.
- [19] A. Ermolov, A. Siarohin, E. Sangineto, and N. Sebe, "Whitening for self-supervised representation learning," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, 2021, pp. 8315–8327.
- [20] A. Bardes, J. Ponce, and Y. LeCun, "Vicreg: Variance-invariance-covariance regularization for self-supervised learning," in *International Conference on Learning Representations (ICLR)*, 2022.
- [21] Y. He, Y. Wen, X. Wang, and T. Ma, "Sc-jepa: Stabilizing latent predictive learning for time-series anomaly prediction," *arXiv preprint arXiv:2602.04643*, 2026.
- [22] K. Woldeesenbet and A. Woldeesenbet, "T-sar-jepa: Self-supervised temporal anomaly detection in sar amplitude stacks via latent prediction," *arXiv preprint arXiv:2606.05700*, 2026.
- [23] J. Lee and S. Sim, "Cf-jepa: Mask-free forward prediction with asym-

- metric encoder utilization for time-series representation learning,” *arXiv preprint arXiv:2606.07031*, 2026.
- [24] A. M. Girgis, A. Valcarce, and M. Bennis, “Time-series jepa for predictive remote control under capacity-limited networks,” *IEEE Internet of Things Journal*, vol. 13, no. 7, pp. 14 617–14 629, 2026.
 - [25] A. Fertig, K. C. Sekaran, L. Balasubramanian, and M. Botsch, “On-line monitoring framework for automotive time series data using jepa embeddings,” *arXiv preprint arXiv:2602.09985*, 2026.
 - [26] M. Raissi, P. Perdikaris, and G. E. Karniadakis, “Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations,” *Journal of Computational Physics*, vol. 378, pp. 686–707, 2019.
 - [27] S. Wang, Y. Teng, and P. Perdikaris, “Understanding and mitigating gradient pathologies in physics-informed neural networks,” *SIAM Journal on Scientific Computing*, vol. 43, no. 5, pp. A3055–A3081, 2021.
 - [28] A. Siffer, P.-A. Fouque, A. Termier, and C. Largouet, “Anomaly detection in streams with extreme value theory,” in *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2017, pp. 1067–1075.
 - [29] M. Bächlin, M. Plotnik, D. Roggen, I. Maidan, J. M. Hausdorff, N. Giladi, and G. Tröster, “Wearable assistant for parkinson’s disease patients with the freezing of gait symptom,” *IEEE Transactions on Information Technology in Biomedicine*, vol. 14, no. 2, pp. 436–446, 2010.
 - [30] S. Schmidl, P. Wenig, and T. Papenbrock, “Anomaly detection in time series: a comprehensive evaluation,” *Proceedings of the VLDB Endowment*, vol. 15, no. 9, pp. 1779–1797, 2022.
 - [31] Y. Zhou, C. Brif, and I. Lourentzou, “mTSBench: A Comprehensive Benchmark for Multi-Domain Time Series Anomaly Detection,” *Transactions on Machine Learning Research (TMLR)*, 2026, arXiv:2506.21550.
 - [32] R. A. Angryk, P. C. Martens, B. Aydin, D. Kempton, S. S. Mahajan, S. Basodi, A. Ahmadzadeh, S. Boubrahimi, S. M. Hamdi, M. A. Schuh *et al.*, “Multivariate time series dataset for space weather data analytics,” *Scientific Data*, vol. 7, no. 1, p. 227, 2020.
 - [33] L. M. Candanedo and V. Feldheim, “Accurate occupancy detection of an office room from light, temperature, humidity and co2 measurements using statistical learning models,” *Energy and Buildings*, vol. 112, pp. 28–39, 2016.
 - [34] A. Asuncion and D. Newman, “UCI Machine Learning Repository: CalIt2 Building People Counts Data Set,” <https://archive.ics.uci.edu/dataset/179/calit2+building+people+counts>, 2007.
 - [35] S. Kim, K. Choi, J. Choi, B. Lee, and S. Yoon, “Towards a rigorous evaluation of time-series anomaly detection,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 36, no. 7, 2022, pp. 7194–7201.
 - [36] R. Wu and E. Keogh, “Current time series anomaly detection benchmarks are flawed and are creating the illusion of progress,” *IEEE Transactions on Knowledge and Data Engineering*, vol. 35, no. 3, pp. 2421–2434, 2023.