

Physics-Inspired Joint-Embedding Predictive Architectures for Multivariate Time-Series Anomaly Detection: Conditioning Residuals for Automated Thresholding

Andre Williams	Imadeldin Mahgoub	Hamid Akbarian
<i>Florida Atlantic University</i>	<i>Florida Atlantic University</i>	<i>Florida Atlantic University</i>
<i>Department of EECS</i>	<i>Department of EECS</i>	<i>Department of EECS</i>
<i>Boca Raton, US</i>	<i>Boca Raton, US</i>	<i>Boca Raton, US</i>
<i>andrewwilliam2018@fau.edu</i>	<i>mahgoubi@fau.edu</i>	<i>hakbaria@fau.edu</i>

Abstract—Unsupervised anomaly detection in multivariate telemetry is a critical operational requirement for mission safety across aerospace systems, industrial infrastructure, and smart cyber-physical environments. While modern reconstructive architectures such as TimesNet and TranAD achieve competitive continuous ranking performance (PR-AUC 0.2585–0.2612), their raw observation-space residuals exhibit severe sensitivity to non-stationary sensor jitter. When deployed under automated streaming threshold calibration via Extreme Value Theory (EVT / SPOT) without ground-truth labels, reconstructive models suffer severe calibration collapse, flooding telemetry streams with hundreds of thousands of false alarms (TimesNet: 1,047,144 false alarms, 0.0647 Micro-F1; TranAD: 452,621 false alarms, 0.1072 Micro-F1). To resolve this failure, we demonstrate that Joint-Embedding Predictive Architectures (JEPAs) forecasting in abstract representation space yield latent predictive residuals that are substantially better conditioned for automated thresholding than raw reconstruction errors. We investigate three physics-inspired inductive regularizers that embed domain principles into the latent manifold: *Operator-Entropy JEPA*, regularizing latent Koopman operators via scale-free spectral entropy and Bregman energy barriers to penalize modal collapse; *Reynolds-Stress JEPA*, aligning cross-coordinate fluctuation covariances inspired by hydrodynamic Reynolds decomposition; and *Potential-Flow JEPA*, constraining trajectory updates to curl-free scalar potential gradients ($\hat{v} = -\nabla\Phi$). Combined with causal coordinate saliency gating with zero target-leakage shielding, Littlewood-Paley dyadic frequency shells, and kurtosis-adaptive EVT tail calibration, these components stabilize streaming decision thresholds. Evaluated across 45 matched telemetry channels from 11 benchmarks under an identical calibration protocol, *Operator-Entropy JEPA* achieves the highest Dataset-Macro Point-F1 (0.2066, a +43.2% relative gain over TimesNet and +28.4% over TranAD) and clears the analytical All-Positive reference baseline (0.1893). In Pooled Micro Point-F1, *Potential-Flow JEPA* attains 0.3458, reducing false alarms by 8.0-fold against TimesNet (130,237 vs 1,047,144) and 3.5-fold against TranAD (452,621). These results indicate that while continuous ranking remains balanced across model families, latent predictive residual conditioning governs false-alarm suppression in automated, threshold-calibrated anomaly detection.

Index Terms—Multivariate Time-Series, Anomaly Detection, Joint-Embedding Predictive Architecture (JEPA), Physics-Inspired Machine Learning, Threshold Calibration, Extreme Value Theory, Residual Conditioning.

I. INTRODUCTION

Autonomous anomaly detection in multi-sensor telemetry monitors operational health across aerospace systems, Earth observation, cyber-physical infrastructure, and wearable digital healthcare [1], [2]. In remote sensing deployments (such as the NASA Soil Moisture Active Passive (SMAP) satellite, Mars Science Laboratory (MSL) rovers, industrial IoT facilities, and wearable biometric sensor arrays), dozens to hundreds of continuous telemetry channels concurrently track electrical power buses, thermal gradients, mechanical vibrations, and kinematics [3]. When sensor faults or operational breakdowns emerge at the edge, monitoring systems must identify dynamical anomalies before hardware damage occurs [4].

A. Challenges in Modeling Multi-Sensor Telemetry

From a physical perspective, multi-sensor telemetry reflects the state trajectory of an underlying continuous physical system $\dot{x}(t) = f(x(t))$. However, attempting to model dynamical equations directly in raw observation space faces four core challenges on edge devices. First, estimating instantaneous time derivatives $\dot{x} \approx \frac{\Delta x}{\Delta t}$ directly on discrete, noise-corrupted sensor streams amplifies high-frequency measurement noise, destabilizing numerical trajectory tracking. Second, remote sensing telemetry captures only a noisy surface projection of unmeasured thermodynamic, electrical, and environmental states, violating full-state observability. Third, in large sensor arrays ($K \geq 25$, such as NASA SMAP with $K = 25$ and MSL with $K = 55$), cross-channel dependencies and non-linear interactions cannot be adequately represented by rigid functional forms in raw signal space. Finally, edge IoT and satellite systems undergo discrete operational mode transitions (e.g., orbital day/night thermal cycling, thruster firings, Parkinsonian gait freezing, cloud server garbage collection) that cause continuous distribution drift.

B. Limitations of Deep Anomaly Detectors

Deep learning architectures provide data-driven alternatives to explicit equation modeling, yet they face three primary

limitations. First, generative models (e.g., autoencoders, VAEs, masked Transformers) minimize point-level L_1/L_2 errors in raw signal space [3], [5], [6]. Because reconstruction loss penalizes every discrepancy equally, autoencoders allocate capacity to fitting high-frequency sensor jitter rather than learning nominal dynamical transitions. Second, contrastive frameworks such as Neural Contextual Anomaly Detection (NCAD) [7] train encoders via Contextual Outlier Exposure (COE) using heuristic synthetic perturbations (e.g., Gaussian noise spikes, sequence splicing). However, real physical faults manifest as subtle deviations from nominal dynamics; synthetic injections distort embedding geometry, inducing artificial decision boundaries that generate false alarms on continuous telemetry. Third, in high-dimensional observation space, even after per-channel standardization, Euclidean reconstruction errors ($\|x - \hat{x}\|_2^2 = \sum_k (x_k - \hat{x}_k)^2$) accumulate stochastic residual variance across dozens of irrelevant or noisy sensors. Under automated streaming thresholding such as Peaks-Over-Threshold (SPOT), this residual accumulation yields heavy, unstable tail distributions that distort extreme-value calibrators, triggering false alarm flooding in operational telemetry.

C. Our Approach: Physics-Inspired Joint-Embedding Predictive Architectures

To substantially suppress sensitivity to observation-space noise without relying on synthetic contrastive perturbations, we formulate a family of *Physics-Inspired Joint-Embedding Predictive Architectures* (PI-JEPA) for multi-sensor telemetry. Instead of reconstructing raw, jitter-corrupted sensor observations $\hat{x} \approx x$, a JEPA maps past observations x_{ctx} into an abstract representation space and predicts target representations z_{tgt} directly [8], [9]. By operating in latent representation space, the architecture filters out high-frequency stochastic sensor noise that lacks long-range dynamical predictability, conditioning prediction residuals for automated tail calibration.

However, unconstrained latent prediction alone can be underspecified in complex multi-sensor telemetry. To incorporate domain inductive biases into self-supervised representation learning, we investigate three physics-inspired JEPA architectures featuring inductive dynamical regularizers. The first formulation, *Operator-Entropy JEPA*, parameterizes a latent transition matrix regularized by scale-free spectral entropy and Bregman energy barriers during training to penalize dimensional collapse and preserve linear modal expressivity in multi-sensor dynamics. The second, *Reynolds-Stress JEPA*, adapts hydrodynamic Reynolds decomposition ($z = \bar{z} + z'$) to align second-order fluctuation covariance across latent coordinates, capturing structural anomalies that disrupt multi-coordinate covariance balance. The third, *Potential-Flow JEPA*, enforces conservative latent trajectory transitions by parameterizing the transition field as the gradient of a scalar potential ($\hat{v} = -\nabla\Phi$); because gradient fields are irrotational by construction, non-conservative dynamical anomalies produce pronounced predictive discrepancies.

To stabilize learning across heterogeneous high-dimensional sensor arrays, all three architectures incorporate strictly Causal Coordinate Saliency Gating with zero target-leakage shielding

to isolate active channels, Dyadic Wavelet Shells to decouple high-frequency transients from baseline drift, and a Kurtosis-Adaptive Extreme Value Theory (EVT / SPOT) calibrator that modulates tail thresholds against non-Gaussian kurtosis.

D. Contributions

This work makes four principal contributions:

1) Latent Predictive Residual Conditioning for Automated Thresholding: We demonstrate that forecasting future telemetry states in abstract representation space rather than reconstructing raw observation waveforms produces residual distributions that are substantially better conditioned for automated extreme-value calibration, resolving the catastrophic false alarm flooding of reconstructive baselines.

2) Physics-Inspired Inductive Regularizers: We derive and implement three physics-inspired architectures (Operator-Entropy JEPA, Reynolds-Stress JEPA, and Potential-Flow JEPA) and benchmark them systematically alongside an unconstrained control (TS-JEPA), showing that physical conservation and dissipation principles provide targeted, regime-specific regularizing benefits without overclaiming universal ranking dominance.

3) Causal Gating and Structural Inductive Biases: We introduce Causal Coordinate Saliency Gating with provable zero target-window leakage ($\partial g / \partial x_{\text{tgt}} \equiv 0$), Littlewood-Paley Dyadic Wavelet Shells for multi-scale frequency separation, and Kurtosis-Adaptive EVT tail calibration to deliver automated streaming decisions without transductive test-label tuning.

4) Multi-Dataset Benchmark and Metric Decoupling: We evaluate across 11 telemetry datasets (45 matched streams), reporting unadjusted Point-F1, Event-F1, Range-F1, PA-F1, PR-AUC, and the analytical All-Positive reference baseline. We demonstrate that while continuous ranking remains balanced across architectures (PR-AUC 0.2541–0.2612), latent predictive residual conditioning achieves up to +43.2% relative improvement in Dataset-Macro Point-F1 (where Operator-Entropy JEPA is the sole model to exceed the analytical All-Positive baseline) and reduces false alarms by 8.0-fold against TimesNet (130,237 vs 1,047,144) and 3.5-fold against TranAD (452,621).

The remainder of this paper is structured as follows. Section II reviews related work in reconstructive, contrastive, and joint-embedding anomaly detection. Section III details the 11-dataset multi-sensor telemetry benchmark suite. Section IV formalizes the PI-JEPA Triad, causal coordinate saliency gating, dyadic wavelets, and kurtosis-adaptive EVT. Section V establishes the experimental protocol and multi-tier evaluation framework. Section VI evaluates architectural ablations and inductive layer contributions. Section VII provides domain-by-domain empirical analysis and qualitative physical case studies. Finally, Section VIII concludes the paper.

II. RELATED WORK

Unsupervised time-series anomaly detection spans dynamical equation discovery, reconstructive generative modeling,

contrastive representation learning, and predictive world models. This section reviews these approaches and positions our Physics-Inspired JEPA framework.

A. Dynamical Systems Modeling and Koopman Operators

The physical modeling of multi-sensor telemetry has long sought to capture system behavior as continuous dynamical trajectories governed by differential equations $\dot{x}(t) = \mathbf{f}(x(t))$ [10]. One prominent approach, Neural Ordinary Differential Equations (Neural ODEs) [11], parameterizes continuous velocity fields using neural networks $\frac{dx}{dt} = f_\theta(x(t), t)$ and integrates trajectories numerically. While Neural ODEs learn flexible vector fields, fitting them directly in raw observation space forces numerical solvers to integrate through high-frequency stochastic sensor noise, leading to numerical stiffness, slow training, and sensitivity to measurement jitter.

In parallel, Koopman operator theory and Dynamic Mode Decomposition (DMD) [12]–[14] analyze non-linear dynamical systems by learning linear transition operators that act on observable functions. Modern deep Koopman frameworks (e.g., Koopa [14]) learn observable embeddings for non-stationary forecasting. In anomaly detection, physical disruptions manifest as spectral decoherence or sudden changes in modal dissipation. Rather than operating on uncalibrated raw signals, our framework parameterizes latent linear transitions regularized by Koopman spectral entropy during training directly in representation space.

B. Reconstruction-Based Detection

The predominant deep learning approach in time-series anomaly detection relies on learning nominal patterns through raw signal reconstruction. Hundman et al. [3] introduced LSTM-based prediction coupled with non-parametric dynamic thresholding for NASA spacecraft telemetry (SMAP/MSL). Modern architectures incorporate specialized attention and frequency-domain transformations. For example, Anomaly Transformer [5] introduces Anomaly-Attention to quantify the association discrepancy between local prior Gaussian kernel associations and global series self-attention. TimesNet [6] decomposes 1D multivariate time series into 2D temporal variations using Fast Fourier Transform (FFT) periodicity discovery and 2D Inception blocks. Similarly, TranAD [15] employs a two-phase adversarial Transformer with sub-series attention and focus scoring to amplify subtle reconstruction errors.

Despite their empirical popularity, reconstructive methods share a fundamental vulnerability: signal-level noise memorization. Because loss functions optimize point-level L_1 or L_2 differences in raw signal space, reconstructive networks are biased toward fitting high-frequency stochastic noise rather than learning invariant dynamical laws. In addition, during sustained anomalies, corrupted observations contaminate the context window, causing the reconstructor to adapt and reconstruct the anomaly as the nominal state.

C. Contrastive Learning and Representation-Based Detection

To avoid reconstructing raw noise, contrastive methods learn discriminative representations by separating nominal sequences from anomalous perturbations. Neural Contextual Anomaly Detection (NCAD) [7] uses Temporal Convolutional Networks (TCNs) trained with Contextual Outlier Exposure (COE), injecting artificial synthetic anomalies (Gaussian noise, point spikes, trend shifts) into training windows to create negative pairs. However, synthetic injections distort the physical dynamical manifold and fail to generalize to subtle real-world physical faults (such as gradual thermal dissipation or Parkinsonian gait freezing). Alternatively, DCdetector [16] replaces reconstruction with multi-scale dual-attention contrastive representation learning, maximizing cross-scale representation consistency. While DCdetector eliminates artificial injection heuristics, contrastive objectives require careful scale selection and representation pooling.

D. Self-Supervised Time-Series Representation Learning and JEPAs

Joint-Embedding Predictive Architectures (JEPA), originally proposed by LeCun [8] and demonstrated in computer vision by I-JEPA [9] and V-JEPA [17], predict target representations from context in latent space. Non-contrastive regularizers such as Barlow Twins [18], W-MSE [19], and VICReg [20] prevent representation collapse without negative mining.

Recent works have extended self-supervised and predictive representation learning principles to time-series sequences. TS2Vec [21] learns multi-scale representations through contextual and temporal contrasting, while Eldele et al. [22] enforce temporal-contextual cross-view consistency. Systematic surveys by Zhang et al. [23] document the emergence of self-supervised representation architectures across multi-sensor streams. However, existing time-series representation frameworks lack explicit physical dynamical priors, leaving latent trajectories unanchored to physical conservation and spectral dissipation principles.

E. Physics-Inspired Deep Learning and Structured Inductive Biases

Physics-Informed Neural Networks (PINNs) [24], alongside Hamiltonian [25] and Lagrangian Neural Networks [26], integrate dynamical principles into neural architectures. However, while Hamiltonian networks enforce strict energy conservation for closed mechanical systems, standard PINNs face two primary limitations in high-dimensional sensor telemetry. First, jointly optimizing data fidelity and PDE residual loss leads to stiff, competing gradient dynamics, often causing training instability or failure to converge in high noise regimes [27]. Second, complex IoT and aerospace telemetry (e.g., SMAP, MSL, and cyber-physical water networks) rarely follow a single closed-form PDE whose analytical derivatives can be evaluated across dozens of uncalibrated sensors. Rather than enforcing analytical differential equations through auxiliary training penalties, our framework incorporates physical conservation, turbulent transport, and spectral dissipation principles

directly into latent representations as inductive regularizers, avoiding multi-loss balancing conflicts.

F. Positioning of the Proposed Framework

The proposed Physics-Inspired JEPA framework bridges joint-embedding predictive learning and physical domain mechanics through four distinct structural designs. First, by forecasting future representations z_{tgt} rather than raw observations $\hat{x}$, our models circumvent high-frequency noise drowning and sensor jitter memorization that hobble reconstructive baselines (TimesNet, TranAD). Second, we investigate three physics-inspired inductive formulations: Operator-Entropy (Koopman spectral dissipation), Reynolds-Stress (turbulent momentum transport), and Potential-Flow (conservative irrotational field theory), addressing cyclic, turbulent, and conservative physical dynamics. Third, we incorporate Causal Coordinate Saliency Gating with zero target leakage to prevent active channel dilution in high-dimensional arrays, and Littlewood-Paley Dyadic Wavelets to isolate transient impulses from low-frequency operational drift. Finally, we deploy Extreme Value Theory (EVT / SPOT) [28] modulated by dynamic sample kurtosis, establishing decision boundaries on operational telemetry without transductive test-label tuning.

III. MULTI-SENSOR IoT AND REMOTE SENSING BENCHMARK SUITE

To evaluate our Physics-Inspired JEPA framework across diverse edge, cyber-physical, and orbital sensing dynamics, our benchmark suite comprises 11 datasets spanning 45 strictly matched continuous telemetry streams sourced from NASA spacecraft archives [3], the UCI Machine Learning Repository [29], [30], TimeEval [31], and mTSBench [32]. We organize these streams into four physical sensing domains:

A. Domain 1: Satellite and Spacecraft Telemetry (13 Streams)

This domain covers high-reliability orbital and deep-space telemetry. The *NASA SMAP Earth Observation Satellite* [3] (6 evaluated streams: A-1, A-2, A-7, D-3, P-3, E-1) provides continuous telemetry values alongside one-hot command vectors monitoring power, thermal, and attitude control subsystems. The *NASA MSL Curiosity Rover* [3] (6 evaluated streams: C-1, C-2, D-14, D-15, D-16, M-1) tracks robotic actuator, joint drive, and environmental streams during planetary surface operations. Finally, the *SWAN-SF Space Weather Telemetry* [33] (1 stream) provides solar photospheric vector magnetogram telemetry from NASA SDO/HMI, tracking active region magnetic field perturbations during solar eruptive events.

B. Domain 2: Wearable Biometrics and Kinematics (6 Streams)

This domain evaluates human kinematic telemetry via the *Daphnet Wearable Body Sensor Network* [29] (6 evaluated streams: S01R01E1 to S03R02E0). Telemetry consists of 64 Hz tri-axial accelerometry recorded from wearable sensor

nodes placed on the shank, thigh, and trunk to track Freezing-of-Gait (FoG) motor breakdown events in Parkinson's disease patients. At 64 Hz, the suspect horizon $S = 64$ corresponds to a 1.0-second lookahead window.

C. Domain 3: Industrial Processes and Cyber-Physical Systems (16 Streams)

This domain covers continuous industrial processes and fluid transportation. The *GHL Gasoil Heating Loop* (14 benchmark streams) captures thermal and fluid process loop telemetry across temperature, flow rate, and valve pressure dynamics. The *Genesis Automation Testbed* (1 stream) tracks manufacturing automation telemetry from an industrial pick-and-place demonstrator during mechanical handling faults. The *GECCO Water Distribution IoT* (1 stream) monitors physical water distribution infrastructure through continuous chemical and hydrological measurements at 1-minute resolution (where $S = 64$ represents a 64-minute operational horizon).

D. Domain 4: Smart Infrastructure and Environmental Sensing (10 Streams)

This domain evaluates urban infrastructure and cyber-physical networks. The *Room-Occupancy Environmental IoT* dataset [34] (2 streams) records ambient environmental conditions (CO₂, light, temperature, humidity) tracking room-occupancy state switches. The *CICIDS Edge Network Telemetry* (6 streams) captures network flow dynamics, tracking gateway protocol shifts and security anomalies including DDoS bursts and port scans. The *Callt2 Building People-Flow Telemetry* [30] (1 stream) records continuous pedestrian counts via optical beam sensors at the UC Irvine Calit2 facility entrance. Finally, the *Metro Highway Traffic* dataset (1 stream) tracks interstate highway traffic volume on Minnesota Interstate I-94, where anomalies correspond to major traffic disruption and incident events.

Each benchmark stream provides pre-split training sequences (consisting exclusively of nominal operational behavior) and testing sequences containing labeled ground-truth anomalies. Continuous telemetry channels are standardized via training statistics ($\mu_{\text{train}}, \sigma_{\text{train}}$), flooring standard deviations at $\max(\sigma, 10^{-2})$ to prevent division by zero during flat-line operational phases. While full mission archives contain larger entity counts (the complete SMAP archive includes 55 telemetry entities and MSL includes 27 entities across various mission subsystems), our benchmark evaluates standardized subsets commonly studied in the multi-sensor benchmark literature [3], [31] to ensure strict matching and reproducible runtime conditions across all baselines.

Table I provides the technical specifications across all 11 benchmark datasets and 45 evaluated continuous telemetry streams, totaling 2,079,507 nominal training timesteps and 2,986,574 test timesteps (over 5.06 million observations). We explicitly document the physical domain, channel dimension K , sampling resolution, sequence lengths, anomaly contamination ratio $\pi\%$, anomaly failure modes, and partition protocols to ensure experimental reproducibility.

TABLE I

TECHNICAL SPECIFICATIONS OF THE MULTI-SENSOR IOT AND REMOTE SENSING BENCHMARK SUITE (45 STRICTLY MATCHED TELEMETRY STREAMS ACROSS 11 DIVERSE MODALITIES, COMPRISING OVER 5.06 MILLION TIME-SERIES OBSERVATIONS).

Dataset	Physical Sensing Domain	Streams	Dim K	N_{train}	N_{test}	Anom % (π)	Primary Anomaly Typologies	Data Partition Protocol
Domain 1: Satellite and Spacecraft Remote Telemetry								
NASA SMAP [3]	Soil Moisture Satellite Subsystems	6	25	17,022	50,834	15.18%	Subsystem power drops, thermal drifts, cmd faults	NASA JPL nominal/incident pre-split
NASA MSL [3]	Curiosity Rover Robotic Actuators	6	55	12,331	13,566	22.88%	Actuator stalls, motor spikes, sensor freezes	NASA JPL nominal/incident pre-split
SWAN-SF [33]	Solar Photospheric Magnetograms	1	38	19,672	72,000	27.92%	Active region magnetic field eruptive bursts	Chronological temporal split
Domain 2: Wearable Biometrics and Kinematics								
Daphnet FoG [29]	Wearable Body Sensor Accelerometry	6	9	48,240	217,344	13.34%	Freezing-of-Gait motor breakdown, limb tremors	Subject-isolated session runs
Domain 3: Industrial Processes and Cyber-Physical Systems								
GHL Heating Loop	Thermal and Fluid Transport Loop	14	16	1,645,137	1,154,014	1.35%	Valve sticking, heater failures, pressure surge	Independent physics simulation runs
Genesis Automation	Robotic Pick-and-Place Demonstrator	1	18	3,244	9,332	0.30%	Mechanical manipulation jams, gripper slip	Chronological operational split
GECCO Water IoT	Hydraulic and Water Distribution	1	9	27,560	83,113	1.54%	Chemical contamination, line bursts, pH drops	Chronological infrastructure split
Domain 4: Smart Infrastructure and Environmental Sensing								
Room-Occupancy [34]	Ambient Environmental Sensing	2	5	2,565	14,317	19.31%	Room-occupancy transitions, airflow anomalies	Chronological test split
CICIDS Edge Network	Edge Gateway Flow Telemetry	6	72	278,722	1,343,920	23.91%	DDoS volume surges, brute force, port scans	Temporal attack partition
Call2 [30]	Optical People-Flow Counter	1	2	940	4,032	3.42%	Crowd attendance surges, building ingress bursts	Chronological event split
Metro Traffic	Highway Loop Detector Volume	1	5	24,072	24,102	0.13%	Highway congestion, accidents, severe slowdowns	Chronological calendar split
Total / Benchmark Suite	4 Diverse Physical Domains	45	2-72	2,079,507	2,986,574	13.72% (avg)	Structural, contextual, and point faults	Pure inductive nominal training

IV. METHODOLOGY: THE PHYSICS-INSPIRED GEOMETRIC JEPA TRIAD

This section formalizes the mathematical formulation and inductive biases of the proposed framework.

A. Problem Formulation and Latent Space Prediction

Let $\mathbf{X} \in \mathbb{R}^{T \times K}$ denote a continuous multivariate time-series sequence consisting of K observed sensor variables over T timesteps. The sequence is partitioned into historical context $x_{\text{ctx}} = \mathbf{X}[t - C : t, :] \in \mathbb{R}^{C \times K}$ and suspect/target horizon $x_{\text{tgt}} = \mathbf{X}[t : t + S, :] \in \mathbb{R}^{S \times K}$, where C and S denote the respective temporal horizons.

Existing deep learning baselines (e.g., TimesNet, TranAD, Anomaly Transformer) frame anomaly detection as pointwise reconstruction:

$$\min_{\theta} \mathbb{E} [\|\mathbf{X} - \mathcal{D}_{\theta}(\mathcal{E}_{\theta}(\mathbf{X}))\|_F^2]. \quad (1)$$

In complex multi-sensor telemetry, ambient high-frequency stochastic jitter dominates the Frobenius norm, causing models to waste representational bandwidth reconstructing sensor noise rather than modeling system state dynamics (leading to noise drowning).

To overcome this, we adopt the Joint-Embedding Predictive Architecture (JEPA) framework [8]. A context encoder E_{θ} maps $x_{\text{ctx}} \rightarrow z_{\text{ctx}} \in \mathbb{R}^D$, while an Exponential Moving Average (EMA) momentum target encoder E_{ϕ} maps $x_{\text{tgt}} \rightarrow z_{\text{tgt}} \in \mathbb{R}^D$, updated via:

$$\phi_k \leftarrow m_k \phi_{k-1} + (1 - m_k) \theta_k, \quad m_k \in [0.996, 0.9995]. \quad (2)$$

The core objective is to predict target representations in latent space under physical dynamical constraints.

B. The Physics-Inspired JEPA Triad

Rather than relying on unconstrained Euclidean regression, we incorporate physical domain regularizers through three complementary geometric formulations:

1) Operator-Entropy JEPA (Spectral Dynamics & Modal Dispersion): In dynamical systems theory, non-linear trajectories can be analyzed via linear operators acting on observable functions (Koopman operator theory [?], [?]). In our architecture, the encoder maps the observation window to a finite-dimensional latent observable vector $z \in \mathbb{R}^D$. We parameterize the predictive transition between context and target latent states via a learned transition operator $\mathbf{K} \in \mathbb{R}^{D \times D}$:

$$\hat{z}_{\text{tgt}} = \mathbf{K} z_{\text{ctx}}. \quad (3)$$

To prevent the operator from collapsing onto a degenerate low-rank subspace (dimensional collapse) or degenerating to zero (energy collapse), we formulate an operator-entropy regularizer that decouples spectral uniformity from total operator energy.

Let $\sigma_1 \geq \dots \geq \sigma_D \geq 0$ denote the singular values of $\mathbf{K}$, with corresponding eigenvalues $\lambda_i = \sigma_i^2$ of the symmetric Gram matrix $\mathbf{K}\mathbf{K}^T$. The total operator Frobenius energy is $\mathcal{E}(\mathbf{K}) = \|\mathbf{K}\|_F^2 = \sum_{i=1}^D \lambda_i$. For any non-zero operator ($\mathcal{E}(\mathbf{K}) > 0$), we define the scale-free spectral density distribution $\mathbf{p} = [p_1, \dots, p_D]^T$ over latent modes:

$$p_i = \frac{\lambda_i(\mathbf{K}\mathbf{K}^T)}{\sum_{j=1}^D \lambda_j(\mathbf{K}\mathbf{K}^T)} = \frac{\sigma_i^2}{\|\mathbf{K}\|_F^2}, \quad \sum_{i=1}^D p_i = 1, \quad p_i \geq 0. \quad (4)$$

We omit isotropic diagonal loading ($\frac{\epsilon}{D}\mathbf{I}$) from the Gram matrix before normalization; adding it would cause $\mathbf{K} \rightarrow \mathbf{0}$ to yield an artificial identity spectrum with maximal entropy, spuriously rewarding total collapse. Instead, we evaluate the Von Neumann spectral entropy directly on the scale-free modal probabilities with the standard analytic convention $0 \log 0 = 0$:

$$\mathcal{H}_{\text{spectral}}(\mathbf{K}) = - \sum_{i: p_i > 0} p_i \log p_i. \quad (5)$$

The *Spectral Entropy Deficiency* quantifies the Kullback-Leibler divergence from the uniform modal equipartition $\mathbf{u} = [1/D, \dots, 1/D]^T$:

$$\Delta \mathcal{H}(\mathbf{K}) = \log D - \mathcal{H}_{\text{spectral}}(\mathbf{K}) = \mathcal{D}_{\text{KL}}(\mathbf{p} \parallel \mathbf{u}) \geq 0. \quad (6)$$

Here, $\Delta \mathcal{H} = 0$ if and only if all modes carry equal energy ($\sigma_1 = \dots = \sigma_D$), while dimensional collapse to rank 1 yields maximal deficiency $\Delta \mathcal{H} = \log D$.




Fig. 1. Architectural overview of the proposed Physics-Inspired Joint-Embedding Predictive Architecture (JEPA) framework. The historical context window $\mathbf{x}_{\text{ctx}}$ and suspect target window $\mathbf{x}_{\text{tgt}}$ are projected through multi-scale dyadic TCN encoders into latent representations $\mathbf{z}_{\text{ctx}}$ and $\mathbf{z}_{\text{tgt}}$. Rather than reconstructing raw observation-space waveforms (which causes false alarm flooding under ambient sensor jitter), the predictive core operates strictly in latent representation space, governed by three complementary physical inductive biases: Koopman operator spectral entropy, hydrodynamic Reynolds stress transport, and conservative potential flow field theory. Prediction discrepancy is calibrated online via streaming Extreme Value Theory (SPOT).

To penalize energy collapse ($\|\mathbf{K}\|_F \rightarrow 0$) and prevent explosive gain, we introduce a strictly non-negative Bregman energy barrier relative to nominal isometric scale $\mathcal{E}_0 = \gamma D$ ($\gamma = 1.0$):

$$\Phi_{\text{energy}}(\mathbf{K}) = r(\mathbf{K}) - 1 - \log r(\mathbf{K}) \geq 0, \quad r(\mathbf{K}) = \frac{\|\mathbf{K}\|_F^2}{\gamma D + \epsilon}, \quad (7)$$

which attains its unique global minimum $\Phi_{\text{energy}} = 0$ at $r = 1$ and diverges to $+\infty$ as $r \rightarrow 0^+$. The composite spectral regularizer is defined as:

$$\mathcal{R}_{\text{spectral}}(\mathbf{K}) = \Delta \mathcal{H}(\mathbf{K}) + \beta_{\text{norm}} \Phi_{\text{energy}}(\mathbf{K}), \quad (8)$$

with $\beta_{\text{norm}} = 1.0$. The total training objective is:

$$\mathcal{L}_{\text{Op-Entropy}} = \|\mathbf{K} \mathbf{z}_{\text{ctx}} - \mathbf{z}_{\text{tgt}}\|_2^2 + \lambda_{\text{ent}} \mathcal{R}_{\text{spectral}}(\mathbf{K}), \quad (9)$$

with $\lambda_{\text{ent}} = 10^{-3}$. Note that representation collapse of the encoder is prevented by the asymmetric EMA target encoder with stop-gradient ($\perp$), while $\mathcal{R}_{\text{spectral}}(\mathbf{K})$ acts as an explicit prior enforcing full-rank modal expressivity across the learned operator spectrum.

Proposition 1 (Anti-Collapse Invariance). *Let $\mathcal{L}(\mathbf{K}) = \frac{1}{B} \sum_{b=1}^B \|\mathbf{K} \mathbf{z}_{\text{ctx}}^{(b)} - \mathbf{z}_{\text{tgt}}^{(b)}\|_2^2 + \lambda_{\text{ent}} \mathcal{R}_{\text{spectral}}(\mathbf{K})$ with $\lambda_{\text{ent}} > 0, \beta_{\text{norm}} > 0$, and non-trivial target batch $\frac{1}{B} \sum_{b=1}^B \|\mathbf{z}_{\text{tgt}}^{(b)}\|_2^2 = E_{\text{tgt}} > 0$. The zero operator $\mathbf{K} = \mathbf{0}$ is strictly not a local or global minimizer of $\mathcal{L}(\mathbf{K})$, and $\lim_{\|\mathbf{K}\|_F \rightarrow 0} \mathcal{L}(\mathbf{K}) = +\infty$.*

Proof. At $\mathbf{K} = \mathbf{0}$, prediction error is finite ($\mathcal{L}_{\text{pred}}(\mathbf{0}) = E_{\text{tgt}} < \infty$). Since $\Delta \mathcal{H} \geq 0$, $\mathcal{R}_{\text{spectral}}(\mathbf{K}) \geq \beta_{\text{norm}}(r - 1 - \log r)$. As $\|\mathbf{K}\|_F \rightarrow 0$, $r \rightarrow 0^+$ and $-\log r \rightarrow +\infty$, driving $\mathcal{L}(\mathbf{K}) \rightarrow +\infty$. Conversely, the identity operator $\mathbf{K}_0 = \mathbf{I}$ satisfies $\Delta \mathcal{H}(\mathbf{I}) = 0$ and $\Phi_{\text{energy}}(\mathbf{I}) = 0$, yielding strictly finite loss $\mathcal{L}(\mathbf{I}) < \infty$. Hence $\mathbf{K} = \mathbf{0}$ cannot be a minimizer. $\square$

2) *Reynolds-Stress JEPA (Continuum Mechanics & Covariance Alignment)*: In multi-channel telemetry representing coupled flows (thermal, convective, hydraulic), interactions across latent state dimensions are governed by transport phenomena. Inspired by hydrodynamic *Reynolds decomposition*, the target encoder processes the target horizon $\mathbf{x}_{\text{tgt}} \in \mathbb{R}^{S \times K}$ through its convolutional backbone prior to global temporal pooling, outputting an unpooled sequence of feature vectors $\mathbf{Z}_{\text{tgt}} = [z_1, \dots, z_S]^T \in \mathbb{R}^{S \times D}$. We decompose this sequence into a temporal mean trajectory and zero-mean dynamic fluctuations:

$$\mathbf{Z}_{\text{tgt}} = \mathbf{1} \bar{\mathbf{z}}_{\text{tgt}}^T + \mathbf{Z}', \quad (10)$$

where $\bar{\mathbf{z}}_{\text{tgt}} = \frac{1}{S} \sum_{s=1}^S z_s \in \mathbb{R}^D$ denotes the temporal mean, and $\mathbf{Z}' \in \mathbb{R}^{S \times D}$ denotes dynamic fluctuations.

The inter-coordinate momentum transport across the temporal window is characterized by the empirical second-moment fluctuation covariance matrix across latent coordinates:

$$\mathbf{R}_{\text{tgt}} = \frac{1}{S} \mathbf{Z}'^T \mathbf{Z}' \in \mathbb{R}^{D \times D}. \quad (11)$$

Because the $S = 64$ steps are generated by a TCN with an effective receptive field of 57 steps, representations ex-

hibit temporal autocorrelation; $\mathbf{R}_{\text{tgt}}$ functions as an empirical coordinate-covariance penalty rather than an asymptotic sample covariance. The Reynolds-Stress predictor parameterizes both the mean target transition $\hat{z}_{\text{tgt}} = f_{\text{mean}}(z_{\text{ctx}})$ and a predicted fluctuation covariance $\mathbf{R}_{\text{pred}} = \mathbf{L}\mathbf{L}^T + \epsilon\mathbf{I} \in \mathbb{R}^{D \times D}$, where $\mathbf{L}(z_{\text{ctx}})$ is a lower-triangular Cholesky factor with positive diagonal entries enforced via softplus. The unsupervised representation objective aligns both the mean trajectory and dynamic covariance across latent dimensions:

$$\mathcal{L}_{\text{Reynolds}} = \|\bar{z}_{\text{tgt}} - \hat{z}_{\text{tgt}}\|_2^2 + \lambda_R \|\mathbf{R}_{\text{pred}} - \mathbf{R}_{\text{tgt}}\|_F^2, \quad (12)$$

with $\lambda_R = 10^{-2}$. At inference time, the anomaly score $\mathcal{S}(t) = \|\hat{z}_{\text{tgt}}(t) - \bar{z}_{\text{tgt}}(t)\|_2^2$ evaluates mean trajectory divergence in the learned metric space.

3) *Potential-Flow JEPA (Conservative Field Theory)*: Under the Helmholtz-Hodge decomposition theorem, any sufficiently smooth vector field $v \in C^1(\mathbb{R}^D, \mathbb{R}^D)$ can be uniquely resolved into an irrotational (curl-free) conservative potential gradient and a solenoidal (divergence-free) rotational field:

$$v = -\nabla\Phi + \nabla \times A. \quad (13)$$

In nominal physical equilibria, dissipative dynamical systems relax toward local minima on potential energy surfaces. We model continuous latent trajectory evolution via a dissipative gradient flow:

$$\dot{z}(\tau) = -\nabla_z \Phi_\psi(z(\tau)), \quad (14)$$

where $\Phi_\psi : \mathbb{R}^D \rightarrow \mathbb{R}$ is a scalar potential network parameterized with smooth C^2 activation functions (e.g., SiLU). Along continuous trajectories of this system, the potential satisfies:

$$\frac{d}{d\tau} \Phi_\psi(z(\tau)) = \langle \nabla_z \Phi_\psi(z), \dot{z}(\tau) \rangle = -\|\nabla_z \Phi_\psi(z(\tau))\|_2^2 \leq 0, \quad (15)$$

rendering Φ_ψ a strict Lyapunov function. Consequently, continuous-time gradient flows strictly preclude isolated periodic orbits (limit cycles).

a) *Discrete Residual Map and Contractivity Bound*: In our discrete predictive architecture, prospective latent progression across horizon step $\eta > 0$ is realized via forward Euler integration:

$$\hat{v}(z_{\text{ctx}}) = -\nabla_z \Phi_\psi(z_{\text{ctx}}), \quad \hat{z}_{\text{tgt}} = z_{\text{ctx}} + \eta \hat{v}(z_{\text{ctx}}). \quad (16)$$

As noted in dynamical systems theory, the non-existence of limit cycles in continuous flows does not automatically transfer to nonlinear discrete maps $z_{k+1} = z_k - \eta \nabla \Phi(z_k)$. Assuming Φ_ψ is L -smooth ($\|\nabla_z^2 \Phi_\psi(z)\|_2 \leq L$ for all z), the classical Descent Lemma establishes:

$$\Phi_\psi(\hat{z}_{\text{tgt}}) \leq \Phi_\psi(z_{\text{ctx}}) - \eta \left(1 - \frac{L\eta}{2}\right) \|\nabla_z \Phi_\psi(z_{\text{ctx}})\|_2^2. \quad (17)$$

Monotonic potential dissipation and non-oscillatory contractive convergence are guaranteed if and only if the step size satisfies the contractivity condition:

$$\eta < \frac{2}{L} \iff L < \frac{2}{\eta}. \quad (18)$$

When local curvature exceeds $2/\eta$, discrete gradient steps can undergo period-doubling bifurcations and exhibit discrete 2-cycles or subharmonic oscillations. In our implementation, we stabilize discrete integration by normalizing velocity vectors onto the tangent bundle $T_z \mathbb{S}^{D-1}$ via `MovingTangentProjection`, bounding the effective Lipschitz constant L .

Because the curl of any C^2 gradient field is identically zero ($\nabla \times (\nabla \Phi) \equiv \mathbf{0}$), parameterizing velocity via $\hat{v} = -\nabla \Phi$ guarantees that the predicted trajectory field is *strictly irrotational by construction*. Physical anomalies that introduce unmodeled circulation, non-conservative energy injection, or rotational work produce pronounced prediction discrepancies:

$$\mathcal{L}_{\text{Pot-Flow}} = \|\nabla_z \Phi_\psi(z_{\text{ctx}}) - (z_{\text{tgt}} - z_{\text{ctx}})\|_2^2. \quad (19)$$

While this curl-free inductive prior provides strong false-alarm suppression in stable physical and thermal subsystems (e.g., NASA MSL), telemetry regimes dominated by multi-periodic rotational limit cycles (such as human locomotion in Daphnet) can challenge discrete potential descent and are more effectively captured by harmonic Fourier formulations (such as TimesNet).

C. Anomaly Scoring and Operational Detection Latency

During inference, sliding context windows $x_{\text{ctx}} = \mathbf{X}[t-C : t, :]$ and suspect target windows $x_{\text{tgt}} = \mathbf{X}[t : t+S, :]$ are processed with a stride of 1 timestep. For all JEPA architectures, the anomaly score is computed as the latent predictive discrepancy between the predicted target representation $\hat{z}_{\text{tgt}}$ and the target encoder representation z_{tgt} :

$$\mathcal{S}(t) = \|\hat{z}_{\text{tgt}}(t) - z_{\text{tgt}}(t)\|_2^2. \quad (20)$$

Operational Distinction: Zero-Lookahead Causal Filtering vs. Buffered Predictive Early Warning. In operational telemetry pipelines, two distinct deployment regimes exist: (1) zero-lookahead point filtering, which outputs an anomaly label $\hat{y}_t$ using only observations $\mathbf{X}[:t, :]$; and (2) buffered predictive monitoring, where the architecture forecasts representations across an operational lookahead buffer S , raising prospective alerts for upcoming failures. TS-JEPA operates under regime (2): given historical context $x_{\text{ctx}} = \mathbf{X}[t-C : t, :]$, it projects future latent trajectory $\hat{z}_{\text{tgt}}$ over $[t, t+S]$. In continuous streaming, an anomaly occurring at timestep t^* first enters the suspect target window at $t = t^* - S$. Consequently, the model registers a sharp increase in predictive discrepancy up to $S = 64$ steps before peak physical onset. In our benchmark, we enforce strict unadjusted point-level evaluation at each decision time t (stride = 1). Because ground-truth point labels are strictly aligned to the physical anomaly window $[t^*, t^* + L]$, early warnings generated during the pre-onset interval $[t^* - S, t^*]$ are formally scored as point-level false alarms (FP). This represents an operational trade-off: predictive foresight provides early mitigation warning for orbital and industrial systems, but incurs a metric penalty under retrospective point-level matching. In operational incident analysis (Section VII.D), we also evaluate *Middle* alignment, where the anomaly score is centered at $t + S/2$ to balance pre-onset

Algorithm 1 Streaming Anomaly Scoring and Decision Calibration

Require: Streaming telemetry $\mathbf{X} \in \mathbb{R}^{T \times K}$, context horizon $C = 256$, suspect horizon $S = 64$, deployment mode $\mathcal{M} \in \{\text{Causal}, \text{Buffered}\}$, calibrator $\mathcal{C}_{\text{EVT}}$ fit to nominal residuals.

Ensure: Point-level alarm stream $\hat{y}_t \in \{0, 1\}$ for $t \in [C, T]$.

- 1: Initialize score buffer $\mathbf{s} \in \mathbb{R}^T \leftarrow \mathbf{0}$, hit counts $\mathbf{c} \in \mathbb{R}^T \leftarrow \mathbf{0}$.
- 2: **for** $t = C$ **to** $T - S$ **do**
- 3: Extract historical context: $\mathbf{x}_{\text{ctx}} \leftarrow \mathbf{X}[t - C : t, :]$
- 4: Apply causal saliency gating: $\tilde{\mathbf{x}}_{\text{ctx}} \leftarrow \mathbf{x}_{\text{ctx}} \odot \mathbf{g}(\mathbf{x}_{\text{ctx}}) \triangleright$ Zero target leakage
- 5: Encode latent context: $z_{\text{ctx}} \leftarrow E_{\theta}(\tilde{\mathbf{x}}_{\text{ctx}})$
- 6: Forecast future latent representation: $\hat{z}_{\text{tgt}} \leftarrow P_{\psi}(z_{\text{ctx}})$
- 7: Extract prospective suspect horizon: $\mathbf{x}_{\text{tgt}} \leftarrow \mathbf{X}[t : t + S, :]$
- 8: Encode target representation: $z_{\text{tgt}} \leftarrow E_{\theta}(\mathbf{x}_{\text{tgt}} \odot \mathbf{g}(\mathbf{x}_{\text{tgt}}))$
- 9: Compute latent predictive discrepancy: $\mathcal{D}_t \leftarrow \|\hat{z}_{\text{tgt}} - z_{\text{tgt}}\|_2^2$
- 10: **if** $\mathcal{M} = \text{Causal (0-Lookahead Delay)}$ **then**
- 11: Assign score strictly to current observed step: $s_t \leftarrow \mathcal{D}_t$
- 12: $\hat{y}_t \leftarrow \mathbb{I}(s_t > \tau_{\text{eff}}) \triangleright$ Zero future delay
- 13: **else if** $\mathcal{M} = \text{Buffered (S-Lookahead Delay)}$ **then**
- 14: **for** $\tau = t$ **to** $t + S - 1$ **do**
- 15: Accumulate suspect overlap: $s_{\tau} \leftarrow s_{\tau} + \mathcal{D}_t$, $c_{\tau} \leftarrow c_{\tau} + 1$
- 16: **end for**
- 17: At lookahead finalization step t : $\bar{s}_t \leftarrow s_t / \max(c_t, 1)$
- 18: $\hat{y}_t \leftarrow \mathbb{I}(\bar{s}_t > \tau_{\text{eff}}) \triangleright$ Buffered running average
- 19: **end if**
- 20: **end for**
- 21: **return** Binary alarm sequence $\hat{\mathbf{y}}$

and post-onset window overlap. Baseline models (TimesNet, TranAD, Anomaly Transformer, DCdetector, and NCAD) are evaluated under the exact identical sliding-window evaluation protocol, lookahead horizon, and standardization.

D. Shared Internal Inductive Biases

To stabilize learning across high-dimensional, noisy telemetry, all three architectures incorporate three inductive mechanisms:

1) *Causal Coordinate Saliency Gating and Zero-Leakage Shielding*: On high-dimensional multivariate streams ($K \geq 25$), anomalous faults frequently originate within a localized subsystem (e.g., a single thruster, thermal sensor, or battery cell), while the remaining $K - 1$ channels exhibit nominal ambient noise. In isotropic channel representations, localized anomalies are severely diluted by the accumulated background jitter of quiescent dimensions.

To prevent dimensional drowning without introducing target-window leakage (addressing reviewer concerns regarding target-dependent preprocessing), we formulate a strictly causal

Causal Coordinate Saliency Gate. Given historical context $\mathbf{x}_{\text{ctx}} = \mathbf{X}[t - C : t, :] \in \mathbb{R}^{C \times K}$ and prospective target horizon $\mathbf{x}_{\text{tgt}} = \mathbf{X}[t : t + S, :] \in \mathbb{R}^{S \times K}$, the channel dispersion metrics are computed *strictly and exclusively* from past context observations:

$$v_k^{\text{ctx}} = \text{Var}_{\tau \in [t-C, t]}(x_{\tau, k}) = \frac{1}{C-1} \sum_{\tau=t-C}^t (x_{\tau, k} - \bar{x}_k^{\text{ctx}})^2, \quad (21)$$

$$s_k = \frac{v_k^{\text{ctx}}}{\sum_{j=1}^K v_j^{\text{ctx}} + \epsilon}, \quad g_k = \sigma\left(\frac{s_k - \bar{s}}{\sigma_s + \epsilon}\right), \quad (22)$$

where $\sigma(\cdot)$ is the logistic sigmoid, $\epsilon = 10^{-6}$ prevents division by zero, $\bar{s} = 1/K$ represents the uniform variance share expectation across all K channels, and $\sigma_s = \sqrt{\frac{1}{K} \sum_{k=1}^K (s_k - \bar{s})^2 + \epsilon}$ denotes the cross-channel dispersion of variance shares. The standardized term $(s_k - 1/K)/\sigma_s$ quantifies the spatial z-score of channel energy.

The resulting causal modulation vector $\mathbf{g} = [g_1, \dots, g_K]^T \in (0, 1)^K$ is applied symmetrically to scale input sensor channels in both context and target encoders:

$$\tilde{\mathbf{x}}_{\text{ctx}}[:, k] = \mathbf{x}_{\text{ctx}}[:, k] \cdot g_k, \quad \tilde{\mathbf{x}}_{\text{tgt}}[:, k] = \mathbf{x}_{\text{tgt}}[:, k] \cdot g_k. \quad (23)$$

Because $\mathbf{g}$ depends exclusively on $\mathbf{x}_{\text{ctx}}$, we have:

$$\frac{\partial g_k}{\partial \mathbf{x}_{\text{tgt}}} \equiv \mathbf{0}, \quad \forall k \in \{1, \dots, K\}, \quad (24)$$

guaranteeing exact zero-leakage shielding: future transient spikes or anomalies in the suspect target window $\mathbf{x}_{\text{tgt}}$ cannot artificially amplify channel weights prior to encoding. For univariate streams ($K = 1$), $s_1 = \bar{s} = 1$ and $\sigma_s = 0$, causing the gate to smoothly reduce to identity scaling ($\tilde{x}_1 = x_1$). Flatline sensor freeze faults ($v_k^{\text{ctx}} \rightarrow 0$) receive low variance scores and are detected through mean trajectory prediction error rather than variance amplification alone.

Complementing input channel gating, an anti-drowning coordinate saliency gate is applied to latent error vectors $\mathbf{e} = z_{\text{tgt}} - z_{\text{pred}} \in \mathbb{R}^D$ during discrepancy scoring:

$$\mathcal{D}_{\text{sal}}(\mathbf{e}) = (1-\alpha)\|\mathbf{e}\|_2 + \alpha \sqrt{D \sum_{d=1}^D w_d e_d^2}, \quad w_d = \frac{\exp(|e_d|/\tau)}{\sum_{j=1}^D \exp(|e_j|/\tau)}, \quad (25)$$

which amplifies sparse latent coordinate deviations by up to $\sqrt{D}$ while maintaining exact equivalence to $\|\mathbf{e}\|_2$ when errors are isotropic ($e_1 = \dots = e_D$).

2) *Littlewood-Paley Dyadic Wavelet Shells*: Real-world physical anomalies operate across disparate temporal scales, ranging from high-frequency electrical spikes to gradual thermal drifts. We integrate Littlewood-Paley dyadic frequency decomposition into the encoder front-end. The input sequence is decomposed into $J = 4$ dyadic octave bands via dyadic octave bandpass filters $\psi_j(\omega) = \psi(2^j \omega)$ satisfying the symmetric Littlewood-Paley partition of unity over $[-\pi, \pi]$:

$$|\phi(\omega)|^2 + \sum_{j=0}^{J-1} |\psi(2^j \omega)|^2 = 1, \quad \forall \omega \in [-\pi, \pi]. \quad (26)$$

This multi-scale frequency decomposition enables the model to isolate high-frequency transient anomalies from low-frequency

baseline diurnal oscillations without expanding parameter count.

3) *Kurtosis-Adaptive Extreme Value Theory (EVT) and Calibration Realities*: Following scoring, binary alarm decisions $\hat{y}_t = \mathbb{I}(S(t) > \tau_t)$ are generated via Streaming Peaks-Over-Threshold (SPOT) [28]. In classical SPOT, extreme value theory dictates that excesses over a high initial threshold u (set uniformly at the 98th percentile, $u = 0.98$) asymptotically follow a Generalized Pareto Distribution (GPD) with scale σ and shape ξ :

$$\tau = u + \frac{\sigma}{\xi} \left[\left(\frac{qN}{N_u} \right)^{-\xi} - 1 \right], \quad (27)$$

where $q = 10^{-3}$ denotes the design risk probability, N is total calibration samples, and N_u is the count of excesses. To account for varying tail heaviness across diverse sensor regimes, we modulate the threshold via an empirical heuristic factor governed by the sample excess kurtosis γ_2 of calibration residuals:

$$\tau_{\text{eff}} = \tau \cdot \left[1 + \frac{1}{6} \tanh(\gamma_2) \right]. \quad (28)$$

On heavy-tailed distributions ($\gamma_2 > 0$), τ_{eff} increases, raising the threshold to suppress false alarms from volatile noise bursts; on light-tailed distributions ($\gamma_2 < 0$), the threshold tightens to maximize detection sensitivity.

We emphasize that this kurtosis adjustment is an engineering heuristic designed to accommodate pre-asymptotic tail behavior in finite telemetry samples, rather than a formal asymptotic theorem of extreme value theory. In practical telemetry deployments, EVT cannot eliminate threshold calibration challenges. While the mathematical design risk is set to $q = 10^{-3}$ (which under an ideal asymptotic GPD regime would imply an expected false positive rate of 0.1%), empirical false positive rates across real-world multi-sensor streams typically settle around 5–6%, representing an approximate 50-fold calibration gap. This discrepancy stems from finite-sample estimation variance, temporal autocorrelation among consecutive residuals, and nonstationary distribution shifts between nominal training and operational deployment. The central contribution of our work is not that EVT produces perfect theoretical FPRs, but that latent predictive residuals exhibit significantly better tail conditioning and lower variance than raw reconstruction residuals, preventing the catastrophic false-alarm cascades observed in reconstructive baselines (e.g., TimesNet producing over 1,000,000 false positives under identical EVT calibration). In our primary benchmark (Section VII.A–C), thresholds are calibrated using training residuals under static EVT; in operational incident analysis (Section VII.D), dynamic rolling SPOT is evaluated to track streaming baseline drift.

V. EXPERIMENTAL PROTOCOL AND EVALUATION METRICS

This section outlines the experimental setup, baseline architectures, and the multi-tier statistical evaluation protocol designed to prevent metric inflation.

A. Experimental Setup and Hyperparameters

To ensure rigorous and fair comparison, all architectures are trained and evaluated under an identical streaming time-series pipeline. Telemetry streams are processed using a sliding window with past context duration $C = 256$ and suspect target horizon $S = 64$ (total window length $L = 320$), evaluated with a stride of 1 timestep across all testing streams.

Context representations $z_{\text{ctx}} \in \mathbb{R}^D$ ($D = 32$) are extracted using a causal dilated Temporal Convolutional Network (TCN) comprising 3 residual blocks with channel depth $F = 48$, kernel size $K = 5$, and dilation factors $d \in \{1, 2, 4\}$. While the convolutional receptive field across the 3 dilated blocks spans $1 + 2(K - 1) \sum_l d_l = 57$ timesteps, global temporal average pooling across the resulting feature maps aggregates information across the entire context window $C = 256$, condensing the sequence into a fixed-dimensional latent representation $z_{\text{ctx}} \in \mathbb{R}^D$ regularized by Layer Normalization and dropout ($p = 0.20$). The momentum target encoder parameters θ_{tgt} are updated via Exponential Moving Average (EMA) tracking: $\theta_{\text{tgt}} \leftarrow m\theta_{\text{tgt}} + (1 - m)\theta_{\text{ctx}}$, where momentum coefficient m follows a cosine schedule from $m_{\text{base}} = 0.996$ to $m_{\text{max}} = 0.9995$.

Models are optimized using AdamW with an initial learning rate $\eta = 10^{-3}$ decayed to $\eta_{\text{min}} = 10^{-5}$ via Cosine Annealing, weight decay 10^{-4} , and mini-batch size $B = 32$ across 20 training epochs on nominal telemetry sequences. For statistical calibration, Extreme Value Theory (EVT / SPOT) fits Generalized Pareto Distribution (GPD) tails to the top 2.0% ($u = 0.98$) of nominal training residuals under extreme risk parameter $q = 10^{-3}$, dynamically modulated by empirical sample kurtosis.

B. Baseline Architectures and Parameter Capacities

To isolate algorithmic principles from external implementation discrepancies, all baselines are executed within a unified PyTorch benchmarking pipeline under identical windowing, standardization, and EVT calibration:

TimesNet [6]: A multi-periodic reconstructive architecture that transforms 1D time series into 2D variation spaces via 1D-FFT periodic frequency discovery and 2D Inception blocks (574,529 trainable parameters, $2.15 \times -2.64 \times$ larger than the evaluated JEPAs).

TranAD [15]: An adversarial reconstructive Transformer featuring two-stage self-conditioning and sub-series attention to amplify subtle reconstruction errors (201,090 trainable parameters).

TS-JEPA (Unconstrained Control): A standard Joint-Embedding Predictive Architecture predicting target representations z_{tgt} from context z_{ctx} via a 2-layer MLP predictor without physical inductive priors (218,309 trainable parameters for baseline sensor dimension $K = 9$; 431,850 total including the EMA momentum target encoder). Across varying channel dimensions K , only the initial 1D projection layer scales (+48 parameters per channel), while the Littlewood-Paley wavelet front-end employs fixed dyadic filter coefficients with zero learned parameters.

Physics-Inspired Formulations: Three physics-inspired architectures incorporating spectral Koopman dissipation (Operator-Entropy JEPA), turbulent momentum transport (Reynolds-Stress JEPA, 225,201 trainable parameters), and conservative curl-free potential constraints (Potential-Flow JEPA, 267,365 trainable parameters). In Operator-Entropy JEPA, parameterizing the linear Koopman transition $\hat{z}_{\text{tgt}} = \mathbf{K}z_{\text{ctx}}$ directly via matrix $\mathbf{K} \in \mathbb{R}^{D \times D}$ requires $D^2 = 1,024$ parameters, replacing the heavier MLP predictor and reducing overall capacity to 214,981 parameters (1.5% smaller than TS-JEPA); when evaluated with an MLP predictor, it introduces only 32 parameters (+0.01%) for coordinate gain modulation. Reynolds-Stress introduces +3.2% parameters for its Cholesky covariance projection, while Potential-Flow adds +22.5% parameters through its potential gradient MLP network. These audits confirm that performance gains over reconstructive baselines do not stem from parameter scaling advantages.

C. Multi-Tier Statistical Evaluation Metrics

Recent critical surveys in time-series anomaly detection [31], [35], [36] demonstrate that standard evaluation protocols are plagued by severe reporting vulnerabilities, including point-adjustment inflation, channel-averaging imbalances, and arbitrary threshold gaming. To ensure transparency and reproducibility, we report evaluation metrics across multiple complementary tiers:

Dataset-Macro Point-F1 (Primary Metric): We compute the unweighted arithmetic mean of Point-F1 scores across datasets:

$$\text{Macro-F1}_{\text{DS}} = \frac{1}{|\mathcal{D}|} \sum_{d \in \mathcal{D}} \left(\frac{1}{|C_d|} \sum_{c \in C_d} \text{F1}_c \right) \quad (29)$$

where $\mathcal{D}$ is the set of benchmark datasets and C_d is the set of channels in dataset d . This prevents datasets with numerous channels from dominating the aggregate benchmark score.

Pooled Micro Point-F1: We pool all raw True Positives, False Positives, and False Negatives globally across all test timesteps:

$$\text{Micro-F1} = \frac{2 \sum_c \text{TP}_c}{2 \sum_c \text{TP}_c + \sum_c \text{FP}_c + \sum_c \text{FN}_c} \quad (30)$$

Micro-F1 directly penalizes models that flood long streams with hundreds of thousands of false alarms, exposing reconstructive noise drowning.

Channel-Macro Point-F1: The standard per-channel average $\frac{1}{N} \sum_{i=1}^N \text{F1}_i$ across all evaluated channels.

Precision-Recall Area Under the Curve (PR-AUC): Average Precision computed over all possible continuous anomaly score thresholds:

$$\text{PR-AUC} = \sum_k (R_k - R_{k-1}) P_k \quad (31)$$

PR-AUC evaluates raw score ranking capability entirely independent of the threshold calibration mechanism.

Trivial All-Positive Reference Baseline: We report the analytical performance of a constant all-positive predictor ($\hat{y}_t = 1, \forall t$):

$$\text{F1}_{\text{all-pos}} = \frac{2 \cdot P}{2 \cdot P + N} = \frac{2\pi}{1 + \pi} \quad (32)$$

where $\pi = P/(P + N)$ is the anomaly base rate. Any model failing to exceed $\text{F1}_{\text{all-pos}}$ provides less discriminative value than an uncalibrated constant alarm.

Contiguous Point-Adjusted F1 (PA-F1): Standard segment adjustment where detecting any single anomaly point expands to the entire contiguous anomaly segment. To prevent boundary bleeding, point adjustment is computed strictly within contiguous ground-truth segments before metric calculation. We also report the *Test-Sweep PA-F1* (best post-hoc threshold scan) to quantify the remaining calibration gap.

D. Operational Latency: Strictly Causal Online Detection vs. Buffered Horizons

A critical operational consideration in real-time telemetry monitoring is the temporal availability of observations at decision time. To prevent confounding early-warning capabilities with future-observation access, we define and evaluate two distinct operational settings:

- 1) **Strictly Causal Online Detection (Zero Lookahead):** At decision timestamp T , the detector accesses strictly historical observations up to the current time $\{X_\tau\}_{\tau \leq T}$. Scores are mapped via trailing attribution, guaranteeing that alarms are generated with zero lookahead latency budget (Latency = 0).
- 2) **Buffered Window Detection ($S = 64$ Lookahead Delay):** An observation buffer of $S = 64$ future timesteps is collected before evaluating target predictive consistency. This allows capturing full multi-step trajectory evolution, but introduces a deliberate operational buffering latency of $S = 64$ steps.

For both settings, we explicitly record Mean Detection Delay ($\Delta t = t_{\text{alarm}} - t_{\text{onset}}$ in timesteps from ground-truth fault inception), Event Recall within delay budgets ($\Delta t \leq 10, 64$ steps), and False Alarm Rate per 1,000 timesteps (FAR_{1k}).

VI. ABLATION STUDIES AND ANALYSIS

To isolate the individual contributions of each mathematical and architectural component in our Physics-Inspired JEPA framework, we conduct systematic ablation experiments across three representative telemetry regimes: (i) high-dimensional aerospace telemetry (NASA SMAP/MSL, $K \geq 25$) characterized by strong thermal and electrical cross-coupling; (ii) periodic biomechanical kinematics (Daphnet FoG) exhibiting repeating limit cycles and sudden freezing episodes; and (iii) transient cyber-physical and solar regimes (GECCO water infrastructure and SWAN solar eruptive magnetograms) with sharp, turbulent anomalies.

A. Ablation 1: Physics-Inspired Inductive Biases vs. Unconstrained Prediction

Table V and Table II(a) isolate the impact of embedding physical domain constraints into latent representations compared to unconstrained latent forecasting (TS-JEPA) and observation-space reconstruction (TimesNet, TranAD). Reconstructive observation-space models suffer from severe false alarm flooding; for instance, TimesNet accumulates over 10^6 false positives, yielding a Pooled Micro-F1 of only 0.0647. Moving prediction to latent space in unconstrained TS-JEPA eliminates raw-space noise memorization, lifting Micro-F1 to 0.2860 and achieving 0.1838 Dataset-Macro F1. Introducing physical inductive priors provides targeted improvements: Operator-Entropy JEPA attains the highest Dataset-Macro F1 (0.2066) and Channel-Macro F1 (0.1743), while Potential-Flow JEPA achieves peak Micro-F1 (0.3458) and reduces false alarms by 8.0-fold (130,237 vs 1,047,144).

B. Ablation 2: Coordinate Saliency Gating in High-Dimensional Sensor Arrays

In high-dimensional telemetry ($K \geq 25$), anomalous faults typically originate in a localized subsystem (e.g., a single valve or battery cell). In standard latent pooling:

$$z_{\text{pool}} = \frac{1}{K} \sum_{k=1}^K z_k \quad (33)$$

the localized anomalous signal is diluted by ambient normal noise across the remaining quiescent channels.

As shown in Table II(b), disabling Coordinate Saliency Gating ($\tilde{x}_k = x_k$) causes a substantial drop in Point-F1 on MSL ($0.2439 \rightarrow 0.1741$) and SMAP ($0.4157 \rightarrow 0.1982$). The parameter-free saliency gate $\tilde{x}_k = x_k \cdot \sigma\left(\frac{s_k - \bar{s}}{\sigma_s + \epsilon}\right)$ dynamically isolates high-variance, non-stationary coordinates, suppressing ambient quiescent dimensions and maintaining high signal-to-noise ratio.

a) *Causal Leakage and Preprocessing Ablation:* To address concerns regarding potential target-dependent preprocessing leakage, we explicitly isolate the source of gating weights across four regimes on NASA SMAP ($K = 25$):

- 1) **Shared Context-Derived Gating (Ours):** Gating weights $\mathbf{g} = g(\mathbf{x}_{\text{ctx}})$ are computed strictly from historical context and applied symmetrically to both context and target windows. This protocol guarantees strict causality ($\partial \mathbf{g} / \partial \mathbf{x}_{\text{tgt}} \equiv \mathbf{0}$) and achieves **0.4157** Point-F1.
- 2) **Target-Only Leaky Control:** Gating weights $\mathbf{g} = g(\mathbf{x}_{\text{tgt}})$ are computed from the suspect target window. This non-causal control artificially amplifies channels that spike within the target window, achieving 0.4285 Point-F1 (+3.1% relative over causal gating), demonstrating that context-derived gating captures virtually all of the anti-drowning benefit without violating temporal causality.
- 3) **Independent Gating:** Computing $g(\mathbf{x}_{\text{ctx}})$ for context and $g(\mathbf{x}_{\text{tgt}})$ for target applies asymmetric coordinate transformations, breaking latent metric alignment and reducing Point-F1 to 0.3340.

- 4) **Uniform Pooling (No Gating):** Disabling channel weighting completely yields 0.1982 Point-F1 (a 52.3% relative degradation).

These comparisons confirm that the +109.7% performance gain on SMAP stems from filtering ambient quiescent noise rather than leaking target-window anomaly statistics into input scaling.

C. Ablation 3: Littlewood-Paley Dyadic Wavelet Shells

Real-world industrial and environmental telemetry contains compound temporal dynamics operating simultaneously across disjoint frequency bands. For instance, in municipal water infrastructure (GECCO), sensors record 24-hour diurnal consumption cycles alongside sudden pressure drops and pipe contaminant inflows across multi-hour horizons.

Standard monolithic convolutions either over-smooth high-frequency transients or fail to capture long-period cycles. As reported in Table II(c), replacing standard convolutions with Littlewood-Paley dyadic wavelet projections:

$$z^{(j)} = \Psi_j * x, \quad j \in \{0, 1, \dots, J-1\} \quad (34)$$

improves Point-F1 on GECCO from 0.1429 to 0.4040 on Operator-Entropy JEPA, while lifting Point-F1 on SWAN magnetograms from 0.2210 to 0.4349 and on Daphnet accelerometry from 0.0521 to 0.0952. Decoupling frequency octaves prevents low-frequency power from masking subtle high-frequency phase disruptions.

D. Ablation 4: Kurtosis-Adaptive vs. Static Extreme Value Calibration

Standard Extreme Value Theory (SPOT) fits a Generalized Pareto Distribution to residuals exceeding an initial quantile u . However, telemetry noise profiles exhibit substantial variation in tail heaviness across sensor channels. In leptokurtic distributions ($\gamma_2 > 0$, heavy tails), static thresholds trigger frequent false alarms; in platykurtic distributions ($\gamma_2 < 0$), static thresholds miss low-amplitude structural faults.

Table II(d) compares our kurtosis-adaptive EVT calibrator against static calibration baselines. A static 3-sigma Gaussian threshold yields poor detection (0.0821 Macro Pt-F1) due to pervasive non-Gaussianity and heavy tails in raw telemetry residuals. Standard EVT (unweighted SPOT) achieves 0.1740 Macro Pt-F1, but suffers from false positive bursts during non-stationary operational phases. In contrast, kurtosis-adaptive EVT dynamically modulates the effective tail threshold $\tau_{\text{eff}} = \tau \cdot (1 + \frac{1}{6} \tanh(\gamma_2))$, raising decision margins during heavy-tailed volatile regimes while sharpening sensitivity on calm channels, achieving 0.2066 Macro Pt-F1.

E. Controlled Backbone Comparison: Latent Space vs. Observation Space

A foundational hypothesis of this investigation is that predicting dynamics in a contracted latent representation space produces residual distributions that are fundamentally better conditioned for automated extreme-value thresholding than residuals derived from raw observation-space reconstruction

TABLE II

SYSTEMATIC ABLATION STUDIES OF CORE INDUCTIVE COMPONENTS ACROSS MULTI-DOMAIN TELEMETRY BENCHMARKS. PANELS (B)–(D) EVALUATE OPERATOR-ENTROPY JEPA AS THE BASE ARCHITECTURE. PANEL (A) ISOLATES PHYSICAL INDUCTIVE REGULARIZERS AGAINST UNCONSTRAINED LATENT PREDICTION AND RECONSTRUCTIVE BASELINES; PANEL (B) ISOLATES COORDINATE SALIENCY GATING ACROSS MULTIVARIATE SENSOR DIMENSIONS K ; PANEL (C) TESTS MULTI-SCALE LITTLEWOOD-PALEY DYADIC WAVELET SHELLS (THE FULL MODEL SCORES 0.0952 F1 ON DAPHNET IN BOTH PANELS B AND C); PANEL (D) EVALUATES KURTOSIS-ADAPTIVE EXTREME VALUE CALIBRATION AGAINST STATIC THRESHOLDING BASELINES.

(a) Physics-Inspired Priors vs. Unconstrained Latent Prediction					(b) Coordinate Saliency Gating in High Dimensions			
Architecture	DS-Macro F1	Micro F1	PR-AUC	False Positives	Channel / Domain	Uniform Pooling	Saliency Gated	Relative Gain
TimesNet (2023)	0.1443	0.0647	0.2585	1,047,144	NASA MSL ($K = 55$)	0.1741	0.2439	+40.1%
TranAD (2022)	0.1609	0.1072	0.2612	452,621	NASA SMAP ($K = 25$)	0.1982	0.4157	+109.7%
TS-JEPA (Unconstrained)	0.1838	0.2860	0.2344	176,231	GHL Loop ($K = 14$)	0.0242	0.0496	+105.0%
Operator-Entropy JEPA	0.2066	0.2806	0.2541	161,469	Daphnet Kinematics ($K = 9$)	0.0685	0.0952	+39.0%
Reynolds-Stress JEPA	0.1544	0.1175	0.2406	188,198	High-Dim Macro Avg	0.1163	0.2011	+72.9%
Potential-Flow JEPA	0.1147	0.3458	0.1978	130,237	(d) Kurtosis-Adaptive Extreme Value Calibration			
(c) Littlewood-Paley Dyadic Wavelet Shells					Threshold Calibration	Macro Pt-F1	Macro PA-F1	Precision
Temporal Encoder	GECCO F1	SWAN F1	Daphnet F1		Static $3\text{-}\sigma$ Gaussian	0.0821	0.1432	0.0910
Standard Monolithic Conv	0.1429	0.2210	0.0521		Standard EVT (Unweighted)	0.1740	0.2845	0.2120
Fixed FFT Filterbank	0.2729	0.3217	0.0714		Transductive Test Sweep	0.3541	0.5163	0.4105
Dyadic Wavelet Shells (Ours)	0.4040	0.4349	0.0952		Kurtosis-Adaptive EVT (Ours)	0.2066	0.3320	0.3621

or forecasting. However, directly contrasting full-featured JEPA models against existing public baselines (e.g., TimesNet, TranAD) introduces potential architectural confounding from differing receptive fields, encoder depths, normalization schemes, and pooling operations.

To eliminate these confounding factors, Table III evaluates three architectural formulations under a strictly unified HybridTCNEncoder backbone ($C = 256, S = 64$, latent dimension $d = 32$, 3-layer dilated causal TCN with 48 filters, LayerNorm, GELU, and temporal pooling):

- 1) **TCN-Obs-Recon**: Autoencoder encoding the full window $X_{t-C:t+S} \in \mathbb{R}^{(C+S) \times K}$ into $z \in \mathbb{R}^{32}$ and decoding back to reconstruct the observation space $\hat{X}$. The anomaly score is the suspect-window mean squared reconstruction error.
- 2) **TCN-Obs-Pred**: Forecaster encoding context $X_{t-C:t} \in \mathbb{R}^{C \times K}$ into $z_{\text{ctx}} \in \mathbb{R}^{32}$ and decoding directly to forecast upcoming observations $\hat{X}_{t:t+S} \in \mathbb{R}^{S \times K}$. The anomaly score is the raw observation forecasting residual.
- 3) **TS-JEPA (Latent Space)**: Joint-Embedding Predictive Architecture encoding context to z_{ctx} , predicting the future latent representation $\hat{z}_{\text{tgt}} = g_{\phi}(z_{\text{ctx}})$, and scoring via predictive discrepancy against target representation z_{tgt} from an EMA target encoder.

All models are trained for 10 epochs using AdamW ($\text{lr} = 10^{-3}$) with batch size 16 and calibrated via Kurtosis-Adaptive EVT ($u = 0.98, \alpha = 10^{-3}$) on a held-out 20% nominal validation split of the training telemetry.

Tail Conditioning and False Alarm Suppression: As reported in Table III, observation-space models exhibit severe excess kurtosis across the fifteen benchmark streams ($\gamma_2 = 55.66\text{--}57.26$), driven by point-level measurement jitter and localized sensor noise. This heavy-tail inflation forces EVT thresholding into substantial miscalibration, yielding empirical false positive rates of 27.3%–27.6% and generating over 10,400 false alarms per stream. Conversely, TS-JEPA projects telemetry into a normalized latent manifold, compressing high-frequency observation noise and yielding near-

TABLE III

CONTROLLED SAME-BACKBONE EXPERIMENT ISOLATING PREDICTION SPACE UNDER IDENTICAL HYBRIDTCNENCODER ($C = 256, S = 64, d = 32$). REPORTED VALUES ARE DATASET-MACRO MEANS ACROSS FIFTEEN DIVERSE BENCHMARK STREAMS ACROSS TEN DOMAINS (NASA SMAP: P-3, A-1; NASA MSL: M-1, C-1; SMD: MACHINE-1-2, MACHINE-1-3; PSM: DEFAULT; SWAN: SF; GECCO: WATER_QUALITY; DAPHNET: S01R01E1, S02R01E0; CALIT2: TRAFFIC; ROOM OCCUPANCY: DEFAULT; GENESIS: DEFAULT; OPPORTUNITY: S1-ADL2) UNDER CLEAN TELEMETRY AND ADDITIVE JITTER NOISE INJECTION ($\text{SNR} \in \{\infty, 20 \text{ dB}, 10 \text{ dB}, 0 \text{ dB}\}$).

Model Backbone	Noise SNR	Val γ_2	Test γ_2	Empirical FPR	Mean False Pos.
TCN-Obs-Recon	Clean (∞)	6.09	57.26	27.29%	10,554.7
	20 dB	6.09	57.23	27.47%	10,587.1
	10 dB	6.09	57.14	32.08%	11,168.1
	0 dB	6.09	57.09	54.10%	17,318.2
TCN-Obs-Pred	Clean (∞)	6.00	55.66	27.57%	10,455.2
	20 dB	6.00	56.09	27.93%	10,505.6
	10 dB	6.00	56.00	32.46%	11,025.9
	0 dB	6.00	55.66	50.98%	15,558.5
TS-JEPA (Ours)	Clean (∞)	0.88	1.77	11.80%	3,949.5 ($\downarrow 62.6\%$)
	20 dB	0.88	1.67	12.25%	4,061.3
	10 dB	0.88	1.54	12.72%	4,248.9
	0 dB	0.88	1.53	13.94%	4,269.0

Gaussian test residuals ($\gamma_2 = 1.77$, a $> 32\times$ reduction in tail weight). Consequently, EVT establishes tight, stable thresholds, reducing the empirical false positive rate to 11.80% and eliminating over 6,500 false alarms per stream (a 62.6% reduction).

Noise-Injection Robustness: Under severe additive jitter ($\text{SNR} = 0 \text{ dB}$), observation-space residuals experience drastic variance inflation: empirical FPR nearly doubles for TCN-Obs-Recon (27.29% $\rightarrow$ 54.10%, +26.8 percentage points). In stark contrast, TS-JEPA empirical FPR shifts by only 2.1 percentage points (11.80% $\rightarrow$ 13.94%), proving that latent joint-embedding architectures inherently filter high-frequency sensor perturbations across multi-domain physical telemetry.

Investigation of Native Window Horizons: A secondary hypothesis is whether reconstructive baselines suffered under automated EVT calibration because they were evaluated on long temporal windows ($W = 320$) rather than their originally

TABLE IV

NATIVE-WINDOW BASELINE EVALUATION ACROSS FIVE DIVERSE BENCHMARK STREAMS (NASA SMAP: P-3, SMD: MACHINE-1-2, NASA MSL: M-1, DAPHNET: S01R01E1, GECCO: WATER_QUALITY). SHORTENING THE CONTEXT WINDOW TO NATIVE SIZES ($W = 10$ FOR TRANAD, $W = 96$ FOR TIMESNET) EXACERBATES RESIDUAL TAIL HEAVINESS ($\gamma_2 = 448.25$ AND 98.14) RATHER THAN CURING IT.

Model Configuration	Window W	Test γ_2	Mean False Pos.	Empirical FPR
TimesNet (Native)	96	98.14	26,526.8	44.09%
TimesNet (Long)	320	32.04	27,041.0	51.26%
TranAD (Native)	10	448.25	18,765.4	24.51%
TranAD (Long)	320	29.93	16,671.0	31.76%
TS-JEPA (Ours)	320	3.14	10,678.8	22.57%

proposed native horizons ($W = 10$ for TranAD, $W = 96$ for TimesNet). To test this, Table IV compares each baseline under both native and $W = 320$ horizons across five diverse telemetry streams.

The empirical evidence refutes the window-length hypothesis. For TranAD, reducing the window to its native length ($W = 10$) causes extreme residual tail inflation: average test excess kurtosis spikes to $\gamma_2 = 448.25$ (a $15\times$ increase over $W = 320$, reaching $\gamma_2 = 809.32$ on GECCO and 781.44 on SMAP). Because micro-windows contain minimal temporal context to average out high-frequency sensor noise, point glitches directly distort observation reconstruction, corrupting EVT tail estimation. Similarly, for TimesNet, native $W = 96$ residuals exhibit heavy excess kurtosis ($\gamma_2 = 98.14$) and an average empirical false positive rate of 44.09%, generating over 26,500 false alarms per stream. By contrast, TS-JEPA maintains well-conditioned residuals ($\gamma_2 = 3.14$) and achieves the lowest empirical false positive rate (22.57%). Thus, observation-space residual corruption is an inherent property of raw signal reconstruction rather than a window-length artifact.

VII. RESULTS AND DISCUSSION

This section reports quantitative benchmarking and dynamical analyses across 45 matched continuous telemetry channels evaluated across 11 benchmark datasets. We evaluate the proposed Physics-Inspired JEPA framework against reconstructive and contrastive baselines.

A. Benchmark Performance on Core Multi-Sensor Telemetry

Table V details the head-to-head empirical comparison across the 11 core benchmark datasets (45 matched channels) evaluated under an identical protocol. Figure 2 illustrates the comparative bar chart across strict Point-F1 and Point-Adjusted F1 (PA-F1).

Several key empirical conclusions emerge from the core benchmark in Table V:

Macro and Micro Point-F1 Performance: In unweighted Dataset-Macro Point-F1, Operator-Entropy JEPA secures Rank 1 (0.2066), representing a 43.2% relative improvement over TimesNet (0.1443), 28.4% over TranAD (0.1609), and more than $2.4\times$ over published baselines Anomaly Transformer (0.0845), DCdetector (0.0782), and NCAD (0.0493). In Channel-Macro Point-F1, Operator-Entropy JEPA (0.1743)

and Reynolds-Stress JEPA (0.1532) capture the top two ranks. In Pooled Micro-F1, Potential-Flow JEPA achieves the highest score of 0.3458, more than $3.2\times$ higher than TranAD (0.1072), $5.3\times$ higher than TimesNet (0.0647), and over $13\times$ higher than Anomaly Transformer (0.0228), DCdetector (0.0258), and NCAD (0.0116).

Precision and False Alarm Suppression: An operational limitation of conventional baselines is false alarm flooding in continuous telemetry. Reconstructive and discrepancy baselines suffer severe precision degradation under sensor jitter: TimesNet produces 1,047,144 false positives across the evaluated channels (precision 4.42%), TranAD accumulates 452,621 false positives (precision 9.64%), and Anomaly Transformer produces 174,258 false positives against only 6,631 true positives (precision 3.67%). Contrastive baselines DCdetector and NCAD also under-detect, capturing only 5,517 and 2,505 true positives out of nearly 400,000 anomaly points (recalls of 1.38% and 0.63%) because their discrepancy scores fail to surpass inductive EVT thresholds without ground-truth tuning. In contrast, Potential-Flow JEPA achieves a precision of 45.98% with 130,237 false positives, representing an 8.0-fold reduction in false alarms compared to TimesNet, while reliably detecting 110,866 true positives. Operator-Entropy JEPA similarly achieves 36.21% precision with 91,651 true positives and 161,469 false positives.

The Point-Adjustment Mirage in Deep Baselines: Comparing strict Point-F1 against Point-Adjusted F1 (PA-F1) in Table V exposes the empirical vulnerability of prior benchmarking protocols. Under PA-F1, Anomaly Transformer (0.2350), DCdetector (0.2602), and NCAD (0.2376) appear competitive with TranAD (0.2689) and TimesNet (0.2373). However, as soon as point adjustment is stripped and models are held to pure inductive EVT calibration, their actual point-level Micro-F1 collapses to 0.0116–0.0258. This demonstrates that the high scores historically reported for reconstructive and contrastive deep models were largely artifacts of point adjustment masking extreme false-alarm rates and missed detections.

Continuous Ranking (PR-AUC) vs. Calibrated Decisions (Point-F1): Comparing threshold-free continuous ranking against calibrated decisions reveals an important empirical insight. In PR-AUC (Average Precision), TranAD (0.2612) and TimesNet (0.2585) achieve top aggregate ranking scores, outperforming JEPA models on 5 of the 11 datasets (Daphnet, Genesis, CICIDS, Metro, and Room Occupancy; see Table VII). Under an oracle test threshold sweep evaluated on ground-truth labels (Oracle PA-F1), all six architectures achieve comparable performance within a tight band: Reynolds-Stress (0.5343), Operator-Entropy (0.5163), TimesNet (0.4886), TS-JEPA (0.4820), TranAD (0.4816), and Potential-Flow (0.4620). This confirms that reconstructive representations retain strong continuous anomaly ranking fidelity.

However, when automated extreme-value thresholding (SPOT) is applied without ground-truth labels, reconstructive models suffer severe threshold calibration collapse: TimesNet plunges to a Micro-F1 of 0.0647 (generating 1,047,144 false positives, flagging nearly 40% of nominal points), and TranAD drops to 0.1072 (452,621 false alarms). This collapse does not stem from score aggregation or thresholding discrepan-

TABLE V

MULTI-TIER BENCHMARK ON CORE MULTI-SENSOR DATASETS (45 STRICTLY MATCHED CHANNELS ACROSS 11 DIVERSE MODALITIES). WE REPORT UNADJUSTED POINT-F1 ACROSS THREE AGGREGATION REGIMES: UNWEIGHTED DATASET-MACRO (EQUAL $1/N$ DATASET WEIGHT), POOLED MICRO-F1 ($\sum TP, \sum FP, \sum FN$), AND CHANNEL-MACRO F1, ALONGSIDE THRESHOLD-INDEPENDENT PR-AUC AND POINT-ADJUSTED F1 (PA-F1). THE BEST RESULT IN EACH COLUMN IS IN **BOLD**, AND THE SECOND BEST IS UNDERLINED. ALL MODELS EVALUATED UNDER IDENTICAL PURE INDUCTIVE EVT CALIBRATION WITHOUT TRANSDUCTIVE TEST-TAIL CLAMPS.

Architecture	Physical Inductive Bias / Formulation	Strict Point-Level Evaluation (Unadjusted)			Ranking	Point-Adjusted (Literature Ref.)	
		DS-Macro $\uparrow$	Micro-F1 $\uparrow$	Chan-Macro $\uparrow$	PR-AUC $\uparrow$	DS-Macro $\uparrow$	Chan-Macro $\uparrow$
<i>Proposed Physics-Inspired Geometric JEPA Triad</i>							
Operator-Entropy JEPA	Koopman Spectral Entropy ($\mathcal{H}_K$)	0.2066	0.2806	0.1743	0.2541	0.3320	0.3205
Reynolds-Stress JEPA	Hydrodynamic Stress Tensor ($\mathbf{R}$)	0.1544	0.1175	<u>0.1532</u>	0.2406	0.3153	<u>0.3226</u>
Potential-Flow JEPA	Curl-Free Potential Field ($-\nabla\Phi$)	0.1147	0.3458	0.1259	0.1978	0.3053	0.3436
<i>Ablation Control (Unconstrained Latent Prediction)</i>							
Vanilla TS-JEPA	Latent Prediction (Unconstrained)	<u>0.1838</u>	<u>0.2860</u>	0.1442	0.2344	0.3449	0.3066
<i>Published State-of-the-Art Deep Baselines (Unified Reimplementations)</i>							
TranAD (VLDB 2022)	Reconstructive Transformer + Adversarial	0.1609	0.1072	0.1319	0.2612	0.2689	0.2537
TimesNet (ICLR 2023)	Multi-Period 2D Temporal Inception	0.1443	0.0647	0.1331	<u>0.2585</u>	0.2373	0.2072
Anomaly Transformer (ICLR 2022)	Association Discrepancy (Prior vs. Series)	0.0845	0.0228	0.0576	—	0.2350	0.2228
DCdetector (KDD 2023)	Dual Attention Contrastive Discrepancy	0.0782	0.0258	0.0697	—	0.2602	0.2406
NCAD (KDD 2021)	Contextual Contrastive Representation	0.0493	0.0116	0.0362	—	0.2376	0.1716
<i>Theoretical Reference Limit</i>							
All-Positive Trivial Predictor	Constant Positive ($\hat{y} \equiv 1$)	0.1893	0.2372	0.1942	—	—	—




Fig. 2. Benchmark evaluation of the proposed Physics-Inspired JEPA Triad (operator_entropy_jepa, reynolds_stress_jepa, potential_flow_jepa) and unconstrained control (ts_jepa) against published deep learning baselines across matched telemetry channels.

cies: all baselines and JEPA variants are evaluated under the exact same 64-step suspect region mean aggregation, identical window-to-point temporal reduction, and identical kurtosis-adaptive EVT calibrator ($q = 10^{-3}$, $u = 0.98$). Instead, observation-space residuals exhibit high temporal variance and sensor jitter, causing test residuals to systematically violate the Generalized Pareto Distribution (GPD) scale estimated from nominal training splits. In contrast, predicting in a regularized latent representation space ($\mathbb{R}^{32}$) filters out high-frequency observation noise, yielding better-conditioned residual distributions where TS-JEPA, Operator-Entropy, and Potential-Flow maintain Micro-F1 scores of 0.2806–0.3458 (while Reynolds-Stress achieves 0.1175) while generating 2.4-fold to 8.0-fold fewer false alarms compared to reconstructive baselines (5.6-

fold to 8.0-fold fewer compared to TimesNet). Nonetheless, with empirical false positive rates of approximately 5%–6% on nominal points, even latent architectures overshoot the theoretical design risk ($q = 10^{-3} = 0.1\%$) by roughly 50-fold, demonstrating that real-world non-stationary telemetry residuals continue to present heavy-tailed calibration challenges.

Statistical Significance and Cross-Stream Consistency: All 45 continuous streams are evaluated under a unified deterministic seed (42) with fixed initializations and identical inductive splits across all architectures. Conducting paired Wilcoxon signed-rank tests across the 45 streams confirms that channel-level Point-F1 varies systematically by modality: on periodic human kinematics (Daphnet), TimesNet’s 2D Fourier decomposition matches or edges JEPA ($p = 0.81$), whereas on

high-dimensional orbital telemetry (NASA SMAP and MSL) and municipal water loops (GECCO), Operator-Entropy and Potential-Flow JEPA establish dominant margins ($p < 0.05$). The primary statistical advantage of JEPA lies in its global stability across the entire 5.06-million-point corpus, where it achieves an 8-fold reduction in false alarm volume and a 5-fold gain in pooled Micro-F1 over TimesNet.

Role of Unconstrained Latent Prediction (TS-JEPA Control): Unconstrained TS-JEPA achieves 0.1838 Dataset-Macro F1, 0.2860 Micro-F1, and the highest Dataset-Macro PA-F1 (0.3449). Across all 18 metric comparisons in Table V, physics-inspired variants outperform the unconstrained TS-JEPA control in 9 instances; on strict unadjusted point-level metrics (excluding PA-F1), physics variants lead in 6 of 12 comparisons (50%). This confirms that unconstrained TS-JEPA is a remarkably strong baseline, and that predicting in latent representation space accounts for the foundational improvement in threshold calibration stability. The physical inductive priors do not provide uniform dominance across all metrics, but act as targeted regularizers offering domain-aligned advantages in specialized regimes (e.g., Operator-Entropy in cyclic and spectral telemetry, Potential-Flow in conservative mechanical subsystems).

Comparison to Analytical All-Positive Reference: In Table V, the analytical All-Positive baseline achieves a Dataset-Macro F1 of 0.1893 and a Channel-Macro F1 of 0.1942. On Dataset-Macro, Operator-Entropy JEPA (0.2066) is the sole model to surpass this trivial reference limit, while on Channel-Macro, all deep models fall below the constant alarm baseline due to severe channel-level class imbalance. Table VI details the per-dataset Point-F1 matrix alongside the analytical All-Positive baseline for each benchmark. Across the 11 individual datasets, Operator-Entropy JEPA outperforms the All-Positive baseline on 7 datasets (CallIt2, GECCO, GHL, Genesis, metro, room-occupancy, and SMAP), establishing substantial margins on municipal water infrastructure (GECCO: 0.4040 vs 0.0304) and satellite telemetry (SMAP: 0.4157 vs 0.2433), confirming that its macro performance does not arise from degenerating into a constant-positive detector.

B. Domain-Specialized Strengths of the Flagship Triad

Figure 3 illustrates the radar profile of the flagship models across five distinct telemetry domains, and Table VII details the per-dataset PR-AUC matrix across all 11 benchmark datasets.

Each model in the triad exhibits specialized mathematical inductive advantages:

1) *Operator-Entropy JEPA (Koopman Spectral Mechanics):* By regularizing the spectral entropy of the latent Koopman transition matrix during unsupervised representation training, `operator_entropy_jepa` encourages high-rank modal expressivity across non-linear dynamics. It achieves strong ranking across NASA spacecraft telemetry (MSL: 0.478 PR-AUC; SMAP: 0.467 PR-AUC) and industrial water infrastructure (GECCO: 0.351 PR-AUC vs Potential-Flow’s 0.015). However, on highly periodic human gait kinematics (Daphnet), it achieves 0.139 PR-AUC, trailing TimesNet (0.303) whose 2D




Fig. 3. PR-AUC radar comparison across five specialized physical domains: Spacecraft Telemetry (MSL, SMAP, swan), Biomechanics (Daphnet), Cyber-Physical Systems (GHL, GECCO, Genesis), Traffic & Environmental (CallIt2, metro, room-occupancy), and Network Cyber (CICIDS). The physics-inspired triad provides broad multi-domain coverage.

Fourier decomposition explicitly models multi-periodic harmonics. Similarly, on swan solar telemetry, Operator-Entropy (0.436) trails unconstrained TS-JEPA (0.651), confirming that the spectral regularizer acts as a targeted inductive bias rather than a universally dominant prior across all physical regimes.

2) *Reynolds-Stress JEPA (Continuum Mechanics & Turbulent Transport):* Grounded in hydrodynamic Reynolds decomposition ($z = \bar{z} + z'$), `reynolds_stress_jepa` aligns cross-channel fluctuation covariance between context and target horizons. This formulation is designed for systems characterized by coupled flows. On the SMAP satellite telemetry dataset, it achieves Rank 1 in PR-AUC (0.512), outperforming TimesNet (0.479) and TranAD (0.345). On SWAN solar telemetry, it achieves 0.522 PR-AUC and strong precision.

3) *Potential-Flow JEPA (Conservative Field Theory):* By constraining nominal dynamics to curl-free potential gradients ($\hat{v} = -\nabla\Phi$), `potential_flow_jepa` enforces conservative trajectory tracking. It achieves the highest aggregate Micro Point-F1 (0.3458) and highest Channel PA-F1 (0.3436) across evaluated models, with a precision of 45.98%. On NASA spacecraft subsystems (MSL), it scores the top PR-AUC of 0.551 (vs TimesNet’s 0.380 and TranAD’s 0.332).

4) *Role of Vanilla TS-JEPA (Ablation Control):* To isolate the contribution of the physical inductive biases from the predictive latent architecture itself, we evaluate `ts_jepa` as an unconstrained control. `ts_jepa` achieves 0.1838 Dataset-Macro Point-F1, 0.1442 Channel-Macro Point-F1, and 0.3066 Channel PA-F1. While substantially outperforming reconstructive baselines (+27.4% in Dataset-Macro F1 over TimesNet), it

TABLE VI

PER-DATASET UNADJUSTED POINT-F1 ACROSS ALL 11 TELEMETRY BENCHMARKS (45 MATCHED STREAMS) ALONGSIDE THE ANALYTICAL ALL-POSITIVE TRIVIAL PREDICTOR ($\hat{y} \equiv 1$). BEST RESULT FOR EACH DATASET IS IN **BOLD**, AND SECOND BEST IS UNDERLINED. THE BOTTOM ROW REPORTS UNWEIGHTED DATASET-MACRO AVERAGES MATCHING TABLE V.

Dataset	Physical Sensing Domain	Operator-Entropy	Reynolds-Stress	Potential-Flow	Vanilla TS-JEPA	TranAD	TimesNet	All-Positive
Callt2	Building Pedestrian Flow	0.0804	0.0000	0.0000	0.0000	0.1441	0.1524	0.0519
cicids	Network Intrusion Dynamics	<u>0.1576</u>	0.0335	0.1703	0.1182	0.0835	0.0884	0.3059
Daphnet	Wearable Parkinsonian Kinematics	0.0952	0.0327	0.0824	0.0974	<u>0.1401</u>	0.2158	0.2554
GECCO	Municipal Water Infrastructure	0.4040	<u>0.3843</u>	0.0045	0.2729	0.0304	0.0304	0.0304
Genesis	Pick-and-Place Automation Testbed	<u>0.0227</u>	0.0000	0.0000	0.0606	0.0000	0.0000	0.0062
GHL	Gasoil Heating Loop Dynamics	<u>0.0496</u>	0.1367	0.0183	0.0222	0.0154	0.0242	0.0259
metro	Interstate Highway Traffic (I-94)	<u>0.0038</u>	0.0031	0.0094	0.0034	0.0000	0.0000	0.0026
MSL	NASA Mars Science Laboratory	0.2439	0.1484	0.1365	0.1423	<u>0.1920</u>	0.1741	0.3929
room-occupancy	Smart Environmental Multi-Sensor	0.3643	0.2407	0.3395	0.3405	0.5808	<u>0.5282</u>	0.3303
SMAP	NASA Soil Moisture Active Passive	<u>0.4157</u>	0.4205	0.3764	0.4102	0.2614	0.2334	0.2433
swan	Solar Flare Magnetic Telemetry	<u>0.4349</u>	0.2984	0.1247	0.5540	0.3217	0.1402	0.4375
Dataset-Macro Mean	Across All 11 Benchmarks	0.2066	0.1544	0.1147	<u>0.1838</u>	0.1609	0.1443	0.1893

TABLE VII

PER-DATASET PRECISION-RECALL AREA UNDER THE CURVE (PR-AUC / AVERAGE PRECISION) ACROSS ALL 11 TELEMETRY BENCHMARKS (45 MATCHED CHANNELS). THE BEST RESULT FOR EACH DATASET IS IN **BOLD**, AND THE SECOND BEST IS UNDERLINED. DATASET-MACRO MEAN MATCHES THE AGGREGATE RANKING REPORTED IN TABLE V.

Dataset	Physical Telemetry Domain	Operator-Entropy	Reynolds-Stress	Potential-Flow	Vanilla TS-JEPA	TranAD	TimesNet
Callt2	Building Pedestrian Flow	0.097	0.028	0.028	0.030	<u>0.076</u>	<u>0.076</u>
Daphnet	Wearable Parkinsonian Kinematics	0.139	0.105	0.179	0.152	<u>0.242</u>	0.303
GECCO	Municipal Water Infrastructure	0.351	<u>0.318</u>	0.015	0.285	0.284	0.280
GHL	Gasoil Heating Loop Dynamics	<u>0.084</u>	0.170	0.021	0.052	0.017	0.014
Genesis	Pick-and-Place Automation Testbed	0.005	0.003	0.002	<u>0.010</u>	0.009	0.017
MSL	NASA Mars Science Laboratory	<u>0.478</u>	0.424	0.551	0.390	0.332	0.380
SMAP	NASA Soil Moisture Active Passive	0.467	0.512	<u>0.481</u>	0.440	0.345	0.479
cicids	Network Intrusion Dynamics	0.357	0.280	<u>0.365</u>	0.300	0.366	0.318
metro	Interstate Highway Traffic (I-94)	0.001	0.001	<u>0.002</u>	0.001	0.002	0.001
room-occupancy	Smart Environmental Multi-Sensor	0.380	0.283	0.233	0.267	<u>0.645</u>	0.659
swan	Solar Flare Magnetic Telemetry	0.436	0.522	0.299	0.651	<u>0.556</u>	0.315
Dataset-Macro Mean	Across All 11 Benchmarks	0.2541	0.2406	0.1978	0.2344	0.2612	<u>0.2585</u>

trails `operator_entropy_jepa` in Dataset-Macro Point-F1 (0.1838 vs 0.2066) and Channel Point-F1 (0.1442 vs 0.1743), and trails `potential_flow_jepa` in Micro-F1 (0.2860 vs 0.3458) and Channel PA-F1 (0.3066 vs 0.3436). This indicates that while latent representation prediction mitigates noise drowning, domain-aligned physical conservation laws provide beneficial inductive structure.

C. Impact of Architectural Enhancements

Ablation of the three internal inductive layers confirms their collective role in stabilizing high-dimensional inference. First, *Coordinate Saliency Gating* dynamically prioritizes high-variation informative channels in high-dimensional telemetry ($K \geq 25$ in NASA SMAP and MSL), preventing localized subsystem faults from being diluted by quiescent ambient sensors without requiring manual channel sub-selection. Second, *Littlewood-Paley Dyadic Wavelet Shells* decouple fast transient spikes from low-frequency baseline drifts within the encoder; on GECCO (water infrastructure exhibiting multi-hour operational demand variations alongside sharp chemical contaminant inflows and pipe pressure drops), dyadic filtering improved `operator_entropy_jepa` performance from 0.1429 to 0.4040 Point-F1. Third, *Kurtosis-Adaptive Extreme Value Theory* modulates the tail threshold via excess kurtosis γ_2 , dynamically raising decision margins during heavy-tailed volatile bursts while sharpening sensitivity on calm channels to suppress streaming false alarms.

D. Operational Incident Evaluation: Event Metrics and Streaming Adaptation

While unadjusted Point-F1 prevents score inflation from naive point adjustment, point-level scoring can penalize slight temporal boundary offsets in operational telemetry. In physical systems, an anomaly represents a continuous operational incident; if an algorithm detects the failure with a minor phase delay or slight boundary misalignment, point-level metrics penalize every misaligned point as a false positive or false negative, depressing nominal precision even when incidents are detected.

To evaluate operational incident detection, we analyze model performance under three complementary criteria: (i) strict unadjusted Point-F1, (ii) incident-level Event-F1 ($P_{\text{event}}, R_{\text{event}}$), and (iii) Tatbul et al. (NeurIPS 2018) Range-Aware F1 ($F1_{\text{range}}$) evaluated with balanced existence and overlap reward ($\alpha = 0.5$), uniform positional overlap, and reciprocal cardinality factor. Detections operate with an operational lookahead buffer of $S = 64$ timesteps.

Table VIII examines operational trade-offs across three representative continuous telemetry streams: Room Occupancy (stream: default), NASA SMAP (stream: P-3), and Callt2 (stream: traffic). Unlike Table VI which reports aggregate dataset-macro averages across all constituent streams, this study isolates individual continuous streams to contrast static calibration with dynamic adaptive tracking under three operational configurations:

1) **Standard Configuration:** Evaluates centered middle-

window alignment ($t+S/2$) with a 5-step moving average filter and static EVT calibration fit to training residuals.

- 2) **Sharp Configuration:** Evaluates dense un-lagged scoring (t) with a minimal 3-step moving average filter and static EVT calibration.
- 3) **Adaptive Configuration:** Evaluates dense un-lagged scoring (t) with a 3-step filter and dynamic rolling SPOT thresholding to track non-stationary baseline drift.

As shown in Table VIII, evaluating anomaly detection at the incident level reveals three operational characteristics:

Incident Sensitivity vs. False Alarm Volume: On Room Occupancy (default, 11 ground-truth events, $\pi = 20.0\%$ base rate) and SMAP (P-3, 1 continuous event), static EVT thresholding achieves 100% Event Recall across all tested configurations. However, static calibration generates over 4,000 point-level false alarms on Room Occupancy due to unmodeled diurnal baseline drift, yielding a 20.0% point precision that matches the background anomaly base rate.

Streaming Adaptation Pareto Trade-off: When dynamic rolling SPOT thresholding is applied to track baseline drift, false alarms decrease from 4,705 to 223 points on Room Occupancy (a $21\times$ reduction) and precision increases to 29.9%. However, this introduces an operational trade-off: event recall drops from 100% to 63.6% (7 of 11 events detected), demonstrating the classic operating curve between alarm fatigue and fault sensitivity on non-stationary streams.

Boundary Alignment in Operational Telemetry: On SMAP (P-3), both tabulated JEPAs capture the single continuous thruster fault segment (100% event recall). On Callt2 (traffic, 21 rapid pedestrian burst events), models detect between 4.8% and 23.8% of events under dynamic SPOT, demonstrating that short intermittent pulses remain challenging for fixed-horizon window architectures.

E. Operational Latency: Strictly Causal Online Detection vs. Buffered Horizons

Standard benchmark evaluations of sliding-window models evaluate window discrepancy over a target suspect region S , incurring an effective lookahead delay of S steps ($S = 64$ in our core benchmark, or 1.06 seconds at 60 Hz). In mission-critical industrial operations, anomaly detection often requires strictly causal latency constraints where no future telemetry is observed ($\tau = 0$ lookahead).

To quantify the operational trade-off between lookahead buffering and strictly causal detection, Table IX reports performance across five matched telemetry streams (NASA SMAP P-3, SMD machine-1-2, NASA MSL M-1, Daphnet S01R01E1, and GECCO water_quality) under two operational protocols:

- 1) **Strictly Causal Online Detection (0-Lookahead Delay):** Window scores are assigned strictly to the final observed point t , enforcing zero future information leakage.
- 2) **Buffered Window Detection (64-Lookahead Delay):** Window scores are assigned across the 64-step target suspect horizon, incurring a 64-step lookahead latency before score aggregation.

Under strictly causal online streaming, TS-JEPA maintains an empirical false-positive rate of 4.18% ($\text{FAR}_{1k} = 41.76$ false alarms per 1,000 points), generating an average of 1,685.6 false alarms across streams. In contrast, TimesNet and TranAD generate 16,952.0 and 16,637.6 false alarms (empirical FPRs of 30.06% and 20.81%, respectively)—over $10\times$ the false alarm volume of TS-JEPA. On GECCO municipal water telemetry (113,000 test points), TimesNet and TranAD trigger over 81,500 false alarms (99.6% FPR), whereas TS-JEPA triggers zero false alarms. When a 64-step lookahead buffer is permitted, TS-JEPA's mean detection delay decreases from 130.3 to 89.6 steps ($2.6\times$ faster than TimesNet's 235.3 steps).

F. Empirical Validation of Physical Inductive Formulations

To address whether the proposed mathematical regularizers provide demonstrable inductive properties rather than formal metaphors, we conduct controlled empirical validations for each component of the triad.

1) **Operator-Entropy: Anti-Collapse Spectral Rank Preservation:** A failure mode in joint-embedding architectures is dimensional collapse, where the learned transition operator projects representations onto a low-dimensional subspace. In Operator-Entropy JEPAs, the transition operator $K \in \mathbb{R}^{D \times D}$ ($D = 32$) is regularized via Von Neumann spectral entropy $S(\rho) = -\text{Tr}(\rho \log \rho)$ on the normalized Gram matrix $\rho = KK^T / \text{Tr}(KK^T)$ alongside a logarithmic energy barrier.

We evaluate the effective rank using the RankMe metric ($\exp(\mathcal{H})$ where $\mathcal{H}$ is the spectral entropy of singular values), representation covariance conditioning, and downstream calibrated Point-F1 across five telemetry streams. Table X compares Operator-Entropy JEPAs against unconstrained Vanilla TS-JEPAs, non-contrastive VICReg TS-JEPAs, and Representation-Entropy JEPAs (entropy applied directly to representation covariance rather than the transition operator).

As shown in Table X, unconstrained Vanilla TS-JEPAs, VICReg, and Representation-Entropy collapse across 7–8 singular modes, yielding an effective operator rank of 24.3–24.8. In contrast, Operator-Entropy JEPAs achieve an effective operator rank of 31.996 out of a theoretical maximum of 32.0, demonstrating that the spectral entropy regularizer strictly prevents modal collapse in the transition dynamics. This modal preservation yields a Point-F1 of 0.0710 ($2.35\times$ higher than Vanilla TS-JEPAs' 0.0301 and $16.5\times$ higher than VICReg's 0.0043).

2) **Reynolds-Stress: Covariance Alignment vs. Marginal Perturbations:** Reynolds-Stress JEPAs align second-order latent fluctuation statistics ($\Sigma_{z'} = \frac{1}{B} \sum z'(z')^T$) between predicted and target horizons. To test whether this loss specifically detects cross-channel covariance breakdowns versus marginal variance shifts, we construct a controlled synthetic benchmark with two isolated anomaly regimes:

- 1) **Pure Cross-Channel Covariance Anomaly:** Marginal distributions remain identical ($\mu = 0, \sigma^2 = 1$), while cross-channel correlation $\rho(x_1, x_2)$ flips from $+0.85$ to -0.85 .
- 2) **Pure Marginal Variance Anomaly:** Cross-channel correlation is preserved ($\rho = +0.85$), while marginal variance scales by $3.0\times$.

TABLE VIII

OPERATIONAL INCIDENT DETECTION COMPARISON ACROSS REPRESENTATIVE INDIVIDUAL TELEMETRY STREAMS: ROOM OCCUPANCY (STREAM: `DEFAULT`, 11 GROUND-TRUTH EVENTS), NASA SMAP (STREAM: `P-3`, 1 CONTINUOUS THRUSTER FAULT), AND CALIT2 (STREAM: `TRAFFIC`, 21 PEDESTRIAN BURST EVENTS). WHILE TABLE VI REPORTS DATASET-MACRO AVERAGES ACROSS ALL CHANNELS, THIS TABLE ISOLATES INDIVIDUAL CONTINUOUS STREAMS TO STUDY OPERATIONAL TRADE-OFFS AMONG STANDARD (MIDDLE, STATIC EVT), SHARP (DENSE, STATIC EVT), AND ADAPTIVE (DENSE, DYNAMIC SPOT) INFERENCE MODES.

Evaluated Stream	Operational Configuration	Pt-F1	Pt-Prec	Evt-F1	Evt-Prec	Evt-Rec	Rng-F1	False Pos.
<i>Room Occupancy: Stream default (Environmental Sensing, 11 Ground-Truth Incidents)</i>								
operator_entropy	Standard (Middle, Static EVT)	0.3728	22.9%	0.5238	35.5%	100%	0.3708	3,842
operator_entropy	Sharp (Dense, Static EVT)	0.3329	20.0%	0.9167	84.6%	100%	0.3329	4,705
operator_entropy	Adaptive (Dense, Dynamic SPOT)	0.1273	29.9%	0.6364	63.6%	63.6%	0.3699	223 ($\downarrow 21\times$)
potential_flow	Sharp (Dense, Static EVT)	0.3431	20.7%	0.9167	84.6%	100%	0.3431	4,496
potential_flow	Adaptive (Dense, Dynamic SPOT)	0.0441	10.2%	0.5714	60.0%	54.5%	0.1636	291 ($\downarrow 15\times$)
<i>NASA SMAP: Stream P-3 (Spacecraft Telemetry, 1 Continuous Fault Segment)</i>								
operator_entropy	Standard (Middle, Static EVT)	0.0087	14.6%	0.1818	10.0%	100%	0.2266	35
operator_entropy	Sharp (Dense, Static EVT)	0.1274	27.3%	0.1538	8.3%	100%	0.3634	295
potential_flow	Adaptive (Dense, Dynamic SPOT)	0.0110	6.9%	0.3333	20.0%	100%	0.1213	108
<i>Callit2: Stream traffic (Pedestrian Entrance Flow, 21 Transient Event Clusters)</i>								
operator_entropy	Standard (Middle, Static EVT)	0.0131	1.9%	0.0588	7.7%	4.8%	0.0216	53
operator_entropy	Adaptive (Dense, Dynamic SPOT)	0.1368	8.7%	0.3226	50.0%	23.8%	0.1274	357
potential_flow	Adaptive (Dense, Dynamic SPOT)	0.0903	14.3%	0.0909	100.0%	4.8%	0.0714	42

TABLE IX

OPERATIONAL LATENCY AND DETECTION FIDELITY UNDER STRICTLY CAUSAL (0-LOOKAHEAD DELAY) VERSUS BUFFERED WINDOW (64-LOOKAHEAD DELAY) INFERENCE ACROSS 5 REPRESENTATIVE TELEMETRY STREAMS. TS-JEPA ACHIEVES AN EMPIRICAL FALSE-POSITIVE RATE OF 4.18% ($\text{FAR}_{1k} = 41.76$ FALSE ALARMS PER 1K POINTS) UNDER CAUSAL ONLINE STREAMING, COMPARED TO 30.06% FOR TIMESNET AND 20.81% FOR TRANAD.

Inference Setting	Model	Empirical FPR	FAR_{1k} (/1k pts)	False Positives	Point-F1	Precision	Mean Delay (Δt)	Recall ($\tau \leq 64$)
Strictly Causal (0-Lookahead)	TS-JEPA	0.0418	41.76	1,685.6	0.0688	0.1581	130.3	0.1071
	TimesNet	0.3006	300.63	16,952.0	0.0774	0.1525	251.2	0.2286
	TranAD	0.2081	208.13	16,637.6	0.0881	0.1255	43.2	0.3286
Buffered Window (64-Lookahead)	TS-JEPA	0.1514	151.39	2,696.0	0.0431	0.0355	89.6	0.1357
	TimesNet	0.2800	280.03	16,916.8	0.0807	0.1404	235.3	0.2571
	TranAD	0.1998	199.82	16,339.6	0.0197	0.2284	32.8	0.2786

TABLE X

EMPIRICAL SPECTRAL DIAGNOSTICS OF LEARNED TRANSITION OPERATORS AND REPRESENTATIONS ACROSS FIVE BENCHMARK TELEMETRY STREAMS ($D = 32$). OPERATOR-ENTROPY PRESERVES FULL THEORETICAL RANK ($\exp(\mathcal{H}) = 31.996$ OUT OF 32.0), PREVENTING DIMENSIONAL COLLAPSE.

Model Architecture	Operator Rank (RankMe)	Rep. Rank	Point-F1	Empirical FPR
Operator-Entropy JEPA	31.996	3.564	0.0710	0.1497
Representation-Entropy JEPA	24.647	7.579	0.0479	0.0442
Vanilla TS-JEPA (Unconstrained)	24.807	3.993	0.0301	0.0371
VICReg TS-JEPA	24.339	5.881	0.0043	0.0273

TABLE XI

CONTROLLED COVARIANCE VS. MARGINAL ANOMALY BENCHMARK. REYNOLDS-STRESS JEPA IMPROVES MARGINAL VARIANCE DETECTION BY +76% IN PR-AUC OVER UNCONSTRAINED TS-JEPA (0.5838 vs. 0.3313).

Model Architecture	Pure Covariance Shift		Pure Marginal Shift	
	PR-AUC	Point-F1	PR-AUC	Point-F1
Reynolds-Stress JEPA	0.0498	0.0000	0.5838	0.0593
TS-JEPA (Unconstrained)	0.0409	0.0000	0.3313	0.0000
TimesNet	0.9313	0.6558	0.9998	0.6267
TranAD	0.6761	0.7246	0.9343	0.6560

Table XI demonstrates that on pure marginal variance shifts, Reynolds-Stress JEPA achieves a PR-AUC of 0.5838 com-

pared to 0.3313 for unconstrained TS-JEPA (a +76% relative improvement), with an empirical FPR of 0.09%. On pure covariance shifts without marginal change, both latent models show reduced sensitivity (0.0409–0.0498 PR-AUC) relative to observation-space reconstructors (0.6761–0.9313), because temporal global pooling in standard 1D encoders averages out instantaneous phase alignments unless explicit cross-channel attention or relational graphs are utilized.

3) *Potential-Flow & Helmholtz Decomposition: Conservative vs. Circulating Dynamics:* Potential-Flow JEPA parameterizes transition velocities as the negative gradient of a scalar potential ($\hat{v} = -\nabla\Phi$). While curl-free gradient fields enforce exact energy conservation ($\oint \hat{v} \cdot d\mathbf{r} = 0$), non-conservative systems often exhibit limit cycles with non-zero circulation.

To evaluate this inductive boundary, we compare Potential-Flow JEPA, Helmholtz-JEPA (which decomposes velocity into scalar potential plus solenoidal rotational circulation: $\hat{v} = -\nabla\Phi + \text{curl}(A)$), and unconstrained TS-JEPA across two controlled dynamical systems:

- 1) **Conservative Duffing Oscillator:** Hamiltonian phase space with symplectic energy conservation; non-conservative damping dissipation injected as anomalies.
- 2) **Non-Conservative Van der Pol Oscillator:** Limit-cycle attractor with continuous phase vorticity; non-linear per-

TABLE XII

DYNAMICAL SYSTEM VALIDATION COMPARING CURL-FREE POTENTIAL-FLOW ($\hat{v} = -\nabla\Phi$), HELMHOLTZ DECOMPOSITION ($\hat{v} = -\nabla\Phi + \text{CURL}(A)$), AND UNCONSTRAINED TS-JEPA ON CONSERVATIVE HAMILTONIAN VS. NON-CONSERVATIVE LIMIT-CYCLE PHASE SPACES.

Model Architecture	Conservative Duffing		Non-Conservative Van der Pol	
	PR-AUC	Point-F1	PR-AUC	Point-F1
Helmholtz JEPA ($\nabla\Phi + \text{curl}(A)$)	0.5985	0.7815	0.5250	0.7574
TS-JEPA (Unconstrained)	0.5855	0.7583	0.4577	0.6660
Potential-Flow JEPA ($-\nabla\Phi$)	0.1701	0.2036	0.3106	0.3548

turbation injected as anomalies.

Table XII confirms that decomposing the transition field into potential and curl components (Helmholtz JEPA) achieves the highest detection accuracy across both conservative (0.5985 PR-AUC, 0.7815 Point-F1) and non-conservative limit-cycle regimes (0.5250 PR-AUC, 0.7574 Point-F1), outperforming unconstrained TS-JEPA. Pure curl-free Potential-Flow achieves a zero false-positive rate on Duffing (0.0000%) and 0.89% on Van der Pol, but lower recall due to the mathematical inability of pure gradient fields to represent closed circulating phase trajectories.

4) *Reconciling Metric Profiles: Micro-F1, Macro-F1, and Continuous Ranking:* A key empirical pattern in Table V is that Potential-Flow JEPA achieves the highest Pooled Micro-F1 (0.3458), yet exhibits lower Dataset-Macro Point-F1 (0.1147) and PR-AUC (0.1978) compared to reconstructive baselines. This apparent contrast stems directly from aggregation weights across heterogeneous telemetry streams:

- 1) **Micro-F1 Weighting by Point Volume:** Pooled Micro-F1 evaluates all 5.06 million test points globally. High-dimensional telemetry streams (e.g., NASA SMAP and MSL with 25–55 channels and long continuous fault segments) contribute the vast majority of positive anomaly points. On these long events, Potential-Flow’s curl-free scalar potential acts as an ultra-conservative detector that suppresses ambient measurement jitter, achieving 45.98% precision (130,237 false positives vs. 1,047,144 for TimesNet).
- 2) **Macro-F1 Sensitivity to Transient Events:** Dataset-Macro Point-F1 weights each of the 11 datasets equally. On short intermittent spike datasets (e.g., Callt2, Metro) or multi-periodic gait kinematics (Daphnet), curl-free gradient dynamics ($\hat{v} = -\nabla\Phi$) lack the rotational degrees of freedom required to track circulating limit-cycle orbits (as confirmed in Table XII), causing lower event recall on periodic streams and depressing unweighted dataset-macro averages.
- 3) **Continuous Ranking vs. Calibrated Thresholds:** Reconstructive models like TimesNet and TranAD achieve higher PR-AUC (0.2585–0.2612) because their raw reconstruction errors scale continuously with signal amplitude, providing smooth continuous ranking across test points. However, because observation residuals exhibit severe tail heaviness ($\gamma_2 > 50$), automated extreme value thresholding causes catastrophic false-alarm flooding at test time.

G. Automated Threshold Calibration: EVT, Kurtosis Adaptation, and Nominal Reliability

A foundational question in unsupervised anomaly detection is whether automated extreme-value thresholding adheres to theoretical design risk parameters, or whether thresholding collapse arises from residual tail heaviness. Table XIII reports a controlled evaluation of calibration procedures across five benchmark streams, measuring empirical False Positive Rate (FPR) on held-out nominal validation sequences versus test detection fidelity.

Three empirical conclusions emerge from Table XIII: First, on held-out nominal validation data, empirical FPR closely tracks the specified design risk: setting $q = 10^{-4}$ yields an empirical nominal FPR of 0.01% (0.0001), setting $q = 10^{-3}$ yields 0.19% (0.0019), and setting $q = 10^{-2}$ yields 0.74% (0.0074). This confirms that the Generalized Pareto Distribution accurately models the upper tail of nominal TS-JEPA residuals under uncorrupted stationary conditions. Second, on test telemetry, all static thresholding methods exhibit a higher empirical FPR (≈ 17 –18%), reflecting non-stationary baseline shifts between training and test sequences. Third, internal GPD parameter modulation (p_{boost} and q_{decay}) achieves the lowest nominal calibration error (0.16% nominal FPR) and stable test Point-F1 (0.1028). In contrast, unconstrained online rolling SPOT without drift suppression suffers substantial degradation (test FPR 44.46%, precision 3.22%), as anomalous bursts pollute the rolling buffer.

H. Computational Profile: Complexity, Latency, and Memory Footprint

To assess operational suitability for edge and real-time streaming deployments, Table XIV benchmarks parameter complexity, GPU memory footprint (VRAM), forward inference latency, and streaming throughput across all evaluated architectures on an NVIDIA RTX 5090 GPU (input dimension $K = 25$, batch size 32, context $C = 256$, suspect $S = 64$).

As shown in Table XIV, Operator-Entropy JEPA has fewer trainable parameters (285,733) than unconstrained TS-JEPA (289,061), because its explicit Koopman transition matrix $K \in \mathbb{R}^{32 \times 32}$ (1,024 weights) replaces the multi-layer MLP predictor. All JEPA variants maintain an ultra-compact peak GPU memory footprint of 183.8–185.3 MB, consuming nearly $7\times$ less memory than TranAD (1,280.0 MB), whose multi-head self-attention attention maps scale quadratically with sequence length. Forward inference latency for the JEPA family ranges between 0.60–0.84 ms per window (1,184–1,647 windows per second), well within sub-millisecond operational streaming deadlines.

VIII. CONCLUSION AND FUTURE WORK

In this work, we presented a Physics-Inspired Joint-Embedding Predictive Architecture (PI-JEPA) framework for unsupervised anomaly detection in multivariate sensor telemetry. By replacing raw observation-space reconstruction with forward prediction in latent representation space, our approach substantially mitigates the threshold calibration collapse that causes reconstructive baselines (TimesNet, TranAD) to flood

TABLE XIII

EMPIRICAL CALIBRATION BENCHMARK COMPARING NON-EVT BASELINES, STANDARD STATIC SPOT, KURTOSIS-ADAPTIVE EVT, AND ROLLING SPOT ON HELD-OUT NOMINAL VALIDATION SPLITS VERSUS TEST PERFORMANCE. ON HELD-OUT NOMINAL SEQUENCES, EMPIRICAL FPR CLOSELY ADHERES TO THEORETICAL DESIGN RISK ($q = 10^{-4} \rightarrow 0.01\%$, $q = 10^{-3} \rightarrow 0.19\%$, $q = 10^{-2} \rightarrow 0.74\%$).

Calibration Category	Thresholding Method	Nominal Val FPR	Test FPR	Point-F1	Precision	Recall
Non-EVT Baselines	Max-of-Validation	0.0000	0.1774	0.1029	0.1466	0.2651
	Empirical 99.5th Percentile	0.0054	0.1837	0.1028	0.1451	0.2665
	Gaussian 3-Sigma ($\mu + 3\sigma$)	0.0088	0.1879	0.1016	0.1435	0.2748
Static SPOT (Varying Design Risk q , $u = 0.98$)	Static SPOT ($q = 10^{-4}$)	0.0001	0.1733	0.1028	0.1464	0.2641
	Static SPOT ($q = 10^{-3}$)	0.0019	0.1764	0.1030	0.1462	0.2655
	Static SPOT ($q = 10^{-2}$)	0.0074	0.1886	0.1022	0.1445	0.2672
Kurtosis-Adaptive EVT ($u = 0.98$, $q = 10^{-3}$)	Internal GPD Adaptive (p_{boost} , q_{decay})	0.0016	0.1756	0.1028	0.1462	0.2651
	Multiplicative Rule ($\beta = 0.083$)	0.2778	0.2241	0.0973	0.1481	0.2709
	Multiplicative Rule ($\beta = 0.167$)	0.3046	0.2686	0.0964	0.1585	0.2767
	Moors Robust Quantile Kurtosis ($\beta = 0.167$)	0.1935	0.1848	0.1024	0.1465	0.2689
Online Streaming	Rolling Dynamic SPOT (Buffer 1,000 steps)	N/A	0.4446	0.0565	0.0322	0.4106

TABLE XIV

COMPUTATIONAL PROFILE, MEMORY CONSUMPTION, LATENCY, AND THROUGHPUT BENCHMARKED ON AN NVIDIA RTX 5090 GPU ($K = 25$, WINDOW SIZE 320).

Model Architecture	Parameters	VRAM (MB)	Latency (ms/win)	Infer (win/s)	Train (win/s)
TS-JEPA	289,061	183.8	0.6072	1,646.8	612.8
Operator-Entropy JEPA	285,733	183.9	0.6939	1,441.0	553.2
Reynolds-Stress JEPA	295,953	184.0	0.7456	1,341.2	535.8
Potential-Flow JEPA	338,117	184.4	0.8445	1,184.2	400.2
Helmholtz JEPA	416,391	185.3	0.8378	1,193.6	526.6
TimesNet	648,553	279.6	0.2908	3,439.4	1,161.4
TranAD	205,746	1,280.0	0.1897	5,272.2	1,554.4

continuous streams with hundreds of thousands of false alarms under ambient sensor noise.

We embedded physical domain principles directly into latent representations through three complementary formulations: Operator-Entropy JEPA (regularizing latent Koopman transition matrices via spectral entropy during unsupervised representation training to prevent modal collapse), Reynolds-Stress JEPA (aligning cross-channel second-moment fluctuation covariance tensors inspired by hydrodynamic Reynolds decomposition), and Potential-Flow JEPA (enforcing conservative, irrotational trajectory dynamics via a scalar potential gradient). Combined with Causal Coordinate Saliency Gating, Littlewood-Paley Dyadic Wavelet Shells, and Kurtosis-Adaptive Extreme Value Theory tail calibration, these components stabilize streaming threshold calibration. Across 11 telemetry datasets (45 matched channels) evaluated under identical protocols, Operator-Entropy JEPA achieves the highest Dataset-Macro Point-F1 (0.2066, representing a 43.2% relative improvement over TimesNet and 28.4% over TranAD) and clears the analytical all-positive baseline (0.1893). In Pooled Micro-F1, Potential-Flow JEPA attains 0.3458 (over $3.2\times$ higher than TranAD and $5.3\times$ higher than TimesNet) while reducing false alarms by 8.0-fold.

Four concrete directions extend this work. First, extending conservative field formulations by incorporating parameter-controlled solenoidal vorticity components will allow modeling dissipative systems with limit-cycle attractors. Second, developing memory-efficient, test-time momentum adaptation will enable tracking orbital thermal cycles and environmental seasonal drift without catastrophic forgetting. Third, expanding

empirical evaluations across extensive multi-seed benchmarks will enable rigorous confidence interval reporting and rank-based statistical testing (such as Friedman-Nemenyi tests). Finally, quantizing PI-JEPA encoders into 8-bit integer (INT8) formats will enable deployment on radiation-hardened spacecraft microcontrollers and ultra-low-power wearable edge devices.

CODE AND DATA AVAILABILITY

To ensure reproducibility, all benchmark dataset extraction scripts, preprocessing pipelines, model checkpoints, PyTorch training code, fixed random seed configurations ($\text{seed} = 42$), and evaluation logs for all six evaluated architectures across the 11 telemetry datasets are open-sourced at an anonymized public repository: https://anonymous.4open.science/r/NCAD_CS-Reproducibility-8F21. All public datasets (NASA SMAP/MSL, Daphnet, CICIDS-2017, GHL, Callt2, GECCO, Genesis, Metro Traffic, and SWAN-SF) are accessible via their respective referenced public repositories.

REFERENCES

- [1] F. Meng, B. Yuan, and F. Gao, "Spacecraft telemetry anomaly detection based on deep learning: A review and benchmark," *IEEE Transactions on Aerospace and Electronic Systems*, vol. 56, no. 5, pp. 3863–3878, 2020.
- [2] R. Chalapathy and S. Chawla, "Deep learning for anomaly detection: A survey," *arXiv preprint arXiv:1901.03407*, 2019.
- [3] K. Hundman, V. Constantinou, C. Laporte, I. Colwell, and T. Soderstrom, "Detecting spacecraft anomalies using lstms and nonparametric dynamic thresholding," *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pp. 387–395, 2018.
- [4] M. P. Cuéllar, M. Santos, J. L. Alonso, M. Fabregas, and R. Farías, "Explainable anomaly detection in spacecraft telemetry," *Engineering Applications of Artificial Intelligence*, vol. 133, p. 108119, 2024.
- [5] J. Xu, H. Wu, J. Wang, and M. Long, "Anomaly transformer: Time series anomaly detection with association discrepancy," in *International Conference on Learning Representations (ICLR)*, 2022.
- [6] H. Wu, T. Hu, Y. Liu, H. Zhou, J. Wang, and M. Long, "Timesnet: Temporal 2d-variation modeling for general time series analysis," in *International Conference on Learning Representations (ICLR)*, 2023.
- [7] C. U. Carmona, F.-X. Aubert, V. Flunkert, and J. Gasthaus, "Neural contextual anomaly detection for time series," in *International Conference on Learning Representations (ICLR)*, 2022.
- [8] Y. LeCun, "A path towards autonomous machine intelligence," *OpenReview Preprint*, 2022.

- [9] M. Assran, Q. Duval, I. Misra, P. Bojanowski, P. Vincent, M. Rabbat, Y. LeCun, and N. Ballas, "Self-supervised learning from images with a joint-embedding predictive architecture," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023, pp. 15 619–15 629.
- [10] S. L. Brunton, J. L. Proctor, and J. N. Kutz, "Discovering governing equations from data by sparse identification of nonlinear dynamical systems," *Proceedings of the National Academy of Sciences*, vol. 113, no. 15, pp. 3932–3937, 2016.
- [11] R. T. Chen, Y. Rubanova, J. Bettencourt, and D. K. Duvenaud, "Neural ordinary differential equations," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 31, 2018, pp. 6571–6583.
- [12] S. L. Brunton, B. W. Brunton, J. L. Proctor, and J. N. Kutz, "Koopman invariant subspaces and finite-dimensional approximations of nonlinear dynamical systems," *Journal of Nonlinear Science*, vol. 26, pp. 1131–1162, 2016.
- [13] B. Lusch, J. N. Kutz, and S. L. Brunton, "Deep learning for universal linear embeddings of nonlinear dynamics," *Nature Communications*, vol. 9, no. 1, p. 4950, 2018.
- [14] Y. Liu, C. Li, J. Wang, and M. Long, "Koopman: Learning non-stationary time series dynamics with Koopman predictors," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 36, 2023, pp. 13 480–13 492.
- [15] S. Tuli, G. Casale, and N. R. Jennings, "Tranad: Deep transformer networks for anomaly detection in multivariate time series data," *Proceedings of the VLDB Endowment*, vol. 15, no. 6, pp. 1201–1214, 2022.
- [16] Y. Yang, C. Zhang, T. Zhou, Q. Gao, Q. Wen, and Y. Liang, "Dcdetector: Dual attention contrastive representation learning for time series anomaly detection," in *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2023, pp. 3033–3045.
- [17] A. Bardes, Q. Duval, C. Feichtenhofer, I. Laptev, D. Haziza, V. Khalidov, N. Ballas, and Y. LeCun, "V-jepa: Video joint-embedding predictive architecture," *arXiv preprint arXiv:2404.08471*, 2024.
- [18] J. Zbontar, L. Jing, I. Misra, Y. LeCun, and S. Deny, "Barlow twins: Self-supervised learning via redundancy reduction," in *International Conference on Machine Learning (ICML)*, 2021, pp. 12 310–12 320.
- [19] A. Ermolov, A. Siarohin, E. Sangineto, and N. Sebe, "Whitening for self-supervised representation learning," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, 2021, pp. 8315–8327.
- [20] A. Bardes, J. Ponce, and Y. LeCun, "Vicreg: Variance-invariance-covariance regularization for self-supervised learning," in *International Conference on Learning Representations (ICLR)*, 2022.
- [21] Z. Yue, Y. Wang, J. Duan, T. Yang, C. Huang, Y. Tong, and B. Xu, "Ts2vec: Towards universal representation of time series," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 36, no. 8, 2022, pp. 8980–8987.
- [22] E. Eldele, M. Ragab, Z. Chen, M. Wu, C. K. Kwok, X. Li, and C. Guan, "Time-series representation learning via temporal and contextual contrasting," in *Proceedings of the International Joint Conference on Artificial Intelligence (IJCAI)*, 2021, pp. 2352–2359.
- [23] K. Zhang, Q. Wen, C. Zhang, R. Cai, M. Wang, L. Liu, and L. Sun, "Self-supervised time series representation learning: A systematic survey," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 46, no. 5, pp. 3241–3261, 2024.
- [24] M. Raissi, P. Perdikaris, and G. E. Karniadakis, "Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations," *Journal of Computational Physics*, vol. 378, pp. 686–707, 2019.
- [25] S. Greydanus, M. Dzamba, and J. Yosinski, "Hamiltonian neural networks," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019, pp. 15 379–15 389.
- [26] M. Cranmer, S. Greydanus, S. Hoyer, P. Battaglia, D. Spergel, and S. Ho, "Lagrangian neural networks," *arXiv preprint arXiv:2003.04630*, 2020.
- [27] S. Wang, Y. Teng, and P. Perdikaris, "Understanding and mitigating gradient pathologies in physics-informed neural networks," *SIAM Journal on Scientific Computing*, vol. 43, no. 5, pp. A3055–A3081, 2021.
- [28] A. Siffer, P.-A. Fouque, A. Termier, and C. Largouet, "Anomaly detection in streams with extreme value theory," in *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2017, pp. 1067–1075.
- [29] M. Bächlin, M. Plotnik, D. Roggen, I. Maidan, J. M. Hausdorff, N. Giladi, and G. Tröster, "Wearable assistant for parkinson's disease patients with the freezing of gait symptom," *IEEE Transactions on Information Technology in Biomedicine*, vol. 14, no. 2, pp. 436–446, 2010.
- [30] A. Asuncion and D. Newman, "UCI Machine Learning Repository: Calit2 Building People Counts Data Set," <https://archive.ics.uci.edu/dataset/179/calit2+building+people+counts>, 2007.
- [31] S. Schmidl, P. Wenig, and T. Papenbrock, "Anomaly detection in time series: a comprehensive evaluation," *Proceedings of the VLDB Endowment*, vol. 15, no. 9, pp. 1779–1797, 2022.
- [32] Y. Zhou, C. Brif, and I. Lourentzou, "mTSBench: A Comprehensive Benchmark for Multi-Domain Time Series Anomaly Detection," *Transactions on Machine Learning Research (TMLR)*, 2026, arXiv:2506.21550.
- [33] R. A. Angryk, P. C. Martens, B. Aydin, D. Kempton, S. S. Mahajan, S. Basodi, A. Ahmadzadeh, S. Boubrahimi, S. M. Hamdi, M. A. Schuh *et al.*, "Multivariate time series dataset for space weather data analytics," *Scientific Data*, vol. 7, no. 1, p. 227, 2020.
- [34] L. M. Candanedo and V. Feldheim, "Accurate occupancy detection of an office room from light, temperature, humidity and co2 measurements using statistical learning models," *Energy and Buildings*, vol. 112, pp. 28–39, 2016.
- [35] S. Kim, K. Choi, J. Choi, B. Lee, and S. Yoon, "Towards a rigorous evaluation of time-series anomaly detection," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 36, no. 7, 2022, pp. 7194–7201.
- [36] R. Wu and E. Keogh, "Current time series anomaly detection benchmarks are flawed and are creating the illusion of progress," *IEEE Transactions on Knowledge and Data Engineering*, vol. 35, no. 3, pp. 2421–2434, 2023.